

WEEKLY MODEL COST BRIEF · PUBLIC

Vendors Are Walking Operators Up the Cost Scale

Week of Mon Sep 29 to Sun Oct 5, 2026 (America/Chicago)

Cheap seats keep disappearing. Named replacements cost more. Flagship stickers can fall while volume traffic is walked up the ladder.

Jerry W. Lambert, Sr.

Executive Briefing · October 2026

<https://jerrywlambert.com>

Week of Mon Sep 29 to Sun Oct 5, 2026 (America/Chicago)


Audience: operators and buyers of LLM inference

Chart: cost walk-up (see figure)


Cheap models keep disappearing — and what's left costs more

List input price, \$ per million tokens · Anthropic, OpenAI, xAI, Perplexity, Google, DeepSeek · official vendor pricing and our rate card only


A When the cheaper options go away



B What you pay next: grey = old price, red = where traffic lands



C Size of the step



Not every move goes up: grok-3 \$3.00 → grok-4.3 \$1.25 · deepseek-v3 \$0.27 → deepseek-flash \$0.15 · Claude Opus 4.1 \$15 → Opus 4.8 \$5 · Sonnet 4.5 \$3 → Sonnet 5.5 \$2 (Nov 30) · gpt-5.4-nano \$0.20 → gpt-6-luna \$0.10 (Apr 2027). Perplexity's Sonar API ended Sep 27; on the Agent API, Sonar input fell \$1 → \$0.25 but output rose \$1 → \$2.50, so it is shown on the timeline only. Claude Haiku 3's price was not captured, so it has no bar. Haiku 4.5: Anthropic lists retirement 'not sooner than Oct 15, 2026' and names no successor. Sonnet 5's cancelled rise still reached our book (\$3, Sep 10-25). DeepSeek peak = 01-04 and 06-10 UTC weekdays. Excluded: Sep 25 rate-card catch-up corrections (kimi-k2.6, qwen-3.7-plus, command-r-plus, deepseek-v4-pro), which were our book catching up, not vendor moves. Other vendors (Z.ai promo end, Moonshot, Mistral) are in the caption. Sources: vendor pricing and deprecation pages captured Sep 17 · Oct 5, 2026 (Gemini 2.0 prices from Google's Vertex AI page); sonar-reasoning and grok-4-fast old prices from our rate card. No aggregator prices.

The Walk-Up

Foundation vendors keep publishing new models. At the cheap end of each lineup, the pattern is consistent: a low-priced seat is retired or repriced, and the next place traffic can land costs more. Operators who stay on “whatever is cheapest this month” get marched up a ladder of list input prices, often +25% to +700% on the successor seat, while flagship and frontier sticker prices stay flat or fall.

That is the walk-up. It is not a claim that every model gets more expensive. It is a claim about **routing pressure**: when the bargain SKU disappears or jumps, default and failover traffic follows the named replacement, and that replacement's list rate is higher.

How to read the chart above:

- **Panel A.** One lane per priority vendor. Grey × marks a cheaper model retired; red ▲ marks a price step on the same model or a forced move to a costlier successor. Hollow markers are announced but not yet in effect; a grey triangle is a rise that was announced and then cancelled.
- **Panel B.** Linear \$/M input axis. Grey dots are the old seat; red dots are where traffic lands next.
- **Panel C.** The same step as a percentage. Solid bars are already in effect; hollow bars are still pending.

Vendors covered, in analysis priority order: **Anthropic, OpenAI, and xAI first**; then **Perplexity, Google (Gemini), and DeepSeek**. Other vendors appear only when they sharpen the same story without crowding those six.

What the Priority Vendors Show

Anthropic. Haiku 3.5 retired at \$0.80/M input; Haiku 4.5 at \$1.00 is the named replacement (+25%). Haiku 4.5 is still the cheapest Claude, with retirement “not sooner than Oct 15, 2026” and no named successor. If it goes, the next-cheapest current Claude is Sonnet 5 at \$2.00 (+100%). Separately, a planned Sonnet 5 rise from \$2 to \$3 on Sep 1 was **cancelled**; Anthropic now treats \$2/\$10 as the standard Sonnet 5 rate. That cancelled step is drawn grey on the chart so the walk-up claim does not lean on a hike that did not stick.

OpenAI. Four announced retirements force traffic onto higher list seats: gpt-5-nano → gpt-5.6-luna (\$0.05 → \$0.20, Dec 11), gpt-5-mini → gpt-5.6-terra (\$0.25 → \$2.00, Dec 11), gpt-4.1-nano → gpt-5.6-luna (\$0.10 → \$0.20, Oct 23), and o4-mini → gpt-5.6-terra (\$1.10 → \$2.00, Oct 23). All four replacements are OpenAI’s own.

xAI. grok-4-fast at \$0.20 retired May 15; traffic auto-redirects and bills at grok-4.3 rates (\$1.25, +525%).

Perplexity. sonar-reasoning → sonar-reasoning-pro was a clean +100% list step when the cheaper seat retired. The broader Sonar → Agent API migration is mixed (input down, output up, plus tool fees) and sits on the timeline rather than as a clean walk-up bar.

Google (Gemini). Gemini 2.0 Flash → 3.6 Flash and 2.0 Flash-Lite → 3.1 Flash-Lite were large steps after the June shutdowns. Separately, Gemini 3.6 / 3.7 / 3.8 Flash still sit on an intro rate of \$0.75/M input through Dec 31, 2026, then step to \$1.50 on Jan 1, 2027 (+100%, announced).

DeepSeek. Peak-hour pricing on deepseek-flash doubles the off-peak rate (\$0.15 → \$0.30 during named UTC weekday windows). That is a time-of-day walk-up rather than a model retirement.

Counter-Examples (Keep the Claim Honest)

The walk-up is concentrated at the **cheap end**. Flagship and successor prices are often flat or falling. The chart footer already calls these out:

- **xAI:** grok-3 fell from \$3 to \$1.25.
- **DeepSeek:** deepseek-v3 fell from \$0.27 to \$0.15.
- **Anthropic:** Opus 4.1 at \$15 → Opus 4.8 at \$5; Sonnet 4.5 at \$3 → Sonnet 5.5 at \$2 (from Nov 30).
- **OpenAI:** gpt-5.4-nano at \$0.20 → gpt-6-luna at \$0.10 (Apr 2027).

If you only watch frontier stickers, the market looks cheaper. If you watch where high-volume “fast / mini / nano / flash” traffic is forced to land after a retirement, the ladder goes up.

Last Week in This Frame (Sep 29 to Oct 5)

No first-party **list price change** on already-billed seats for the priority six during the week. The news that matters for operators is new public SKUs, intro ladders, and looming retirements, not a catalog dump.

Anthropic. Claude Sonnet 5.5 (`claude-sonnet-5-5`) went public at **\$2 / \$10 per million tokens**, the same sticker as Sonnet 5, with a retirement commitment not sooner than Sep 28, 2027. That is a **flat** successor at the mid-tier, consistent with the counter-example lane (Sonnet getting cheaper or staying put at the top of the volume band). Claude Mythos 5 / 5.1 appeared on pricing pages at \$10 / \$50 but remain invite-only (Project Glasswing); they are not a PAYG walk-up you can plan around.

OpenAI. A new billed id, **gpt-6.1-sol**, listed at **\$2 / \$10** with **cached input at \$0.10 (5%)**. The prior **gpt-6-sol** seat stays live on its own page at the same \$2 / \$10 list but **10% cached input (\$0.20)**. That is a new SKU, not a rename of the prior seat. More urgent for routing: on **Oct 23, 2026**, OpenAI shuts down **gpt-4.1-nano**, **o4-mini**, and **gpt-image-1**, with named replacements **gpt-5.6-luna**, **gpt-5.6-terra**, and **gpt-image-2.5**. Two of those text shutdowns are exactly the walk-up moves already drawn on the chart (+100% and +82% on list input). **gpt-5.6-cyber** remains Daybreak / approval-only and is not a general PAYG seat.

Google (Gemini). New public Flash-Lite seats: **gemini-3.5-flash-lite** at **\$0.30 / \$2.50** and **gemini-3.1-flash-lite** at **\$0.25 / \$1.50** (text rates). These sit below full Flash and do not cancel the Jan 1, 2027 intro → list step on 3.6 / 3.7 / 3.8 Flash. **gemini-2.5-flash-image** shut down on Oct 2.

xAI, Perplexity, DeepSeek. No material public list or retirement moves in this window for the walk-up story. Prior moves (grok-4-fast redirect, Sonar → Agent, DeepSeek peak windows) remain the operative facts.

What Operators Should Do With This

1. **Budget the successor, not the retiring id.** When a vendor names a replacement, price the replacement's list rate into failover and default routes before the shutdown date. Oct 23 for the OpenAI nano/mini seats above is the next hard date on the priority board.
2. **Separate frontier stickers from volume seats.** Falling Opus / Sol / Luna prices do not offset a forced nano → luna or flash → flash step on the traffic that actually burns tokens.
3. **Treat intro ladders as timers.** Google's Flash intro through Dec 31, 2026 is a scheduled +100% step on Jan 1, 2027 for 3.6 / 3.7 / 3.8 Flash. Plan the calendar, not only today's page.
4. **Ignore invite-only stickers for capacity planning.** Mythos Glasswing and Daybreak cyber pricing are real on paper and unusable as general PAYG until the vendor opens them.

Prices and dates in the chart and this brief come from vendors' own pricing and deprecation pages (captured Sep 17 to Oct 5, 2026) plus first-party price pins. No aggregator prices.