WHITE PAPER · VERSION 1.0

The 28 Ways AI Models Charge

A Formal Analysis of AI Pricing Mechanisms, From the Per-Token Meter to Reseller Chains

> Every pricing mechanism I can document across 106 models from 18 vendors, ranked from most common to most obscure, with what to watch, how to optimize and what to avoid for each one.

28	106	18
MECHANISMS	TRACKED MODELS	VENDORS

Jerry W. Lambert, Sr.

October 24, 2026 · jerrywlambert.com

Executive Summary

I track 106 priced models from 18 vendors. Across them I count 28 distinct ways a vendor turns usage into a charge. Two are close to universal. Every model prices input and output tokens, and nearly every model discounts cached input. Everything else is a twist on those two meters, and the twists are where budgets break.

The common twists change the price of the same token. A cache write costs 1.25x input on Anthropic, the same as input at most vendors, and nothing at all on many paid models. Batch cuts the price in half. Speed tiers push it the other way. OpenAI Fast is 2x standard, Anthropic fast mode takes Claude Opus 5.5 from $4/$20 to $8/$40, and OpenAI Ultrafast reaches 6x. On some models a prompt-size band reprices the whole request once the prompt crosses a line: above 272,000 tokens on OpenAI, at 200,000 on xAI and Gemini 2.5 Pro, above 100,000 on Claude Haiku 5.5. Data residency adds 10 percent at Anthropic and OpenAI, and on Claude served through regional Amazon Bedrock and Google Cloud endpoints.

The less common twists add new meters. Web search costs $10 per 1,000 calls at Anthropic and OpenAI. Z.ai and MiniMax charge $0.01 per call, which works out to the same price. Google charges $14 per 1,000 after 5,000 free requests a month. Code containers bill by the container-hour or by the 20-minute session. Gemini charges $0.50 per million tokens per hour to keep an explicit cache alive. DeepSeek halves its prices outside two weekday UTC windows and treats every Chinese public holiday as off-peak. Prices also move on a calendar. Gemini 4 Argon doubles on December 1, 2026, and Gemini 3.8, 3.7 and 3.6 Flash double on January 1, 2027. Claude 4.7 and later use a tokenizer that produces about 30 percent more tokens for the same text, which raises the bill without touching the price.

At the far end sit mechanisms that never show up as a token price. Claude Pro and Max ration usage in five-hour windows, and Claude Duet, which runs only on Opus 5.5, draws on those same plan limits. Claude Enterprise charges $20 per user per month plus API rates. OpenRouter takes 5 percent when you bring your own key. Enterprise contracts add take-or-pay commits, and some vendors sell prepaid points that expire.

My operating rule is simple. Every mechanism needs a meter, a dated price and a stacking rule. If a billing system cannot meter a mechanism, it should refuse the request that triggers it instead of billing it at the wrong price. Bill a 1-hour cache write on Moonshot Kimi K3.1 Pro at the 5-minute price and you collect $2.00 of the $4.00 owed on that line. Let a web search through unmetered and you lose $10 for every 1,000 calls.

HOW I COUNTED

I worked from my live Tariff as read on October 24, 2026: 106 priced models from 18 vendors (Alibaba, Amazon, Anthropic, Cohere, DeepSeek, Google, Meta, Microsoft, MiniMax, Mistral, Moonshot, OpenAI, StepFun, Tencent, TypeSafe, Xiaomi, xAI, Z.ai).

I ranked each mechanism by how many of the tracked models can produce it, most common first. For mechanisms my Tariff stores directly (token prices, cache prices, prompt-size bands, peak and off-peak bands, $0 lines, scheduled step-ups, per-request fees), that means the models that carry the price line. For mechanisms vendors publish but my Tariff does not store as a separate line (batch, speed tiers, web search, data residency, storage, runtime), it means the models whose vendor's own pricing page offers it.

Mechanisms outside per-model API pricing (media units, training, subscriptions, per-user licenses, commits and similar) come last, in their original order. The ranking measures how widely a mechanism can produce a charge, not how often the charge appears on invoices. The 106 models are the set I track, not the whole market.

Prices are US dollars per million tokens unless stated. A pair such as $4/$20 means $4 per million input tokens and $20 per million output tokens. Any claim I could not confirm against a vendor's own page or my live Tariff carries the label UNVERIFIED.

The 28 Mechanisms, Ranked by Prevalence

#	MECHANISM	VENDORS (EXAMPLES)
1	Per Input and Output Token	All
2	Cached Input Read Discount	All but TypeSafe
3	Per-Call Server Tool Fees	Anthropic, OpenAI, Google, xAI, Z.ai, MiniMax, Xiaomi, DeepSeek
4	Separate Cache Write Charge	Anthropic, OpenAI, Alibaba, Z.ai, Xiaomi, Amazon, DeepSeek, Mistral, Cohere
5	Batch Discount	Anthropic, OpenAI, Google, Xiaomi (Mistral UNVERIFIED)
6	Free Tiers and $0 Price Lines	Google, Z.ai, Xiaomi, Amazon, TypeSafe
7	Speed, Priority and "Highspeed" Tiers	OpenAI, Anthropic, Google, MiniMax, Moonshot, Xiaomi
8	Container and Session Runtime Fees	Anthropic, OpenAI
9	Data Residency and Regional Price Lists	Anthropic, OpenAI, Bedrock, Google Cloud, Xiaomi, Tencent, Alibaba
10	Prompt-Size (Long Context) Bands	OpenAI, xAI, Google, Anthropic, MiniMax, Alibaba
11	Cache Write Priced by Retention Time (TTL)	Anthropic, OpenAI, Moonshot
12	Flex (Slow, Cheap Synchronous) Tier	OpenAI, Google

#	MECHANISM	VENDORS (EXAMPLES)
13	Hidden Token Overhead	Anthropic, OpenAI
14	Reasoning (Thinking) Tokens Billed as Output	Google, Alibaba, xAI and others
15	Cache and File Storage by Time	Google, OpenAI (Z.ai "limited-time free")
16	Promotional Pricing and Dated Step-Ups	Google, OpenAI, MiniMax
17	Time-of-Day and Day-of-Week Bands	DeepSeek
18	Legacy Model Names Billed as a Different Model	DeepSeek (xAI documents a naming rule)
19	Per Image, Per Second, Per Minute, Per Character	OpenAI, Google, xAI, MiniMax, Z.ai, Xiaomi
20	Fine-Tuning, Training and Fine-Tuned Inference	OpenAI
21	Subscription With Bundled, Rationed Usage	Anthropic, MiniMax, Xiaomi, Cursor, Hark
22	Per-User License and Bring-Your-Own-Key Fees	Anthropic, OpenRouter
23	Minimums, Top-Ups and Expiring Prepaid Balances	Moonshot, Anthropic, OpenAI, MiniMax, DeepSeek
24	Commits, Take-or-Pay and Negotiated Volume Discounts	Anthropic, Google, enterprise contracts
25	Marketplace Consumption Units	Anthropic on AWS and Azure
26	Data-for-Discount Tiers	Google, OpenAI
27	Billing-Event Rules	DeepSeek, Anthropic, MiniMax, Meta
28	Resale Markup and Reseller Chains	Resellers and aggregators

The Mechanisms

01 Per Input and Output Token

NAME · MECHANISM 1 OF 28

DESCRIPTION

This is the base meter. The vendor prices input (prompt) tokens and output (completion) tokens separately, per million. Output almost always costs more.

EXAMPLE

Claude Opus 5.5 lists $4.00 input and $20.00 output. Across the 106 models I track, input runs from $0.03 (Z.ai GLM-OCR) to $30.00 (OpenAI gpt-5.5-pro), and output tops out at $180.00, also on gpt-5.5-pro. The output-to-input ratio clusters at a few values. 5x is the most common, then 4x, 6x and 3x.

MODELS THAT USE IT

Every vendor I track.

THINGS TO WATCH

- Output-heavy work such as agents, code generation and long reports is priced on the expensive side. At a 5x ratio, one output token costs as much as five input tokens.
- Token counts do not compare cleanly across vendors, or even across generations of one vendor's models (see Mechanism 13).
- Some vendors price output at $0. TypeSafe Jev bills $0.042 input and $0 output.

WAYS TO OPTIMIZE

- Cap output length with a max-tokens setting and ask for structured, terse output.
- Route by task. Claude Haiku 5.5 at $0.10/$0.50 against Opus 5.5 at $4/$20 is a 40x spread inside one vendor.

THINGS TO AVOID

- Comparing models on input price alone.
- Treating a list price as the effective price. Nearly every mechanism that follows changes it.

NAME · MECHANISM 2 OF 28

02 Cached Input Read Discount

DESCRIPTION

When a request reuses a prompt prefix the vendor already holds in cache, those tokens bill at a cache-hit price below the normal input price.

EXAMPLE

VENDOR AND MODEL	CACHE READ PRICE
Anthropic, most models	0.1x input
Anthropic Opus 5.5 and Sonnet 5.5	0.05x input
Anthropic Fable 5.1	0.025x input ($0.25 against $10.00)
DeepSeek Flash, off-peak	$0.003 cached against $0.15 uncached (98 percent off)

Across tracked models the most common ratio is 0.1x, then 0.2x and 0.25x.

MODELS THAT USE IT

Every vendor except TypeSafe. A few models price a cache read at the full input price (several Alibaba models, plus OpenAI gpt-5.5-pro and o3-pro). Some paid models publish no cache price at all (TypeSafe models and some Z.ai models). Google Gemma 4 12B and 31B carry all four price lines (input, output, cache read and cache write) at $0.

THINGS TO WATCH

- Discounts vary 30-fold, from 0.0083x (Xiaomi MiMo-V2.6-Pro) to 0.25x. Fewer than half of the discounting models charge exactly 0.1x, so a blanket "90 percent off cached input" assumption is wrong for most of them, and more so once you add the at-par, unpriced and free models.
- On Anthropic, a cache hit also refreshes the cache timer.
- Some models cache only implicitly. Google says Gemini 4 Argon supports "implicit caching only (explicit caching is not supported)."

WAYS TO OPTIMIZE

- Put stable content first (system prompt, tool definitions, reference documents) and variable content last, so the prefix matches.
- Keep a conversation on one model. Caches do not carry across models.

THINGS TO AVOID

- Reordering tools or injecting a timestamp at the top of the prompt. Either one breaks the prefix.
- Billing a model with no published cache price as if it had one. A cache price with no first-party basis is an invented number.

NAME · MECHANISM 3 OF 28

03 Per-Call Server Tool Fees

DESCRIPTION

When the vendor runs a tool on its own servers (web search, grounding, file search), it charges per call or per 1,000 calls. That fee sits on top of the tokens the tool's results add to the prompt.

EXAMPLE

VENDOR	TOOL	PRICE	NOTES
Anthropic	Web search	$10 per 1,000 searches plus tokens	A failed search is not billed. Web fetch has no extra fee.
OpenAI	Web search	$10 per 1,000 calls plus content tokens	The non-reasoning preview tool is $25 per 1,000 with free content tokens.
OpenAI	File search	$2.50 per 1,000 calls	
Google	Grounding with Search or Maps (Gemini 3 and newer)	$14 per 1,000 after 5,000 free requests per month	The allowance is shared across models.
Google	Grounding on Gemini 2.5	$35 per 1,000 grounded prompts (Search), $25 per 1,000 (Maps)	After the free daily allowance.
Z.ai	Web search	$0.01 per use	
MiniMax	web_search	$0.01 per request	
Xiaomi	Overseas search	$5 per 1,000	
DeepSeek	Search and Fetch	$0.75 per 1,000 successful requests	Free until 00:00 Beijing time on November 1, 2026 (16:00 UTC, 11:00 AM Central, on October 31).
xAI	Web, X search, code and collections tools	$2.50 to $10 per 1,000	UNVERIFIED. From my October 20 review; I did not fetch the xAI tool page again for this article.

On Google, one customer request can trigger several billed searches.

MODELS THAT USE IT

Z.ai, OpenAI, Anthropic, Google, xAI, Xiaomi, DeepSeek and MiniMax. Moonshot also prices its built-in tools, but I have not checked the rate.

THINGS TO WATCH

- The fee is per call, not per answer. One prompt can fan out into many searches.
- The tool's results come back as input tokens and stay in context on later turns.
- Free allowances are shared across models and reset monthly or daily, depending on the model family.

WAYS TO OPTIMIZE

- Cap tool calls per request.
- Prefer token-only tools where the vendor offers them, such as Anthropic web fetch and Google URL context.

THINGS TO AVOID

- Letting a server tool through a billing layer that cannot count it. If the call count is not metered, refuse the tool instead of billing tokens only.

NAME · MECHANISM 4 OF 28

04 Separate Cache Write Charge

DESCRIPTION

The vendor bills the cache write as its own line. Depending on the vendor, the write costs more than input, the same as input, less than input, or nothing.

EXAMPLE

TREATMENT	EXAMPLE
Premium	Anthropic 5-minute write at 1.25x input: Opus 5.5 $5.00 against $4.00
Premium	OpenAI newer models, 30-minute write at 1.25x: gpt-6-astra $12.50 against $10.00
Discount	DeepSeek V4 Pro write $0.435 against $1.32 peak input. UNVERIFIED. DeepSeek's current page publishes no write price; my Tariff carries the figure from an earlier source.
Free	Z.ai lists cached input storage as "Limited-time Free" on all paid GLM models
Free	Xiaomi lists "Cache Write: Limited-time Free"

MODELS THAT USE IT

Anthropic, OpenAI and Alibaba charge a premium for the write. Some vendors price it below input, as in the DeepSeek example. Z.ai, Xiaomi and Amazon Nova charge nothing for it. Most other models price the write at par with input, and TypeSafe jev-1.13.0 and jev-1.14.0 have no cache-write line at all.

THINGS TO WATCH

- A write premium means caching loses money if nothing reads the prefix again. Anthropic says a 5-minute write pays off after one read and a 1-hour write after two.
- "Limited-time free" has no end date on either the Z.ai page or the Xiaomi page.

WAYS TO OPTIMIZE

- Cache only the prefixes you expect to reuse within the retention window.

THINGS TO AVOID

- Treating a free write as permanent. Record it as a promotion and check it again at every price review.

05 Batch Discount

NAME · MECHANISM 5 OF 28

DESCRIPTION

You submit work asynchronously and accept a delay, typically up to 24 hours. In exchange you pay roughly half.

EXAMPLE

MODEL	STANDARD	BATCH
Anthropic Opus 5.5	$4/$20	$2/$10 (50 percent off input and output)
OpenAI gpt-6-astra	$10/$50	$5/$25
Google Gemini 4 Argon	$2/$10	$1/$5
Xiaomi MiMo-V2.6-Pro	$0.435/$0.87	$0.2175/$0.435

MODELS THAT USE IT

OpenAI, Anthropic, Google and Xiaomi. Mistral's pricing page has a Batch tab, but I did not extract its rates, so Mistral batch pricing is UNVERIFIED. xAI says "not every model accepts Batch API requests," and its batch rates are UNVERIFIED.

THINGS TO WATCH

- On Anthropic, the batch discount stacks with caching.
- Not every model or mode qualifies. Xiaomi Ultraspeed, Anthropic fast mode and Anthropic Managed Agents have no batch option.
- OpenAI's batch table omits cached-input discounts for several older models, gpt-4.1 among them.

WAYS TO OPTIMIZE

- Move evaluations, backfills, classification and nightly reports to batch.

THINGS TO AVOID

- Using batch for anything a user is waiting on.

NAME · MECHANISM 6 OF 28

06 Free Tiers and $0 Price Lines

DESCRIPTION

The vendor prices a model, a unit or a usage band at zero, either permanently or for a limited time.

EXAMPLE

- **Fully free models.** Google Gemma 4 is free tier only and listed as "Not available" on paid. I track two sizes, 12B and 31B. Z.ai also lists GLM-4.7-Flash, GLM-4.5-Flash and GLM-4.6V-Flash as free, but I do not track those three.
- **Free units.** TypeSafe Jev output is $0. Amazon Nova cache writes are $0. Z.ai and Xiaomi cache writes are free for a limited time.
- **Free allowances.** Anthropic code execution includes 1,550 free hours per organization per month. Google grounding includes 5,000 requests per month. OpenAI includes 1 GB of file storage, and OpenAI moderation is free.
- **Free trial credit.** Anthropic gives new users "a small amount of free credits."

MODELS THAT USE IT

Z.ai, TypeSafe, Amazon, Xiaomi and Google carry $0 price lines. Gemma 4 12B and Gemma 4 31B are entirely free. Z.ai's free Flash models are on the vendor's list but not among the models I track.

THINGS TO WATCH

- A free tier can carry a data cost. Google's free tier states "Content used to improve our products."
- A $0 line with no first-party source looks exactly like a missing price.

WAYS TO OPTIMIZE

- Use free allowances for development and evaluation traffic.

THINGS TO AVOID

- Recording an unknown price as $0. A $0 line should carry its source and the vendor's own word "free."

NAME · MECHANISM 7 OF 28

07 Speed, Priority and "Highspeed" Tiers

DESCRIPTION

You pay more for faster output or for priority admission. Vendors sell this either as a request flag on the same model or as a separate model ID.

EXAMPLE

MODEL	TIER	PRICE	STANDARD
OpenAI gpt-6-astra	Fast (renamed from Priority on July 30, 2026), 2x	$20/$100	$10/$50
OpenAI gpt-6-astra	Ultrafast, 6x	$60/$300	$10/$50
Anthropic Opus 5.5	Fast mode, 2x	$8/$40	$4/$20
Google Gemini 4 Argon	Priority, 1.8x	$3.60/$18.00	$2.00/$10.00
MiniMax M3	Priority, 1.5x	$0.45/$1.80	$0.30/$1.20
Moonshot kimi-k3.1-pro-highspeed	Separate model ID	$4.00 input, $0.20 cached, $4.00 5-minute write, $20.00 output	$2.00, $0.20, $2.00, $10.00
Xiaomi mimo-v2.6-pro-ultraspeed	Separate model ID, 10x	$4.35/$8.70	$0.435/$0.87
xAI	Priority, 2x	UNVERIFIED (from my October 20 review)	Not verified

Anthropic fast mode on Opus 5 and Opus 4.8 is $10/$50. It stacks with caching and data residency, and it is not available with batch.

MODELS THAT USE IT

OpenAI, Google, Anthropic and MiniMax sell it as a tier. Moonshot and Xiaomi sell it as separate highspeed model IDs.

THINGS TO WATCH

- The multiplier is not uniform. I see 1.5x, 1.8x, 2x, 6x and 10x in published prices.
- Moonshot Highspeed doubles input and output but keeps the cached read at $0.20.
- A tier flag on the request can silently raise the price if a client library sets it.

WAYS TO OPTIMIZE

- Reserve speed tiers for user-facing paths where latency matters, and keep them off background work.

THINGS TO AVOID

- Passing a tier flag through to a vendor when the platform bills at standard rates. Either price the multiplier or reject the flag.

NAME · MECHANISM 8 OF 28

08 Container and Session Runtime Fees

DESCRIPTION

The vendor charges by time for a sandbox or an agent session, separately from tokens.

EXAMPLE

PRODUCT	PRICE	RULES
Anthropic code execution	$0.05 per container-hour after 1,550 free hours per month	5-minute minimum. Free when used with the newer web search or fetch tools. Files in the request start billing even if the tool is never called.
Anthropic Managed Agents	$0.08 per session-hour	Counted only while the session is "running."
OpenAI Hosted Shell and Code Interpreter	$0.03 (1 GB) to $1.92 (64 GB) per 20-minute session	Billed by the minute, with a 5-minute minimum.
Google code execution	No runtime charge	Tokens only.

MODELS THAT USE IT

OpenAI and Anthropic.

THINGS TO WATCH

- Memory size multiplies OpenAI's price 64-fold.
- Minimum billable durations make many short calls expensive.

WAYS TO OPTIMIZE

- Reuse a container across steps.
- Pick the smallest memory class that does the job.

THINGS TO AVOID

- Attaching files "just in case" on Anthropic. The attachment starts the meter.

NAME · MECHANISM 9 OF 28

09 Data Residency and Regional Price Lists

DESCRIPTION

The vendor adds a surcharge to keep processing inside a named geography, or publishes a separate price list for each region.

EXAMPLE

- Anthropic's inference_geo: "us" setting is 1.1x on every token category for Claude 4.6 and later.
- Claude on Amazon Bedrock regional endpoints and Google Cloud multi-region endpoints carries a 10 percent premium over global endpoints.
- OpenAI regional processing and FedRAMP endpoints carry a 10 percent uplift for models released on or after March 5, 2026.
- Xiaomi publishes separate China (RMB) and overseas (USD) tables. MiMo-V2.6-Pro input is ¥3.00 on one and $0.435 on the other. The Tencent Hunyuan prices in my Tariff are Singapore-region prices. Alibaba Model Studio publishes an International (Singapore) list.

MODELS THAT USE IT

Claude models with US-only inference, plus models priced from a region-specific list at Alibaba, Xiaomi and Tencent. Which OpenAI models carry the uplift depends on each model's release date, and that list is UNVERIFIED.

THINGS TO WATCH

- The surcharge multiplies every other mechanism, including fast mode and long-context bands.

WAYS TO OPTIMIZE

- Use residency only for data that requires it, and route the rest globally.

THINGS TO AVOID

- Accepting a residency parameter on a route that bills global rates.

10 Prompt-Size (Long Context) Bands

NAME · MECHANISM 10 OF 28

DESCRIPTION

When the prompt crosses a threshold, the whole request reprices at a higher band. The higher price applies to every token in the request, not only the tokens above the line.

EXAMPLE

MODEL	THRESHOLD	STANDARD	ABOVE THRESHOLD
OpenAI gpt-5.5	Above 272,000 input tokens	$5/$30	$10/$45
xAI grok-4.7	200,000 prompt tokens or more	$2.00/$6.00	$4.00/$12.00, "for all tokens in the request"
Google Gemini 2.5 Pro	Above 200,000	$1.25/$10	$2.50/$15
Anthropic Claude Haiku 5.5	Above 100,000 (cache reads count toward it)	$0.10/$0.50	$0.50/$2.50 (5x)
MiniMax M3	Above 512,000	$0.30/$1.20	$0.60/$2.40
Alibaba Qwen 3.7 Plus	Above 256,000	$0.40/$1.60	$1.20/$4.80

The counter-example is Anthropic. Claude 4.6 and later, except Haiku 5.5, bill the full 1M-token window at standard rates.

MODELS THAT USE IT

OpenAI, xAI, Google, Anthropic (Haiku 5.5 only), MiniMax and Alibaba.

THINGS TO WATCH

- One token over the line can double the cost of the entire request.
- Thresholds differ by vendor, and so does the boundary. OpenAI's line is ">272K" and xAI's is "≥200K".

WAYS TO OPTIMIZE

- Trim or summarize context to stay under the threshold.
- Retrieve the passages a request needs instead of stuffing whole documents into the prompt.

THINGS TO AVOID

- Estimating cost from a blended average. Price each request against its own band.

NAME · MECHANISM 11 OF 28

11 Cache Write Priced by Retention Time (TTL)

DESCRIPTION

The longer the vendor holds the cache, the more the write costs.

EXAMPLE

MODEL	SHORT RETENTION WRITE	LONG RETENTION WRITE
Anthropic Opus 5.5	5-minute: $5 (1.25x input)	1-hour: $8 (2x input)
OpenAI gpt-6-astra	30-minute: $12.50	12-hour: $25.00 (2x the 30-minute)
Moonshot kimi-k3.1-pro	5-minute: $2.00	1-hour: $4.00
Moonshot kimi-k3.1-pro Highspeed	5-minute: $4.00	1-hour: $8.00
Moonshot kimi-k3	5-minute: $3.00	1-hour: $6.00

Moonshot selects the retention time with a request option, prompt_cache_options.ttl.

MODELS THAT USE IT

Anthropic, OpenAI and Moonshot.

THINGS TO WATCH

- A platform that stores only one write column bills the long-retention write at the short-retention price. On Kimi K3.1 Pro that is $2.00 billed against $4.00 owed.
- Each vendor uses a different request field to select retention, so a check written for one vendor's field misses another's.

WAYS TO OPTIMIZE

- Use the long retention only when the gaps between reuses exceed the short window. On Anthropic it needs two reads to break even.

THINGS TO AVOID

- Metering one retention time and passing the other through unpriced. Either price every retention column or refuse the long-retention option.

12 Flex (Slow, Cheap Synchronous) Tier

NAME · MECHANISM 12 OF 28

DESCRIPTION

The request stays synchronous but accepts slower or interruptible service in exchange for a batch-like discount.

EXAMPLE

MODEL	STANDARD	FLEX
OpenAI gpt-6-astra	$10/$50	$5/$25
Google Gemini 4 Argon	$2/$10	$1/$5

MODELS THAT USE IT

OpenAI and Google. Whether Gemini 3.7 Flash offers batch, Flex and Priority is UNVERIFIED. Google's page does not show it as its own section, so I leave it out.

THINGS TO WATCH

- Flex trades availability for price. Expect slower responses or capacity errors.

WAYS TO OPTIMIZE

- Use Flex for non-urgent synchronous jobs that cannot be restructured as batch.

THINGS TO AVOID

- Retrying a failed Flex request on Fast without a cost check.

NAME · MECHANISM 13 OF 28

13 Hidden Token Overhead

DESCRIPTION

These are tokens you pay for but did not write: tool system prompts, tool definitions, tokenizer changes and fixed blocks.

EXAMPLE

- Anthropic adds a tool-use system prompt of 286 tokens (Opus 5.5) to 675 tokens (Opus 4.7) whenever tools are present. The computer-use toolset adds about 4,500 input tokens and the browser toolset about 6,600.
- Claude 4.7 and later use a tokenizer that "produces approximately 30% more tokens for the same text."
- OpenAI bills web search content for gpt-4o-mini and gpt-4.1-mini as a fixed block of 8,000 input tokens per call.

MODELS THAT USE IT

Every tracked Claude model carries tool overhead, and the newer Claude models also use the newer tokenizer. OpenAI gpt-4.1-mini falls under the fixed 8,000-token search block. The other model named in that OpenAI rule, gpt-4o-mini, is not one I track.

THINGS TO WATCH

- A model with a lower list price can cost more per task if it emits 30 percent more tokens.

WAYS TO OPTIMIZE

- Compare models on cost per completed task, not cost per token.
- Declare only the tools a request needs.

THINGS TO AVOID

- Carrying old per-token benchmarks across a tokenizer change.

14 Reasoning (Thinking) Tokens Billed as Output

NAME · MECHANISM 14 OF 28

DESCRIPTION

Internal reasoning tokens bill at the output price, even when the vendor hides or encrypts them.

EXAMPLE

- Google labels every Gemini output price "including thinking tokens." Gemini 4 Argon output is $10.00.
- Alibaba Qwen 3.6-35B-A3B lists one $2.25 output price for "thinking and non-thinking."
- xAI's Responses API always returns encrypted reasoning content for grok-4.7.

MODELS THAT USE IT

Paid Gemini models, whose prices carry the "including thinking tokens" label, and Alibaba Qwen 3.6-35B-A3B. In practice it covers most models that can reason. No tracked vendor publishes a separate thinking-token price.

THINGS TO WATCH

- Reasoning effort settings change output volume, often by multiples, at the same unit price.
- DeepSeek defaults to thinking mode.

WAYS TO OPTIMIZE

- Set the lowest reasoning effort that passes your quality bar.
- Use non-thinking modes for simple tasks.

THINGS TO AVOID

- Forecasting output cost from the length of the visible answer.

15 Cache and File Storage by Time

NAME · MECHANISM 15 OF 28

DESCRIPTION

The vendor charges to keep data available, priced per token-hour or per GB-day.

EXAMPLE

ITEM	PRICE
Gemini context cache storage	$0.50 per million tokens per hour through December 31, 2026, then $1.00 from January 1, 2027
OpenAI file search storage	$0.10 per GB per day after 1 GB free
OpenAI ChatKit upload storage	$0.10 per GB per day after 1 GB free
Z.ai cached input storage	"Limited-time Free"

MODELS THAT USE IT

Gemini models with explicit caching. OpenAI file storage applies to every OpenAI model that uses file search.

THINGS TO WATCH

- A 1M-token Gemini cache held for 24 hours costs $12 in storage today and $24 from January 1, 2027, before a single read.

WAYS TO OPTIMIZE

- Set short cache lifetimes.
- Delete stores you no longer query.

THINGS TO AVOID

- Long-lived explicit caches on low-traffic prompts.

NAME · MECHANISM 16 OF 28

16 Promotional Pricing and Dated Step-Ups

DESCRIPTION

The vendor publishes a price with an announced end date, or schedules an increase in advance.

EXAMPLE

MODEL	CURRENT PRICE	WHAT CHANGES
Google Gemini 4 Argon	$2/$10 through November 30, 2026	$4/$20 from December 1, 2026
Google Gemini 3.8, 3.7 and 3.6 Flash	$0.75/$3.75 through December 31, 2026	$1.50/$7.50 from January 1, 2027
OpenAI GPT-5.6 Sol	Promotional pricing "available at least through November 21, 2026"	No successor price stated
MiniMax M3	"Permanent 50% off" list	No end date
Anthropic Sonnet 5 (context, not counted) [3]	Introductory $2/$10 became the standard price	The planned September 1 increase to $3/$15 was cancelled

[3] Context, not counted. Anthropic announced Sonnet 5 at an introductory $2/$10 through August 31, 2026, then made it the standard price and cancelled the planned September 1 increase to $3/$15. No dated change remains, so Sonnet 5 is not counted here.

MODELS THAT USE IT

Gemini 4 Argon and Gemini 3.8, 3.7 and 3.6 Flash, which have dated step-ups, plus GPT-5.6 Sol and MiniMax M3. The "Permanent 50% off" label is on M3 only, not M3.1-Flash.

THINGS TO WATCH

- Promotions without end dates, such as the Z.ai and Xiaomi "limited-time free" lines, can end without notice.

WAYS TO OPTIMIZE

- Load the successor price with its start date ahead of time, so billing changes on the right day.

THINGS TO AVOID

- Carrying promotional prices into budgets or contracts past the stated end date.

17

NAME · MECHANISM 17 OF 28

Time-of-Day and Day-of-Week Bands

DESCRIPTION

The same model costs different amounts depending on when the request runs.

EXAMPLE

DeepSeek's off-peak rates are half of peak. Peak runs 01:00 to 04:00 and 06:00 to 10:00 UTC, Monday through Friday, excluding Chinese public holidays. Weekends and Chinese public holidays are off-peak all day.

MODEL	PEAK	OFF-PEAK
deepseek-flash	$0.30/$1.20	$0.15/$0.60
deepseek-v4-pro	$1.32/$3.96	$0.66/$1.98

MODELS THAT USE IT

DeepSeek only.

THINGS TO WATCH

- Peak falls in the US evening and overnight: 8:00 PM to 11:00 PM and 1:00 AM to 5:00 AM Central time, Sunday night through Friday morning, under daylight time.
- The holiday calendar is Chinese, not American. A meter that ignores it overbills on those days.

WAYS TO OPTIMIZE

- Schedule DeepSeek batch-style jobs for US business hours and weekends.

THINGS TO AVOID

- Pricing all DeepSeek traffic at one band.

NAME · MECHANISM 18 OF 28

18 Legacy Model Names Billed as a Different Model

DESCRIPTION

An old model name keeps working, but a different model serves the request and sets the bill.

EXAMPLE

- DeepSeek: "The legacy names deepseek-v4-flash and deepseek-v4-flash-vision-exp are still accepted, but the corresponding models have been retired, their requests are served by the DeepSeek-V4.1-Flash-1017 model and billed at the Flash price."
- xAI points unversioned names and "-latest" names at newer releases.

MODELS THAT USE IT

The legacy name deepseek-v4-flash, which DeepSeek serves with DeepSeek-V4.1-Flash-1017 and bills at the same price as the current deepseek-flash ($0.30/$1.20 peak, $0.15/$0.60 off-peak). xAI documents its naming rule, but I have not checked which tracked xAI names point to newer releases, so I leave them out.

THINGS TO WATCH

- The price can stay the same while the model changes, or the reverse.

WAYS TO OPTIMIZE

- Pin dated model versions in production.

THINGS TO AVOID

- Billing one model's price for another model's output without recording the substitution.

NAME · MECHANISM 19 OF 28

19 Per Image, Per Second, Per Minute, Per Character

DESCRIPTION

The vendor prices media outside the token meter, or converts each asset into a fixed number of tokens.

EXAMPLE

UNIT	PUBLISHED EXAMPLES
Per image	xAI grok-imagine-image $0.02 and image-2.0 $0.04 per image. Z.ai GLM-Image $0.015. MiniMax image-01 $0.0035. Gemini 2.5 Flash Image $0.039 per 1024×1024 image (1,290 tokens).
Per second or per video	xAI video $0.020 to $0.080 per second. MiniMax H3 $0.025 to $0.13 per output second, by resolution. Z.ai CogVideoX-3 $0.20 per video.
Per minute or per hour	OpenAI GPT-Live $0.05 per minute, billed per second. OpenAI gpt-transcribe $0.0045 per minute. xAI speech-to-text $0.10 per hour (REST). Xiaomi ASR $0.074 per hour. Gemini bills audio at 25 tokens per second.
Per character	OpenAI tts-1 $15 per million characters. xAI text-to-speech $15 per million characters. MiniMax speech $60 to $100 per million characters.
Modality uplift inside a token model	OpenAI gpt-realtime-2.1 audio input $32 against text input $4.

MODELS THAT USE IT

OpenAI, Google, xAI, MiniMax, Z.ai and Xiaomi publish media units. My Tariff is text-only, so none of the models I track carry them.

THINGS TO WATCH

- Audio sent to a text-priced model can bill at the vendor's higher audio rate. Gemini Flash audio input runs about 2x text. The rate pair is UNVERIFIED for this article.

WAYS TO OPTIMIZE

- Pick resolution and duration deliberately. Output resolution drives video price.

THINGS TO AVOID

- Routing audio or image content through an endpoint priced only for text.

NAME · MECHANISM 20 OF 28

20 Fine-Tuning, Training and Fine-Tuned Inference

DESCRIPTION

The vendor charges for training, then charges a higher inference price for the customized model.

EXAMPLE

ITEM	PRICE
OpenAI gpt-4.1 fine-tuning	$25 per million training tokens
OpenAI gpt-4.1 fine-tuned inference	$3.00/$12.00 against $2.00/$8.00 for the base model
OpenAI o4-mini reinforcement fine-tuning	$100 per training hour

OpenAI is winding the platform down and has closed it to new users.

MODELS THAT USE IT

OpenAI, as a legacy offering. Fine-tuned model hosting fees at other vendors are UNVERIFIED.

THINGS TO WATCH

- In reinforcement fine-tuning, grader tokens bill at the grader model's rate.
- Fine-tuned models retire with their base model.

WAYS TO OPTIMIZE

- Try prompting and caching before training.

THINGS TO AVOID

- Building on a fine-tuning platform with an announced wind-down.

21 Subscription With Bundled, Rationed Usage

NAME · MECHANISM 21 OF 28

DESCRIPTION

The customer pays a flat monthly fee that includes usage measured in opaque limits instead of tokens.

EXAMPLE

- Claude Pro is $20 per month ($17 on annual billing). Max starts at $100 with 5x or 20x Pro usage per five-hour session, plus weekly limits, and includes $100 to $200 per month of API credits. Overage on paid plans can continue "at standard API rates."
- Claude Duet is a research preview, by waitlist, for US Pro and Max subscribers only. It runs on Claude Opus 5.5 with no model choice. It is included at no extra cost and draws on the plan's five-hour and weekly limits, and its background checks count too.
- MiniMax Audio Subscription runs $5 to $999 per month for 100,000 to 20,000,000 audio points. MiniMax Video Packages run $1,000 to $6,000 per month, and unused points expire.
- Xiaomi says Token Plan quota "is not interoperable" with a pay-as-you-go balance.
- Cursor Cloud Agents now draw on "your plan's usage … at the selected model's rate." Hark Pro is $20 (2x usage) and $100 (10x). Both are UNVERIFIED, from my October 22 market brief.

MODELS THAT USE IT

Anthropic, MiniMax, Xiaomi, Cursor and Hark. These are plans, not per-model API prices.

THINGS TO WATCH

- Limits are expressed as multiples of another plan, not in tokens.
- Background agents spend the same pool as chat.

WAYS TO OPTIMIZE

- Measure the cost per unit of work under the plan against API rates before you commit.

THINGS TO AVOID

- Assuming "included" means unlimited.

NAME · MECHANISM 22 OF 28

22 Per-User License and Bring-Your-Own-Key Fees

DESCRIPTION

The customer pays a fixed fee per user, or pays a platform fee when it supplies its own vendor API key (bring your own key, or BYOK).

EXAMPLE

- Claude Enterprise charges $20 per user per month, plus usage billed at API rates.
- OpenRouter BYOK charges 5 percent of what the same model would cost on OpenRouter, after a pay-as-you-go allowance of $25,000 of list-price inference per month.

MODELS THAT USE IT

Anthropic, OpenRouter, and resellers generally.

THINGS TO WATCH

- With BYOK, the vendor bills the customer directly and the platform's fee sits on top.
- A request that falls back from the customer's key to the platform's shared capacity is billed to the platform account.

WAYS TO OPTIMIZE

- Match the pricing model to volume. Per-user pricing favors heavy users, and a percentage fee favors light ones.

THINGS TO AVOID

- Mixing resale and BYOK billing in the same billing cycle. A BYOK request that lands on the platform's own key should be blocked, not silently billed either way. This is a design principle, not a vendor price.

NAME · MECHANISM 23 OF 28

23 Minimums, Top-Ups and Expiring Prepaid Balances

DESCRIPTION

Vendors set spend floors, require unlock payments and let prepaid balances lapse.

EXAMPLE

- Moonshot Kimi K3.1 Pro is "unlocked after a successful top-up (minimum $1)."
- Anthropic code execution and OpenAI containers bill a 5-minute minimum.
- MiniMax video points "expire with the package and do not roll over."
- DeepSeek deducts from a granted balance before a topped-up balance.

MODELS THAT USE IT

Moonshot, Anthropic, OpenAI, MiniMax and DeepSeek.

THINGS TO WATCH

- Prepaid credit is often non-refundable.
- Expiry turns unused prepay into a price increase.

WAYS TO OPTIMIZE

- Buy prepay in amounts you will consume within the validity window.

THINGS TO AVOID

- Holding large prepaid balances with vendors that apply expiry rules.

NAME · MECHANISM 24 OF 28

24 Commits, Take-or-Pay and Negotiated Volume Discounts

DESCRIPTION

The customer promises to spend a set amount in return for a lower price, and the vendor still bills any unused commitment.

EXAMPLE

- Anthropic's sales-assisted Enterprise plan lists "Tiered incentives on committed spend" and negotiated volume discounts.
- Google Enterprise offers provisioned throughput and volume-based discounts.
- Illustrative example, not a vendor price: a $1,000 monthly take-or-pay commit with $700 used produces a $300 shortfall line, with no rollover.

MODELS THAT USE IT

Enterprise contracts at major vendors, including Anthropic and Google.

THINGS TO WATCH

- Whether the vendor discount reaches the end customer or stays with the reseller as margin.
- When the shortfall is billed, monthly or at the end of the term.

WAYS TO OPTIMIZE

- Commit at the floor of forecast usage, not the midpoint.

THINGS TO AVOID

- Commits tied to a single model generation. Prices fall and models retire within a term.

NAME · MECHANISM 25 OF 28

25 Marketplace Consumption Units

DESCRIPTION

The vendor prices usage in dollars, then converts it into the marketplace's own billing unit.

EXAMPLE

Claude Platform on AWS and Claude in Microsoft Foundry bill in Claude Consumption Units at $0.01 per unit (100 units equal $1.00), metered hourly. Discounts show up as fewer units.

MODELS THAT USE IT

Anthropic on AWS and Azure.

THINGS TO WATCH

- The cloud bill shows a single unit line, so per-model detail lives in the vendor console.

WAYS TO OPTIMIZE

- Use marketplace billing to draw down existing cloud commitments.

THINGS TO AVOID

- Reconciling marketplace units without the vendor's per-model breakdown.

NAME · MECHANISM 26 OF 28

26 Data-for-Discount Tiers

DESCRIPTION

The vendor lowers the price in exchange for the right to use your data.

EXAMPLE

- Google's invite-only Gemini API Contributor Tier prices Gemini 4 Argon at $1/$5, half of standard, with "Used to improve our products: Yes."
- OpenAI o4-mini fine-tuned inference with data sharing is $2/$8 against $4/$16.

MODELS THAT USE IT

Google and OpenAI.

THINGS TO WATCH

- You pay for the discount with data rights. Check contracts and regulated data before you opt in.

WAYS TO OPTIMIZE

- Use these tiers only for public or synthetic workloads.

THINGS TO AVOID

- Sending customer or regulated data through a contributor tier.

27 Billing-Event Rules

NAME · MECHANISM 27 OF 28

DESCRIPTION

These rules decide which events are billable, separate from the price itself.

EXAMPLE

- DeepSeek bills Search and Fetch per successful request.
- Anthropic does not bill a web search that errors.
- MiniMax does not deduct points for failed or security-blocked videos, bills the original inputs again when a video is regenerated, and charges $1.50 per cloned voice and $3 per designed voice at first use.
- Meta bills SAM 3.1 video at a minimum of 50 frames. UNVERIFIED, from my October 22 market brief.

MODELS THAT USE IT

DeepSeek, Anthropic, MiniMax and Meta.

THINGS TO WATCH

- The same price can produce different bills depending on success rules and minimums.

WAYS TO OPTIMIZE

- Count only billable events in your own meter, so invoices reconcile.

THINGS TO AVOID

- Billing your customers for failed calls the vendor did not charge you for.

NAME · MECHANISM 28 OF 28

28 Resale Markup and Reseller Chains

DESCRIPTION

A reseller prices on top of vendor cost, and resellers can sit under other resellers.

EXAMPLE

A worked illustration, not a vendor price.

LEVEL	COST TO THIS LEVEL	SELLS AT	MARGIN
Vendor list		$8.00	
Platform, 25 percent markup	$8.00	$10.00	$2.00
Reseller, 20 percent over its own cost	$10.00	$12.00	$2.00

Each level earns $2.00. Each level's markup applies to the price that level pays, never to vendor list.

MODELS THAT USE IT

Aggregators and resellers.

THINGS TO WATCH

- A model a reseller has not priced should be refused at that level, not fall through to the parent's price.

WAYS TO OPTIMIZE

- Keep a dated price trail from vendor list to end customer for every call.

THINGS TO AVOID

- Applying a markup to vendor list when the reseller's actual cost is the parent's price.

How the Mechanisms Stack

Vendors publish how their multipliers combine. On Anthropic, cache multipliers stack on fast mode, data residency stacks on fast mode, batch and caching combine, and fast mode is not available with batch.

Anthropic's own worked example: a one-hour Managed Agents session on Opus 5.5 with 50,000 input tokens and 15,000 output tokens costs $0.58 ($0.20 input, $0.30 output, $0.08 runtime). If 40,000 of those input tokens are cache reads, the session costs $0.428.

The order I use for a billing system: select the prompt-size band, apply multipliers (speed, residency), apply resale rules, then volume tiers and discounts, then credits, then draw down any commit.

Methodology

Built from vendors' public announcements and pricing pages across 106 tracked models, checked October 2026.

Appendix B. Claims Labeled UNVERIFIED

1. UNVERIFIED. xAI server tool fees of $2.50 to $10 per 1,000 calls (Mechanism 3).
2. UNVERIFIED. xAI priority tier at 2x (Mechanism 7).
3. UNVERIFIED. DeepSeek V4 Pro cache write of $0.435. DeepSeek's current page publishes no write price (Mechanism 4).
4. UNVERIFIED. Mistral batch and priority rates, and xAI batch rates (Mechanism 5).
5. UNVERIFIED. Which OpenAI models carry the 10 percent residency uplift for models released on or after March 5, 2026 (Mechanism 9).
6. UNVERIFIED. The Gemini Flash audio-input rate pair, about 2x text (Mechanism 19).
7. UNVERIFIED. Fine-tuned model hosting fees at vendors other than OpenAI (Mechanism 20).
8. UNVERIFIED. Cursor "plan usage at the selected model's rate" and the Hark $20 and $100 plan multiples (Mechanism 21).
9. UNVERIFIED. Meta SAM 3.1 minimum of 50 frames (Mechanism 27).
10. UNVERIFIED. Gemini 3.7 Flash batch, Flex and Priority availability. It is not shown as its own section on Google's page, and I leave it out of the Gemini figures (Mechanism 12).

JERRYWLAMBERT.COM

Get the Next Analysis First

I publish what AI models cost, how vendors bill and what changed overnight. Subscribe to get it in your inbox.

Subscribe

https://jerrywlambert.com/newsletter/

Jerry W. Lambert, Sr.
jerrywlambert.com · Articles · Briefings · Writing · LinkedIn
© 2026 Jerry W. Lambert Sr.