

\$10 / \$50 is not one product. Cache decides the real bill.

What this means for customers — spend, mix, clocks, and visibility.

Subscriber cut · Newsletter & public archive · On-demand · Scannable bullets

Executive summary

- \$10 / \$50 is not one product. Cache decides the real bill.
- OpenAI workhorse remains Sol promo, not Astra.
- Anthropic production seat remains Opus 5, not Fable 5.1.
- Several hard EOLs land in the next three weeks.
- Cost counters still sit under the \$10 frontier.

Prices

- OpenAI Sol promo: about \$4 in / \$20 out (short context), through at least late November.
- OpenAI Astra: about \$10 / \$50. Cache-read about \$1.
- Anthropic Fable 5.1: about \$10 / \$50. Cache-read about \$0.25.
- Anthropic Opus 5: about \$5 / \$25 (production seat).
- Google gemini-3-flash paid standard: about \$0.50 / \$3.00.
- Google gemini-3.8-flash intro: about \$0.75 / \$3.75 through year end, then steps up.
- Mistral Medium 3.5: about \$1.50 / \$7.50.
- DeepSeek V4: weekday peak / off-peak bands (UTC). Off-peak is the cheap agent ladder.
- xAI Grok Standard: about \$2 / \$6 under 200k tokens. Tools add about \$5 per 1,000 successful search or code calls.
- Cohere Command A+: live product, no clear public per-million chat tariff in this paper.
- Moonshot K3: market signal only here (about \$3 / \$15 at long context).

New models

- OpenAI Astra (‘gpt-6-astra’): frontier chat / API tier at \$10 / \$50. Strong capability lane. Cache-read near \$1 makes it materially more expensive than a twin-looking Fable list price when cache hits.
- Anthropic Fable 5.1: premium street tier at \$10 / \$50 with a low \$0.25 cache-read. Not the Opus production seat. Treat as adopt-on-purpose, not a silent upgrade.
- Google Gemini 3.8 Flash: intro pricing window now, then a step-up after year end. Cost-performance Flash ladder, not a \$10 frontier model.
- xAI Grok 4.6 (street): published workhorse card near \$2 / \$6. Different cache than seated grok-4.5. Do not treat 4.5 and 4.6 as one pin.

Deprecations

- 24 September 2026: OpenAI Videos / Sora.
- 27 September 2026: Perplexity Sonar Chat Completions ends. Move to Agent API. New work should be Agent-first.
- 28 September 2026: OpenAI legacy instruct models.
- 30 September 2026: Mistral Leanstral 1.5.
- 30 September 2026: Google omni preview slug.
- 2 October 2026: Google 2.5 Flash image.
- 2 November 2026: xAI grok-imagine-image-quality.

What this means for customers

Prices that can fool a budget

- The \$10 / \$50 twin trap: Astra and Fable 5.1 both list about \$10 / \$50, but Astra cache-read is about \$1 and Fable 5.1 is about \$0.25. Budget them as twins and you will mis-spend.
- Workhorse vs frontier: Sol promo about \$4 / \$20 is day-to-day spend. Astra at \$10 / \$50 is a layer above, not a Sol replacement. Keep them on separate cost centers.
- Hosted catalogs can lag the vendor street: some Agent tables still show Sol nearer \$5 / \$30. Use the vendor’s own pricing page for that vendor’s API.

Near clocks

- Move Perplexity work to the Agent API before 27 September.
- Retire Leanstral 1.5 traffic before 30 September.
- OpenAI Videos / Sora ends 24 September. Some legacy instruct models end 28 September.
- xAI image-quality slug retires 2 November.

Cheaper ladders

- Google Flash (including 3.8 intro then year-end step-up) for high-volume work.
- DeepSeek V4 peak / off-peak: schedule batch jobs in the printed UTC windows to cut cost.
- Mistral Medium 3.5 (~\$1.50 / \$7.50) when quality allows instead of frontier \$10 / \$50.
- Moonshot K3 is a market signal only until your own lane can bill it.

Visibility and spend control

- Token-only views understate agentic spend. xAI tools add about \$5 per 1,000 successful calls. Perplexity-class products add request fees by search depth.
- One vendor for every use case concentrates risk. Mix by job: batch vs interactive, cache-heavy vs one-shot, frontier vs cheap classification.
- Cohere Command A+ is serious for enterprise, but there is no clear public Chat \$/MTok card like Astra or Fable. Get an explicit rate card. Do not invent one.

Practical checklist this week

- Re-price any \$10 / \$50 comparison with cache-read (and cache-write if you use it).
- Keep Sol-class and Astra / Fable-class work on separate cost centers.
- Start Sonar → Agent and Leanstral migrations before the printed dates.
- Add tool and request fees to forecasts for search- and agent-heavy products.
- Require a metered invoice trail that names the exact model id and rate schedule.

What to watch

- Do not budget Astra and Fable 5.1 as the same \$10 / \$50 product.
- Do not treat Astra as Sol.
- Do not treat Fable 5.1 as the Opus production seat.
- Move Perplexity traffic to Agent before 27 September.

Briefing Board Members

- Chief of Staff · Grok (Grok Bot)
- Claude Seat · claude-opus-5
- OpenAI Seat · gpt-5.6-sol
- Grok Seat · grok-4.5
- Google Seat · gemini-3-flash
- DeepSeek Seat · deepseek-v4-flash
- Perplexity Seat · sonar-pro
- Mistral Seat · mistral-medium-3-5 (pending prove)

Get Jerry W. Lambert's Overnight Model Brief by email

<https://jerrywlambert.com/newsletter/>

[Unsubscribe](#)