Quiet overnight on first-party tariffs. Current picks hold.

Monday, 7 September 2026 · Subscriber cut

Executive summary

- Quiet overnight on first-party tariffs. Current picks hold.
- Hard clocks still run into late September.
- Two new names to watch. Not newly priced models.
- Keep Sol and Opus as production pins. Do not treat Astra or Fable 5.1 as those picks.

Prices

- OpenAI Sol promo: about $4 in / $20 out (short context). Still the OpenAI workhorse pin.
- OpenAI Astra: about $10 / $50. Cache-read about $1. Separate product.
- Anthropic Opus 5: about $5 / $25. Production pick.
- Anthropic Fable 5.1: about $10 / $50. Cache-read about $0.25. Street tier only.
- Google gemini-3-flash paid standard: about $0.50 / $3.00.
- Google gemini-3.8-flash intro: about $0.75 / $3.75 through year end, then steps up.
- Mistral Medium 3.5: about $1.50 / $7.50.
- DeepSeek V4: weekday peak / off-peak bands (UTC). Unchanged.
- xAI Grok Standard: about $2 / $6 under 200k tokens on published street cards.
- Zhipu GLM-5.3 Flash promo: 50% off ends 9 September 24:00 UTC+8 (about two days).

New models

- Muse Spark 1.3: new on the watch list. Names only. No settled public per-million chat tariff in this paper.
- GLM-5.3-Flash: new on the watch list. Open scoping. Promo clock above.
- No new first-party tariff rows overnight for Astra, Fable, Sol, Opus, or Gemini Flash.

Deprecations

- 9 September 2026: Zhipu GLM-5.3 Flash promo end (UTC+8).
- 14 September 2026: Amazon Nova Premier EOL.
- 24 September 2026: OpenAI Videos / Sora.
- 27 September 2026: Perplexity Sonar Chat Completions ends. Move to Agent API.
- 28 September 2026: OpenAI legacy instruct models.
- 30 September 2026: Mistral Leanstral 1.5.
- 2 November 2026: xAI grok-imagine-image-quality.

What this means for customers

- No first-party list-price rewrite overnight. Sol short promo still about $4 / $20. Opus 5 still about $5 / $25. Astra still about $10 / $50 short, a layer above Sol, not a substitute.
- Cache still separates matching headline prices. Astra cache-read about $1 vs Fable 5.1 about $0.25. Budget the cache column or agent workloads mis-forecast.
- About two days: Z.ai GLM-5.3 Flash promo ends 9 September (UTC+8 end-of-day). After that, expect list rates.
- About seven days: Amazon Nova Premier EOL 14 September. Stop new dependency now.
- About three weeks: Perplexity Sonar Chat Completions support ends 27 September. Agent API for new work.
- Month-end: Leanstral 1.5 and related preview EOLs around 30 September.
- Research census moved (+2 names): Muse Spark 1.3 and GLM-5.3-Flash. Awareness only, not a first-party tariff change. Do not budget them until the vendor publishes a price and a billed id your gateway can meter.
- Quiet street days are when clocks and shapes bite. Schedule DeepSeek off-peak only with the printed UTC windows.
- Token-only dashboards still understate search and tool agents (request fees / per-1k tool calls).
- One-vendor concentration risk is unchanged. Mix ladders (Flash / mid-tier / frontier) so one EOL or promo cliff cannot hit every use case.
- Checklist: re-check GLM Flash promo traffic before 9 September. Drain Nova Premier paths before 14 September. Keep Sonar to Agent migration on the calendar for 27 September. Keep Sol-class and Astra/Fable-class work on separate cost centers.

What to watch

- Do not treat Astra as Sol.
- Do not treat Fable 5.1 as the Opus production pick.
- Do not invent a Muse or Seed chat tariff without a first-party price card.

Prices are per million tokens unless the sentence says otherwise.

Get Jerry W. Lambert's Overnight Model Brief by email
https://jerrywlambert.com/newsletter/

Unsubscribe