

Daily Model Brief, October 6, 2026

Mistral Large 4 lands at a discount. Three OpenAI models retire October 22.

Tuesday, October 6, 2026 · Prices as of 1:00 PM CT

Executive Summary

Mistral Large 4 (`mistral-large-4`) is live in Public Preview at Mistral's discounted rate of \$0.68 input, \$0.07 cached and \$2.09 output per million tokens (list \$1.36 / \$0.14 / \$4.18), with a 1M-token context window. The existing `mistral-large` line stays on Large 3 at \$0.50 / \$1.50.

OpenAI will shut down `gpt-4.1-nano`, `o4-mini` and `gpt-image-1` on Thursday, October 22, 2026 at 7:00 PM CT (OpenAI lists October 23 with no time).

Google released Nano Banana 2.1 (`gemini-nano-banana-2.1`), set an October 29 shutdown for `gemini-3.1-flash-image` and set November 17 shutdowns for five preview audio models. Perplexity listed a second Decisions API model, and Alibaba posted a US (Virginia) price row for `qwen3.8-flash`.

No published price changed on the established production models tracked in this brief.

What This Means for Customers

- If you adopt Mistral Large 4, budget the sale price and plan for a return to list. Mistral prints no end date on the discount, and list is double the sale rate.
- Do not confuse Large 4 with the existing `mistral-large` line. That line still bills at Large 3 prices.
- Work still on `gpt-4.1-nano`, `o4-mini` or `gpt-image-1` needs a migration plan before the evening of October 22 (CT). Move to the named successors, not a silent remap.
- Nano Banana 2.1 is an image model. Image output at \$30.00 per million tokens works out to about \$0.0336 per 1K image, \$0.0504 per 2K image and \$0.0756 per 4K image. Batch is half price. Treat it as a new image product, not a chat substitute.
- Perplexity's Fast Search option for `web_search` is \$0.001 per call beside the standard \$0.0025. That is a per-call tool fee, not a token rate, so token-only dashboards will understate search-heavy agents.
- Spread workloads across vendors and tiers so one retirement week (OpenAI October 22, Google image October 29, Google audio November 17) cannot hit every use case at once.

New Models and Price Changes

Prices in USD per 1M tokens.

Vendor	Model ID	Status	Input	Cached	Output
Mistral	mistral-large-4	Public Preview, released Oct 6. 1M-token context.	\$0.68 (list \$1.36)	\$0.07 (list \$0.14)	\$2.09 (list \$4.18)
Google	gemini-nano-banana-2.1	Generally available, released Oct 6. Image model.	\$1.50	Not listed	\$7.50 text \$30.00 image
Perplexity	pplx-decider-v1.1-27b	Decisions API model, new on the pricing page.	\$0.04	Not listed	Free
Alibaba	qwen3.8-flash (US, Virginia)	New regional price row posted Oct 6.	\$0.15	Not listed	\$0.47

- Mistral shows the Large 4 rates as a sale price against the list price, with no end date. Mistral Large 3 stays at \$0.50 input and \$1.50 output.
- Nano Banana 2.1 batch rates are half the standard rates.
- Anthropic, OpenAI, xAI, DeepSeek, Moonshot, Z.ai, MiniMax, Meta and TypeSafe pricing, model and deprecation pages were unchanged since the morning check.

Deprecations and Retirements

Vendor	Model ID	Shutdown (CT)	Replacement	Note
OpenAI	gpt-4.1-nano o4-mini gpt-image-1	Thu Oct 22, 2026, 7:00 PM CT	gpt-5.6-luna gpt-5.6-terra gpt-image-2.5-sunburst or gpt-image-2.5-flare	OpenAI lists Oct 23 with no time; 7:00 PM CT Oct 22 is 00:00 UTC Oct 23.
Google	gemini-omni-flash-preview	Thu Oct 22, 2026 (time not stated)	gemini-omni-1.1-flash	Newly listed today.
Google	veo-3.1-generate-preview veo-3.1-fast-generate-preview veo-3.1-lite-generate-preview	Thu Oct 22, 2026 (time not stated)	gemini-omni-1.1-flash	Unchanged since the morning check.
Google	gemini-3.1-flash-image	Thu Oct 29, 2026 (time not stated)	gemini-nano-banana-2.1	New today.
Google	gemini-3.1-flash-tts-preview gemini-2.5-flash-preview-tts gemini-2.5-pro-preview-tts	Tue Nov 17, 2026 (time not stated)	gemini-3.8-flash-tts or gemini-3.8-flash-lite-tts	New today. Previously no shutdown date.
Google	gemini-3.1-flash-live-preview gemini-2.5-flash-native-audio-preview-12-2025	Tue Nov 17, 2026 (time not stated)	gemini-3.8-live	New today. Previously no shutdown date.
Anthropic	claude-sonnet-4-5-20250929	Mon Nov 30, 2026 (time not stated)	claude-sonnet-5-5	Unchanged.
Anthropic	claude-haiku-4-5-20251001	Not sooner than Thu Oct 15, 2026; no date set	None named	Watch item. Active today.

- Unverified: Google's deprecations page now shows gemini-2.5-flash-image shutting down Mon Mar 15, 2027, with gemini-3.1-flash-lite-image as the replacement, while Google's pricing page still says it shut down Oct 2, 2026. Confirm with Google before you budget either date.
- Perplexity will retire its Agent API and Router API routes to older OpenAI models (openai/gpt-5.4, gpt-5.4-mini, gpt-5.4-nano, gpt-5.2, gpt-5.1, gpt-5 and gpt-5-mini) at 7:00 PM CT on Fri Oct 23, 2026.

Around the Market

The day's first-party move is Mistral putting Large 4 on its pricing page at a clear sale against list, with no end date printed. That is a new model, not a rename of Large 3. OpenAI's October 22 clock for three models is the hard calendar item for anyone still on gpt-4.1-nano, o4-mini or gpt-image-1. Google's update adds an image successor in Nano Banana 2.1 and fills in shutdown dates for image and audio previews that previously had none. Perplexity's Decisions line and Fast Search tool fee are product-shape news more than a frontier price change, and Alibaba's US

Virginia row for qwen3.8-flash is a regional posting. Anthropic, xAI, DeepSeek, Moonshot, Z.ai, MiniMax, Meta and TypeSafe left published rates alone.

Bottom Line

Adopt Large 4 only if you want that class and can absorb a return to list. Drain gpt-4.1-nano, o4-mini and gpt-image-1 paths before the evening of October 22 (CT). Treat Nano Banana 2.1 as a new image product with its own cost shape, and keep Large 3 and Large 4 on separate cost lines.

All times Central. Prices, dates and model IDs are taken from first-party vendor pricing, model and deprecation pages as of 1:00 PM CT on October 6, 2026; performance scores are from the Artificial Analysis Intelligence Index v4.3. Items that could not be confirmed are marked unverified.


© 2026 Jerry W. Lambert Sr.

Model Size vs. Price

JerryWLambert.com AI Briefing

Model Size vs. Price: Active Models in Our Registry

As of Tue Oct 6, 2026 · 20 models with vendor-published parameter counts · published per-1M-token prices



Parameter counts are vendor-published (model pages, docs and official Hugging Face model cards). Prices as of Oct 6, 2026. 20 models plotted; 62 not plotted (no published count).
Not plotted, by vendor: Anthropic 11 · OpenAI 14 · xAI 8 · Google 10 · Z.ai 9 · Cohere 2 · Alibaba 2 · Amazon 3 · TypeSafe 3. Anthropic, OpenAI, xAI and Google publish no parameter counts for these models.
Excluded: 13 retiring or unapproved models (gpt-4.1-nano, o4-mini and gpt-image-1 retire Oct 22; o3 and o3-pro retire Dec 11; 8 unapproved). MoE models are plotted at total parameters, with active parameters in the label.
DeepSeek prices are off-peak. Mistral Large 4 is at Mistral's discounted rate (\$1.36 / \$4.18 list). The command-a-plus price is unconfirmed by Cohere. Same-price, same-size models share one dot

Across the 20 active models whose vendors publish parameter counts, price tracks active parameters far more closely than total size. Mistral Large 4 (1.05T total, 52B active) sits mid-pack at \$0.68 input and \$2.09 output per 1M tokens.

Price vs. Performance


JerryWLambert.com AI Briefing

Price vs. Performance: Active Models in Our Registry


Published per-1M-token prices as of Oct 6, 2026, 1:31 PM CT · Intelligence Index scores fetched Oct 6, 2026, 1:33 PM CT · 44 models with a measured score

Anthropic ● OpenAI ● xAI ● Google ● DeepSeek ● Mistral ● Moonshot ● Z.ai ● Cohere ● Alibaba

Output Price vs. Performance



Input Price vs. Performance



Performance: Artificial Analysis Intelligence Index v4.3, fetched Oct 6, 2026, 1:33 PM CT. Each model uses its headline entry (the highest tested reasoning setting where several are tested).

45 active models plotted as 44 dots (aliases with the same model and price share one dot). Not plotted (37): Anthropic 3 · OpenAI 4 · xAI 4 · Google 4 · Z.ai 13 · Cohere 3 · Amazon 3 · TypeSafe 3; 20 have only an estimated score, 17 have no score.

Excluded: 13 retiring or unapproved models. Prices are published list rates per 1M tokens, not Artificial Analysis prices.

DeepSeek prices are off-peak. Mistral Large 4 is at Mistral's discounted rate (\$1.36 / \$4.18 list). deepseek-flash and deepseek-v4-flash are scored as DeepSeek V4.1 Flash, which DeepSeek says now serves both names.

Across the 45 active models with a measured Artificial Analysis Intelligence Index v4.3 score (fetched Oct 6, 2026, 1:33 PM CT), price rises with performance only loosely (rank correlation 0.45 on output price). Several models scoring near 40 cost \$0.60 or less per 1M output tokens, including GLM-5.3-Flash, GPT-6 Luna and DeepSeek V4.1 Flash, while the top scorer, Claude Opus 5.5 at 57.6, costs \$20.