

Mistral Large 4 lands at a discount. Three OpenAI models retire October 22.

Tuesday, 6 October 2026 · As of 1:00 PM Central · Jerry W. Lambert Sr.

Executive summary

- Mistral Large 4 (`mistral-large-4`) is live in Public Preview at Mistral's discounted rate of \$0.68 input, \$0.07 cached and \$2.09 output per million tokens (list \$1.36 / \$0.14 / \$4.18). Context window: 1M tokens. The existing `mistral-large` line stays on Large 3 at \$0.50 / \$1.50.
- OpenAI will shut down `gpt-4.1-nano`, `o4-mini` and `gpt-image-1` on Thursday, October 22, 2026 at about 7:00 PM Central (OpenAI lists October 23 with no time). Named replacements: `gpt-5.6-luna`, `gpt-5.6-terra`, and `gpt-image-2.5-sunburst` or `gpt-image-2.5-flare`.
- Since the morning check, Google released Nano Banana 2.1 (`gemini-nano-banana-2.1`), set an October 29 shutdown for `gemini-3.1-flash-image`, and set November 17 shutdowns for five preview audio models. Perplexity listed a second Decisions API model. Alibaba posted a US (Virginia) price row for `qwen3.8-flash`.
- No published price moved on the established production lines tracked as billed today.

What this means for customers

- If you adopt Mistral Large 4, budget the sale price and watch for a return to list. There is no printed end date on the discount. List is double the sale input and output.
- Do not confuse Large 4 with the existing `mistral-large` pin. That pin still bills Large 3 dollars.
- OpenAI work still on `gpt-4.1-nano`, `o4-mini` or `gpt-image-1` needs a migration plan before October 22 evening Central. Prefer the named successors above, not a silent remap.
- Google Nano Banana 2.1 is an image model at \$1.50 input and \$7.50 text / \$30.00 image output per million tokens. Image output of \$30.00 per million tokens equals about \$0.0336 per 1K image, \$0.0504 per 2K and \$0.0756 per 4K. Batch is half. Treat it as a new image product, not a chat substitute.
- Perplexity's Decisions model (`pplx-decider-v1.1-27b`) lists \$0.04 input and free output. Separately, Fast Search on `web_search` is \$0.001 per call beside the standard \$0.0025. That is a per-call tool fee, not a token rate.
- Alibaba's new US Virginia row for `qwen3.8-flash` is \$0.15 / \$0.47. Regional rows are not a global tariff rewrite.
- Token-only dashboards still understate tool and search agents. Budget the tool column when Perplexity search is in the path.
- Mix ladders so one retirement week (OpenAI October 22, Google image October 29, Google audio November 17) cannot hit every use case at once.

New models and published prices

Prices in USD per 1M tokens, from first-party vendor pricing pages as of 1:00 PM Central.

Vendor	Model	Status	Input	Cached	Output
Mistral	mistral-large-4	Public Preview, released Oct 6	\$0.68 (list \$1.36)	\$0.07 (list \$0.14)	\$2.09 (list \$4.18)
Google	gemini-nano-banana-2.1	Generally available, released Oct 6. Image model	\$1.50	Not listed	\$7.50 text / \$30.00 image
Perplexity	pplx-decoder-v1.1-27b	Decisions API model, new on the pricing page	\$0.04	Not listed	Free
Alibaba	qwen3.8-flash (US, Virginia)	New regional price row posted Oct 6	\$0.15	Not listed	\$0.47

- Anthropic, OpenAI, xAI, DeepSeek, Moonshot, Z.ai, MiniMax, Meta and TypeSafe pricing, model and deprecation pages were unchanged since the morning check.

Deprecations and retirements

- **OpenAI** — gpt-4.1-nano, o4-mini, gpt-image-1: Thu Oct 22, 2026, 7:00 PM Central. Replacements: gpt-5.6-luna, gpt-5.6-terra, gpt-image-2.5-sunburst or gpt-image-2.5-flare.
- **Google** — gemini-omni-flash-preview: Thu Oct 22, 2026 (time not stated). Replacement: gemini-omni-1.1-flash. Newly listed today.
- **Google** — veo-3.1-generate-preview, veo-3.1-fast-generate-preview, veo-3.1-lite-generate-preview: Thu Oct 22, 2026 (time not stated). Replacement: gemini-omni-1.1-flash.
- **Google** — gemini-3.1-flash-image: Thu Oct 29, 2026 (time not stated). Replacement: gemini-nano-banana-2.1. New today.
- **Google** — gemini-3.1-flash-tts-preview, gemini-2.5-flash-preview-tts, gemini-2.5-pro-preview-tts: Tue Nov 17, 2026 (time not stated). Replacement: gemini-3.8-flash-tts or gemini-3.8-flash-lite-tts. New today.
- **Google** — gemini-3.1-flash-live-preview, gemini-2.5-flash-native-audio-preview-12-2025: Tue Nov 17, 2026 (time not stated). Replacement: gemini-3.8-live. New today.
- **Anthropic** — claude-sonnet-4-5-20250929: Mon Nov 30, 2026 (time not stated). Replacement: claude-sonnet-5-5.
- **Anthropic** — claude-haiku-4-5-20251001: not sooner than Thu Oct 15, 2026; no hard date set. Active today. Watch item.
- Unverified conflict: Google's deprecations page now shows gemini-2.5-flash-image shutting down Mon Mar 15, 2027, while the pricing page still says it shut down Oct 2, 2026. Confirm on Google's pages before you budget either date.
- Perplexity will retire Agent API and Router API routes to older OpenAI models (openai/gpt-5.4, gpt-5.4-mini, gpt-5.4-nano, gpt-5.2, gpt-5.1, gpt-5 and gpt-5-mini) at 7:00 PM Central on Fri Oct 23, 2026.

Around the market


The day's first-party move is Mistral putting Large 4 on the pricing page at a clear sale vs list, with no end date printed. That is a new billed id, not a rename of Large 3. OpenAI's October 22 clock for three models is the hard calendar item for anyone still on nano, o4-mini or gpt-image-1. Google's afternoon update adds an image successor (Nano Banana 2.1) and fills in shutdown dates for image and audio previews that previously lacked them. Perplexity's Decisions line and Fast Search tool fee are product-shape news more than a frontier price rewrite. Alibaba's US Virginia row for qwen3.8-flash is a regional posting. Anthropic, xAI, DeepSeek, Moonshot, Z.ai, MiniMax, Meta and TypeSafe left published rates alone in this window.

Model size vs. price

JerryWLBert.com AI Briefing

Model Size vs. Price: Active Models in Our Registry

As of Tue Oct 6, 2026 · 20 models with vendor-published parameter counts · prices from our live rate card




Across the 20 active models whose vendors publish parameter counts, price tracks active parameters far more closely than total size. Mistral Large 4 (1.05T total, 52B active) sits mid-pack at \$0.68 input and \$2.09 output per 1M tokens.

Price vs. performance

JerryWLBert.com AI Briefing

Price vs. Performance: Active Models in Our Registry

Prices from our live rate card band 0, 95 models in force, queried Oct 6, 2026, 1:31 PM CT · scores fetched Oct 6, 2026, 1:33 PM CT · 44 models with a measured score



Across the 45 active models with a measured Artificial Analysis Intelligence Index v4.3 score (fetched Oct 6, 2026), price rises with performance only loosely (rank correlation 0.45 on output price). Several models scoring near 40 cost \$0.60 or less per 1M output tokens, including GLM-5.3-Flash, GPT-6 Luna and DeepSeek V4.1 Flash, while the top scorer, Claude Opus 5.5 at 57.6, costs \$20.

Bottom line

Adopt Large 4 only if you want that class and can absorb a return to list. Drain OpenAI nano, o4-mini and gpt-image-1 paths before October 22 evening Central. Treat Nano Banana 2.1 as a new image product with its own dollar shape. Leave Large 3 and Large 4 on separate cost centers.