Daily Model Brief, October 12, 2026

Google retired gemini-3.5-flash and gemini-3.7-flash from its price page and now routes them to 3.6 and 3.8 Flash. No other list-price change forces a budget rewrite today; Z.ai is on watch.

Monday, October 12, 2026 · Prices as of 8:54 AM CT

Executive Summary

Google retired two Gemini Flash models. Its pricing page (last updated 2026-10-09 UTC, fetched this morning) dropped the price sections for gemini-3.5-flash and gemini-3.7-flash, and the models page dropped both ids. Google's deprecations page now says requests to gemini-3.5-flash are automatically routed to gemini-3.6-flash, and requests to gemini-3.7-flash go to gemini-3.8-flash, with no shutdown date announced.

`gemini-3.5-flash` was listed at $1.50 input and $9.00 output per million tokens. The model Google now serves for that id (`gemini-3.6-flash`) lists at $0.75 / $3.75, about half the old price. `gemini-3.7-flash` requests now go to `gemini-3.8-flash`.

Z.ai's pricing page dropped 10 model ids (the GLM-4.5 family, GLM-4.7-FlashX, GLM-4-32B, GLM-4.6V / 4.6V-FlashX / 4.5V, and GLM-OCR) with no deprecation notice. Their model guides are still up, so these models look at risk but are not confirmed retired.

No other vendor changed a published token price that forces a budget rewrite. Anthropic's list rates are unchanged, but its pricing page now says Claude 4.7 and later (and Mythos Preview) use a newer tokenizer that produces about 30% more tokens for the same text, which raises effective cost per task at the same list rate. OpenAI listed gpt-rosalind-discovery at $5 / $25, but it is trusted-access only and stays off.

Alibaba list prices for qwen3.7-max and qwen3.7-plus are unchanged. TypeSafe Jev list prices are unchanged. IBM Granite holds at $0.0636 / $0.265. Amazon and Microsoft price pages still fail; ByteDance is still thin; Cohere still publishes no token price.

What This Means for Customers

- If you still call `gemini-3.5-flash` or `gemini-3.7-flash`, Google is already routing those requests to `gemini-3.6-flash` and `gemini-3.8-flash`. Move callers to the model you intend to use and budget at its price.
- If you price Claude 4.7 or later from list rates alone, re-check task cost. Anthropic says the newer tokenizer produces about 30% more tokens for the same text, so effective cost per task rises even when the published $/MTok does not.
- Anthropic also now states that its US-only inference_geo 1.1x multiplier applies to the Haiku 5.5 band for prompts over 100,000 tokens.
- OpenAI's gpt-rosalind-discovery is listed but trusted-access only. Do not budget it as a general-availability seat.
- Z.ai: ten model ids disappeared from its price page with no retirement notice. Treat them as at risk of retirement and have a fallback ready.
- October 22 and 23 hold. `gpt-4.1-nano`, `o4-mini` and `gpt-image-1` (OpenAI) shut down the evening of October 22 (CT), along with Google's `gemini-omni-flash-preview` and Veo 3.1 previews. Google's deep-research-pro preview shuts down October 23. Perplexity retires its routes to older OpenAI models at 7:00 PM CT on October 23.
- Claude Haiku 4.5 is still active, with retirement not sooner than October 15 and no date set. Keep a tested fallback ready.

New Models and Price Changes

Prices in USD per 1M tokens.

Vendor	Model ID	Status	Input	Cached	Output
Google	`gemini-3.5-flash`	Retired from price page. Google routes requests to gemini-3.6-flash, which lists at about half the price.	$1.50 (served: $0.75)	$0.15 ($0.075)	$9.00 ($3.75)
Google	`gemini-3.7-flash`	Retired from price page. Google routes requests to gemini-3.8-flash.	$0.75	$0.075	$3.75
OpenAI	`gpt-rosalind-discovery`	Newly listed beside gpt-rosalind-research; trusted-access only.	$5.00	$0.50	$25.00
Anthropic	`claude-4.7+ / Mythos Preview`	List $/MTok unchanged. Newer tokenizer (~30% more tokens for the same text) raises effective cost per task.	Unchanged	Unchanged	Unchanged
Alibaba	`qwen3.7-max`	Unchanged (international list price).	$2.50	Not listed	$7.50
Alibaba	`qwen3.7-plus`	Unchanged (international list price).	$0.40	Not listed	$1.60
TypeSafe	`jev-1.13.0 / jev-latest / jev-preview`	Unchanged.	$0.042	—	$0
IBM	`granite-4-h-small`	Unchanged.	$0.0636	Not listed	$0.265

- Google's deprecations page says requests to gemini-3.5-flash route to gemini-3.6-flash and requests to gemini-3.7-flash route to gemini-3.8-flash. No shutdown date is announced for either id.
- Z.ai dropped from its pricing page (still no deprecation notice; model guides still 200): glm-4.7-flashx, glm-4.5, glm-4.5-x, glm-4.5-air, glm-4.5-airx, glm-4-32b-0414-128k, glm-4.6v, glm-4.6v-flashx, glm-4.5v, glm-ocr. Remaining Z.ai list prices (glm-5.3 / 5.3-flash / 5.3-flashx / 5.2 / 5.1 / 5 / 4.7 / 4.6) unchanged. On watch.
- Anthropic also revised tool-overhead token counts (text_editor_20250728: 974 / 745 tokens) and now states inference_geo 1.1x applies to the Haiku 5.5 >100k band. Mythos remains limited access.
- xAI, DeepSeek, Moonshot, MiniMax, Mistral, Perplexity and Meta published token prices were unchanged since Friday's brief. Mistral Large 4 stays on sale at $0.68 input and $2.09 output, with no end date.
- Amazon and Microsoft price pages could not be read today, ByteDance published no usable price table, Cohere still publishes no token prices, and IBM docs returned 403 on supported-models. Their last confirmed prices stand.

Deprecations

Vendor	Model ID	Shutdown (CT)	Replacement	Note
Google	`gemini-3.5-flash`	No shutdown date announced	`gemini-3.6-flash (Google auto-route)`	New today. Pricing and models pages dropped the id (page updated 2026-10-09 UTC).
Google	`gemini-3.7-flash`	No shutdown date announced	`gemini-3.8-flash (Google auto-route)`	New today. Same pages.
Google	`deep-research-pro-preview-12-2025`	Fri Oct 23, 2026 (time not stated)	`deep-research-preview-04-2026`	Unchanged.

OpenAI	`gpt-4.1-nano`
	`o4-mini`
	`gpt-image-1`
		Thu Oct 22, 2026, 7:00 PM CT	`gpt-5.6-luna`
			`gpt-5.6-terra`
			`gpt-image-2.5-sunburst or gpt-image-2.5-flare`
				Unchanged. OpenAI lists Oct 23 with no time; 7:00 PM CT Oct 22 is 00:00 UTC Oct 23.

Google	`gemini-omni-flash-preview`	Thu Oct 22, 2026 (time not stated)	`gemini-omni-1.1-flash`	Unchanged.

Vendor	Model ID	Shutdown (CT)	Replacement	Note

Google	`veo-3.1-generate-preview`
	`veo-3.1-fast-generate-preview`
	`veo-3.1-lite-generate-preview`
		Thu Oct 22, 2026 (time not stated)	`gemini-omni-1.1-flash`	Unchanged.

Google	`gemini-3.1-flash-image`	No shutdown date announced	`gemini-nano-banana-2.1`	Unchanged.

Google	`gemini-3.1-flash-tts-preview`
	`gemini-2.5-flash-preview-tts`
	`gemini-2.5-pro-preview-tts`
		Tue Nov 17, 2026 (time not stated)	`gemini-3.8-flash-tts` or `gemini-3.8-flash-lite-tts`	Unchanged.

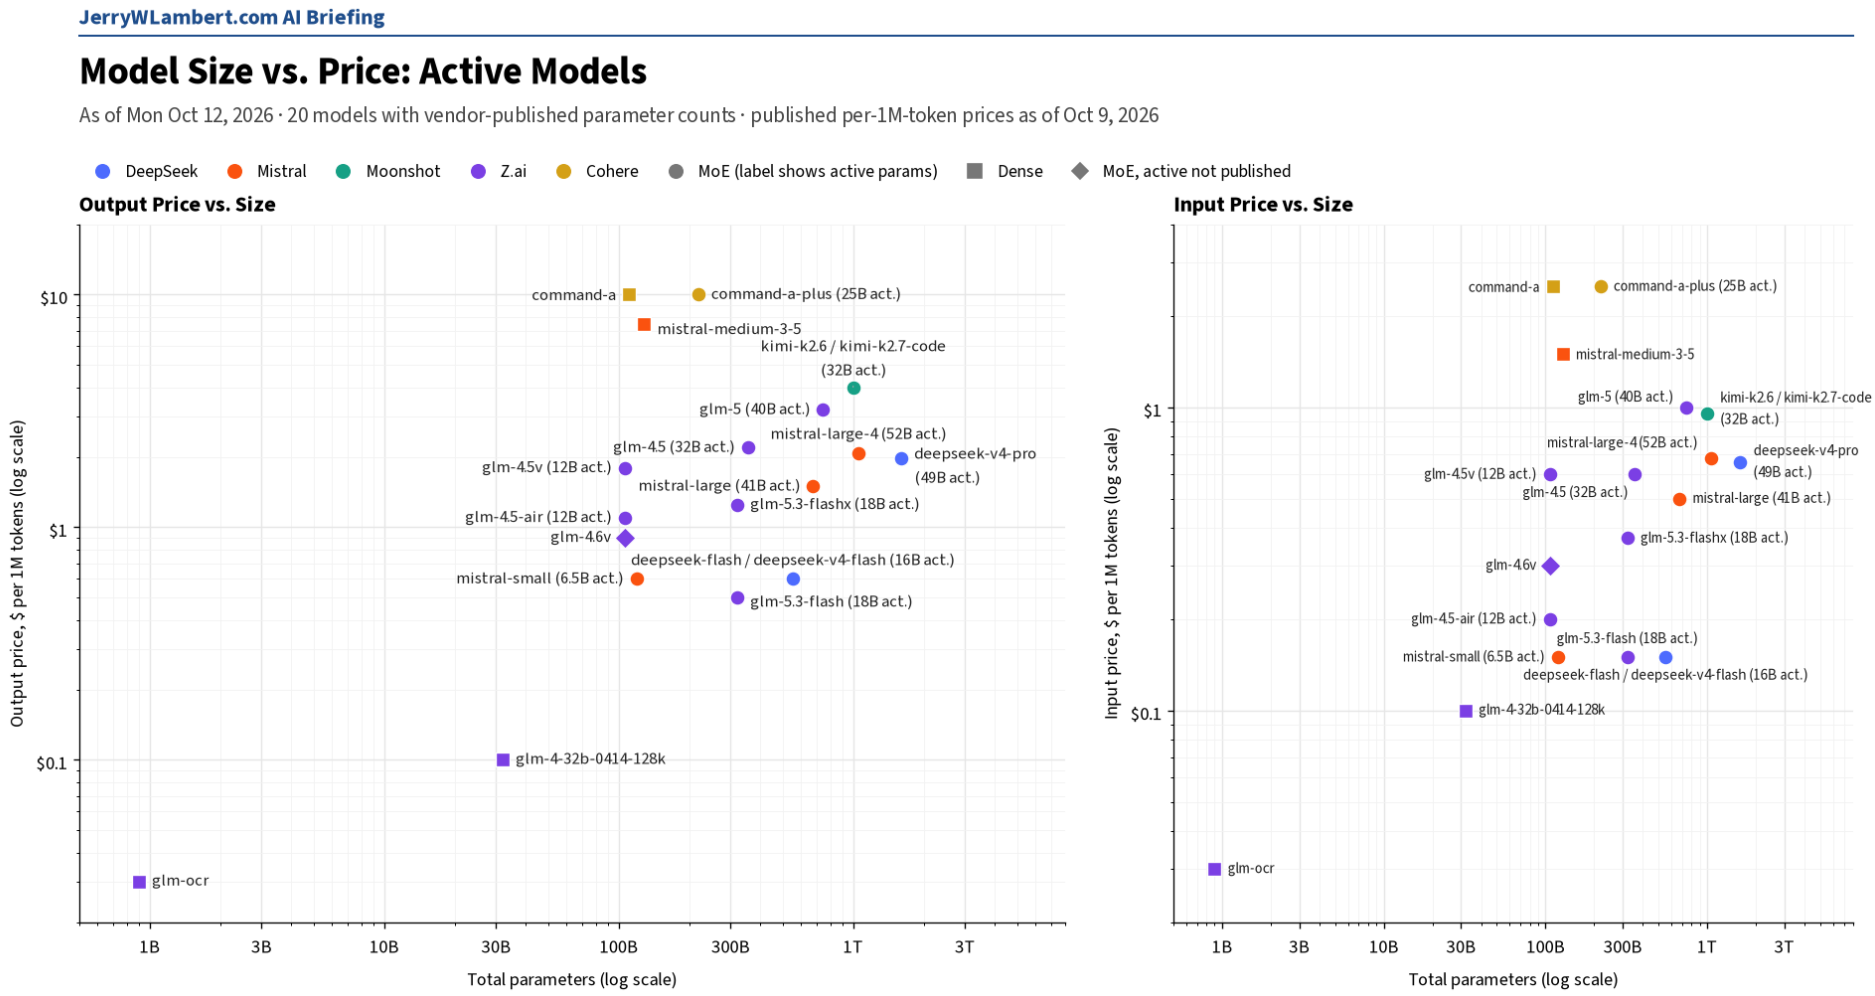
Google	`gemini-3.1-flash-live-preview` `gemini-2.5-flash-native-audio-preview-12-2025`	Tue Nov 17, 2026 (time not stated)	`gemini-3.8-live`	Unchanged.
Anthropic	`claude-sonnet-4-5-20250929`	Mon Nov 30, 2026 (time not stated)	`claude-sonnet-5-5`	Unchanged.
Anthropic	`claude-haiku-4-5-20251001`	Not sooner than Thu Oct 15, 2026; no date set	`None named`	Watch item. Still active.

Cohere	`command-r`
	`command-r-plus`
		Deprecated Sep 15, 2025; no shutdown date set	`command-r-08-2024`
			`command-r-plus-08-2024`
			`command-a-03-2025`
				Closed to new customers; existing users keep access until retirement.

- Perplexity will retire its Agent API and Orchestrator API routes to older OpenAI models (openai/gpt-5.4, gpt-5.4-mini, gpt-5.4-nano, gpt-5.2, gpt-5.1, gpt-5 and gpt-5-mini) at 7:00 PM CT on Fri Oct 23, 2026 (00:00 UTC Oct 24). Unchanged.
- Google also lists deep-research-preview-04-2026 and deep-research-max-preview-04-2026 with no shutdown date announced.
- Unverified: Google's deprecations page still shows gemini-2.5-flash-image shutting down Mon Mar 15, 2027, with gemini-3.1-flash-lite-image as the replacement, while Google's pricing page says it shut down Oct 2, 2026. Confirm with Google before you budget either date.

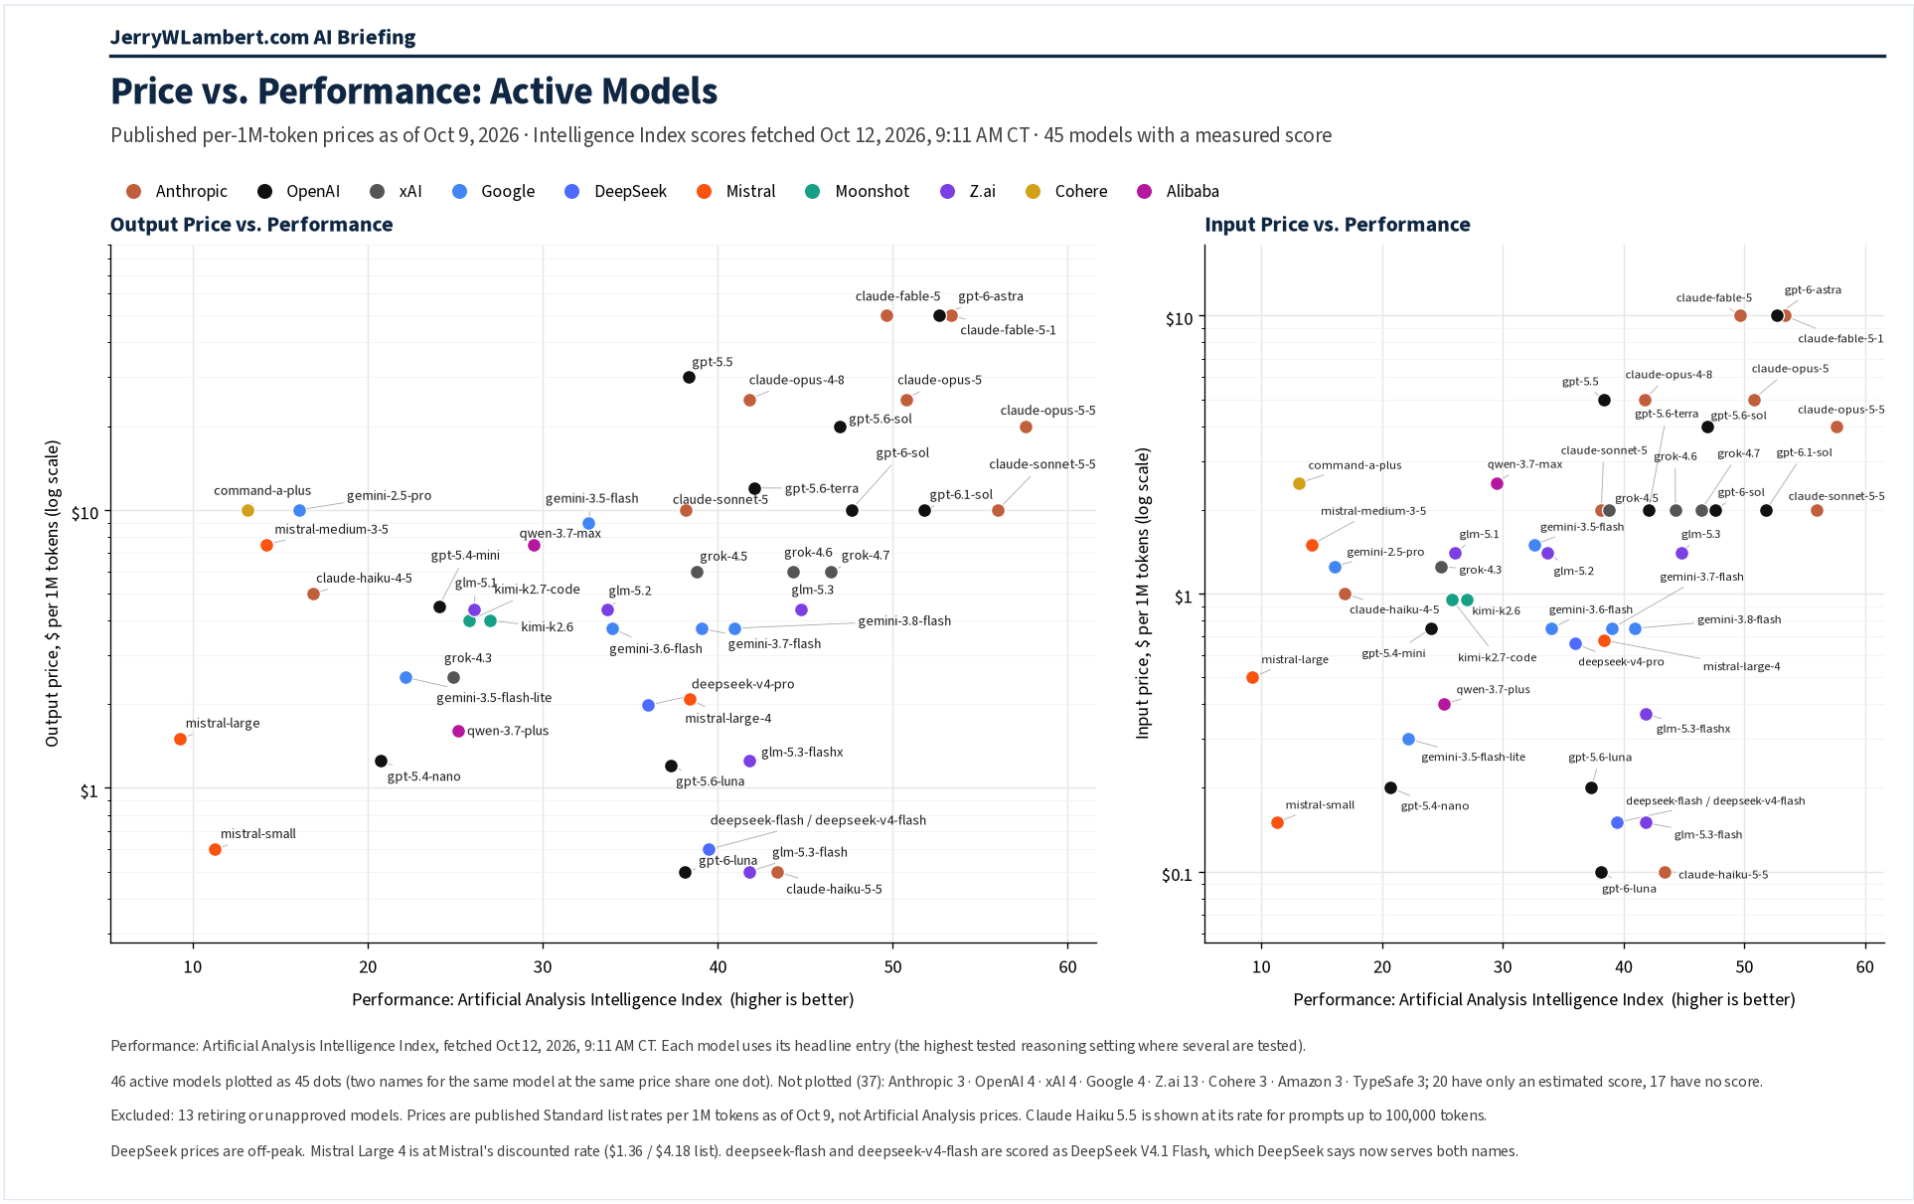
All times Central. Vendor prices, dates and model IDs are taken from first-party vendor pricing, model, deprecation and changelog pages as of 8:54 AM CT on October 12, 2026. Chart prices are as of October 9, 2026. Performance scores are from the Artificial Analysis Intelligence Index, fetched 9:11 AM CT on October 12, 2026. Items that could not be confirmed are marked unverified.
© 2026 Jerry W. Lambert Sr.

Model Size vs. Price



Across the 20 active models whose vendors publish parameter counts, price rises only loosely with total size (rank correlation 0.39 on output price). Anthropic, OpenAI, xAI and Google publish no parameter counts, so their models, including Claude Haiku 5.5 and GPT-6.1 Sol, are not plotted. Prices as of October 9, 2026.

Price vs. Performance



Across the 46 active models with a measured Artificial Analysis Intelligence Index score (fetched Oct 12, 2026, 9:11 AM CT), price rises with performance only loosely (rank correlation 0.43 on output price). Claude Haiku 5.5 still leads at $0.50 or less per million output tokens (43.4). GPT-6.1 Sol scores 51.8 at $10 Standard. Chart prices are as of October 9, 2026.