Daily Model Brief, October 15, 2026

TypeSafe shipped Jev 1.14 at the same price as 1.13, while Meta's Muse Wonder, Google's Gemini 4 Argon launch price, and OpenAI's January 2027 shutdowns for GPT-4.1 and GPT-4o remain on the vendor pages.

Thursday, October 15, 2026 · Prices as of 8:26 AM CT

Executive Summary

TypeSafe shipped Jev 1.14 (`jev-1.14.0`), available from October 14, 2026 at the same price as 1.13: $0.042 per million input tokens, output free. The moving name `jev-preview` now resolves to 1.14; `jev-latest` still resolves to 1.13.

Google's Gemini 4 Argon (`gemini-4-argon`) is unchanged on the pricing page (Last updated October 15, 2026 UTC): $2.00 input, $0.10 cached and $10.00 output per million tokens through November 30, then $4.00 / $0.20 / $20.00 from December 1. Google also added an invite-only Contributor tier at half of Standard for some models; it is not a public price list add.

Meta's Muse Wonder (`muse-wonder-1.0`) is unchanged since it opened on the Model API on October 13, at $4.00 input, $0.20 cached and $20.00 output. Muse Spark is unchanged at $1.25 / $0.15 / $4.25.

OpenAI's January 14, 2027 shutdown for GPT-4.1 (`gpt-4.1-2025-04-14`), GPT-4.1 mini (`gpt-4.1-mini-2025-04-14`), GPT-4o and GPT-4o mini still stands, with `gpt-6-luna` as the named replacement. Prices on those models are unchanged. Shown here as Wednesday, January 13, 2027, 6:00 PM CT (00:00 UTC January 14).

Anthropic, xAI, DeepSeek, Moonshot, Z.ai, Alibaba, MiniMax, Mistral and Perplexity published no chat or token price changes. xAI renamed a voice SKU to grok-voice-3.0 at the same per-minute rate; that is off the text price list.

TypeSafe Jev 1.13 stays $0.042 input and free output. IBM Granite is unchanged at $0.0636 / $0.265 and could not be fully verified today (docs returned 403). Amazon's Bedrock price pages still fail to load, ByteDance still shows no usable token table, Microsoft's Foundry catalog loads but still lists no pay-as-you-go prices, and Cohere still publishes no full token matrix.

What This Means for Customers

- If you use TypeSafe System One classifiers, `jev-1.14.0` is live at the same $0.042 input / free output rate as 1.13. Point `jev-preview` at 1.14 if you want the newest release; keep `jev-latest` if you want to stay on 1.13.
- If you want to try `gemini-4-argon`, the launch price still runs through November 30. Budget at $4.00 input and $20.00 output per million tokens from December 1, and remember that thinking tokens bill as output. Contributor pricing is invite-only at half of Standard for Argon and some Flash models.
- If you call `gpt-4o`, `gpt-4o-mini`, `gpt-4.1` or `gpt-4.1-mini`, you still have until January 14, 2027. Each GPT-4.1 model lists only one snapshot, so the plain names stop working then too. Test `gpt-6-luna` now; it lists at $0.10 input, $0.01 cached and $0.50 output per million tokens (short context), well below gpt-4.1 at $2.00 / $8.00.
- OpenAI still points developers to an optional shutdown_date field on its List models and Retrieve model endpoints, so you can watch for retirement dates programmatically.
- The October 22 and 23 dates are unchanged. `gpt-4.1-nano`, `o4-mini` and `gpt-image-1` (OpenAI) shut down the evening of October 22 (CT), along with Google's `gemini-omni-flash-preview` and Veo 3.1 previews. Google's deep-research-pro preview shuts down October 23. Perplexity retires its routes to older OpenAI models at 7:00 PM CT on October 23.

- Claude Haiku 4.5 is still active. Anthropic still lists retirement not sooner than today, October 15, 2026, with no date set. Keep a tested fallback ready.

New Models and Price Changes

Prices in USD per 1M tokens.

Vendor	Model ID	Status	Input	Cached	Output
TypeSafe	`jev-1.14.0`	New. Available from Oct 14, 2026. Same price as 1.13; output free.	$0.042	—	$0
TypeSafe	`jev-preview`	Now resolves to jev-1.14.0 (was 1.13). Price unchanged.	$0.042	—	$0
TypeSafe	`jev-1.13.0 / jev-latest`	Unchanged. jev-latest still points here.	$0.042	—	$0
Google	`gemini-4-argon`	Unchanged. Stable, paid tier only. Launch price through Nov 30, 2026 (thinking billed as output).	$2.00	$0.10	$10.00
Google	`gemini-4-argon (from Dec 1, 2026)`	Scheduled Standard price printed by Google.	$4.00	$0.20	$20.00
OpenAI	`gpt-4.1`	Shuts down Jan 14, 2027; replacement gpt-6-luna. Price unchanged.	$2.00	$0.50	$8.00
OpenAI	`gpt-4.1-mini`	Shuts down Jan 14, 2027; replacement gpt-6-luna. Price unchanged.	$0.40	$0.10	$1.60
OpenAI	`gpt-4o`	Shuts down Jan 14, 2027; replacement gpt-6-luna. Price unchanged.	$2.50	$1.25	$10.00
OpenAI	`gpt-4o-mini`	Shuts down Jan 14, 2027; replacement gpt-6-luna. Price unchanged.	$0.15	$0.075	$0.60
OpenAI	`gpt-6-luna`	Named replacement. Unchanged (short context).	$0.10	$0.01	$0.50
Meta	`muse-wonder-1.0`	Unchanged since its Oct 13 API opening (preview).	$4.00	$0.20	$20.00
Meta	`muse-spark`	Unchanged.	$1.25	$0.15	$4.25
Anthropic	`claude-opus-5-5`	Unchanged.	$4.00	$0.20	$20.00
IBM	`granite-4-h-small`	Unchanged. Could not fully re-verify today.	$0.0636	Not listed	$0.265

- TypeSafe Jev 1.14: posted as available from October 14, 2026. Input $0.042 per million tokens, output free, the same rate as jev-1.13.0. An optional long-context header raises both context limits to 262k tokens at the same per-token price. System One / classifier use only; not a chat model.
- Gemini 4 Argon: Google describes it as its flagship reasoning model. Context window 1,048,576 tokens (input and output combined); maximum output 65,536 tokens per synchronous request, or up to 1,048,576 with long decode and background execution. Batch and Flex: $1.00 input / $0.05 cached / $5.00 output through Nov 30, then $2.00 / $0.10 / $10.00. Priority: $3.60 / $0.18 / $18.00 through Nov 30, then $7.20 / $0.36 / $36.00. Google's pricing page was last updated October 15, 2026 (UTC). An invite-only Contributor tier at 50% of Standard covers Argon and some Flash models; it is not on the public Standard list.
- OpenAI's October 14 notice covering gpt-4o-2024-08-06, gpt-4o-2024-11-20, gpt-4o-mini-2024-07-18, gpt-4.1-2025-04-14 and gpt-4.1-mini-2025-04-14 is unchanged. Fine-tuned versions follow the same January 14, 2027 shutdown schedule.
- Anthropic's price tables are unchanged. Claude Opus 5.5 cache reads cost 5% of the input price ($0.20 per million tokens).
- Mistral Large 4 stays on sale at $0.68 input and $2.09 output (list $1.36 / $4.18), with no end date. Alibaba's China (Beijing) row for qwen-math-plus-2024-08-16 was removed; the -09-19 row remains. No other chat/token price changes among the vendors above.

- Amazon's Bedrock price pages could not be read today (JavaScript shell only), ByteDance published no usable token price table, Cohere still publishes no full token matrix, Microsoft's Foundry catalog lists no pay-as-you-go prices, and IBM's supported-models docs returned 403. Their last confirmed prices stand.

Deprecations

Vendor	Model ID	Shutdown (CT)	Replacement	Note
OpenAI	`gpt-4.1 (gpt-4.1-2025-04-14)` `gpt-4.1-mini (gpt-4.1-mini-2025-04-14)`	Wed Jan 13, 2027, 6:00 PM CT	`gpt-6-luna`	Unchanged since Oct 14. OpenAI lists Jan 14, 2027 with no time; 6:00 PM CT Jan 13 is 00:00 UTC Jan 14. Each is the only listed snapshot, so the plain names stop too.
OpenAI	`gpt-4o-2024-08-06` `gpt-4o-2024-11-20` `gpt-4o-mini-2024-07-18`	Wed Jan 13, 2027, 6:00 PM CT	`gpt-6-luna`	Unchanged since Oct 14. Same notice and timing. Fine-tuned versions follow the same schedule.
Google	`gemini-3.5-flash`	Not shutting down (Google lists N/A)	`gemini-3.6-flash (Google auto-route)`	Unchanged since Oct 12. Google says the endpoint is not shutting down and routes requests automatically. Not selectable in Google AI Studio.
Google	`gemini-3.7-flash`	Not shutting down (Google lists N/A)	`gemini-3.8-flash (Google auto-route)`	Unchanged since Oct 12. Same wording as gemini-3.5-flash.
Google	`deep-research-pro-preview-12-2025`	Fri Oct 23, 2026 (time not stated)	`deep-research-preview-04-2026`	Unchanged.
OpenAI	`gpt-4.1-nano` `o4-mini` `gpt-image-1`	Thu Oct 22, 2026, 7:00 PM CT	`gpt-5.6-luna` `gpt-5.6-terra` `gpt-image-2.5-sunburst` or `gpt-image-2.5-flare`	Unchanged. OpenAI lists Oct 23 with no time; 7:00 PM CT Oct 22 is 00:00 UTC Oct 23.
Google	`gemini-omni-flash-preview`	Thu Oct 22, 2026 (time not stated)	`gemini-omni-1.1-flash`	Unchanged.
Google	`veo-3.1-generate-preview` `veo-3.1-fast-generate-preview` `veo-3.1-lite-generate-preview`	Thu Oct 22, 2026 (time not stated)	`gemini-omni-1.1-flash`	Unchanged.
Google	`gemini-3.1-flash-image`	No shutdown date announced	`gemini-nano-banana-2.1`	Unchanged.
Google	`gemini-3.1-flash-tts-preview` `gemini-2.5-flash-preview-tts` `gemini-2.5-pro-preview-tts`	Tue Nov 17, 2026 (time not stated)	`gemini-3.8-flash-tts` or `gemini-3.8-flash-lite-tts`	Unchanged.
Google	`gemini-3.1-flash-live-preview` `gemini-2.5-flash-native-audio-preview-12-2025`	Tue Nov 17, 2026 (time not stated)	`gemini-3.8-live`	Unchanged.
Anthropic	`claude-sonnet-4-5-20250929`	Mon Nov 30, 2026 (time not stated)	`claude-sonnet-5-5`	Unchanged.
Anthropic	`claude-haiku-4-5-20251001`	Not sooner than Thu Oct 15, 2026; no date set	`None named`	Watch item. Still active on Anthropic's deprecations page today.

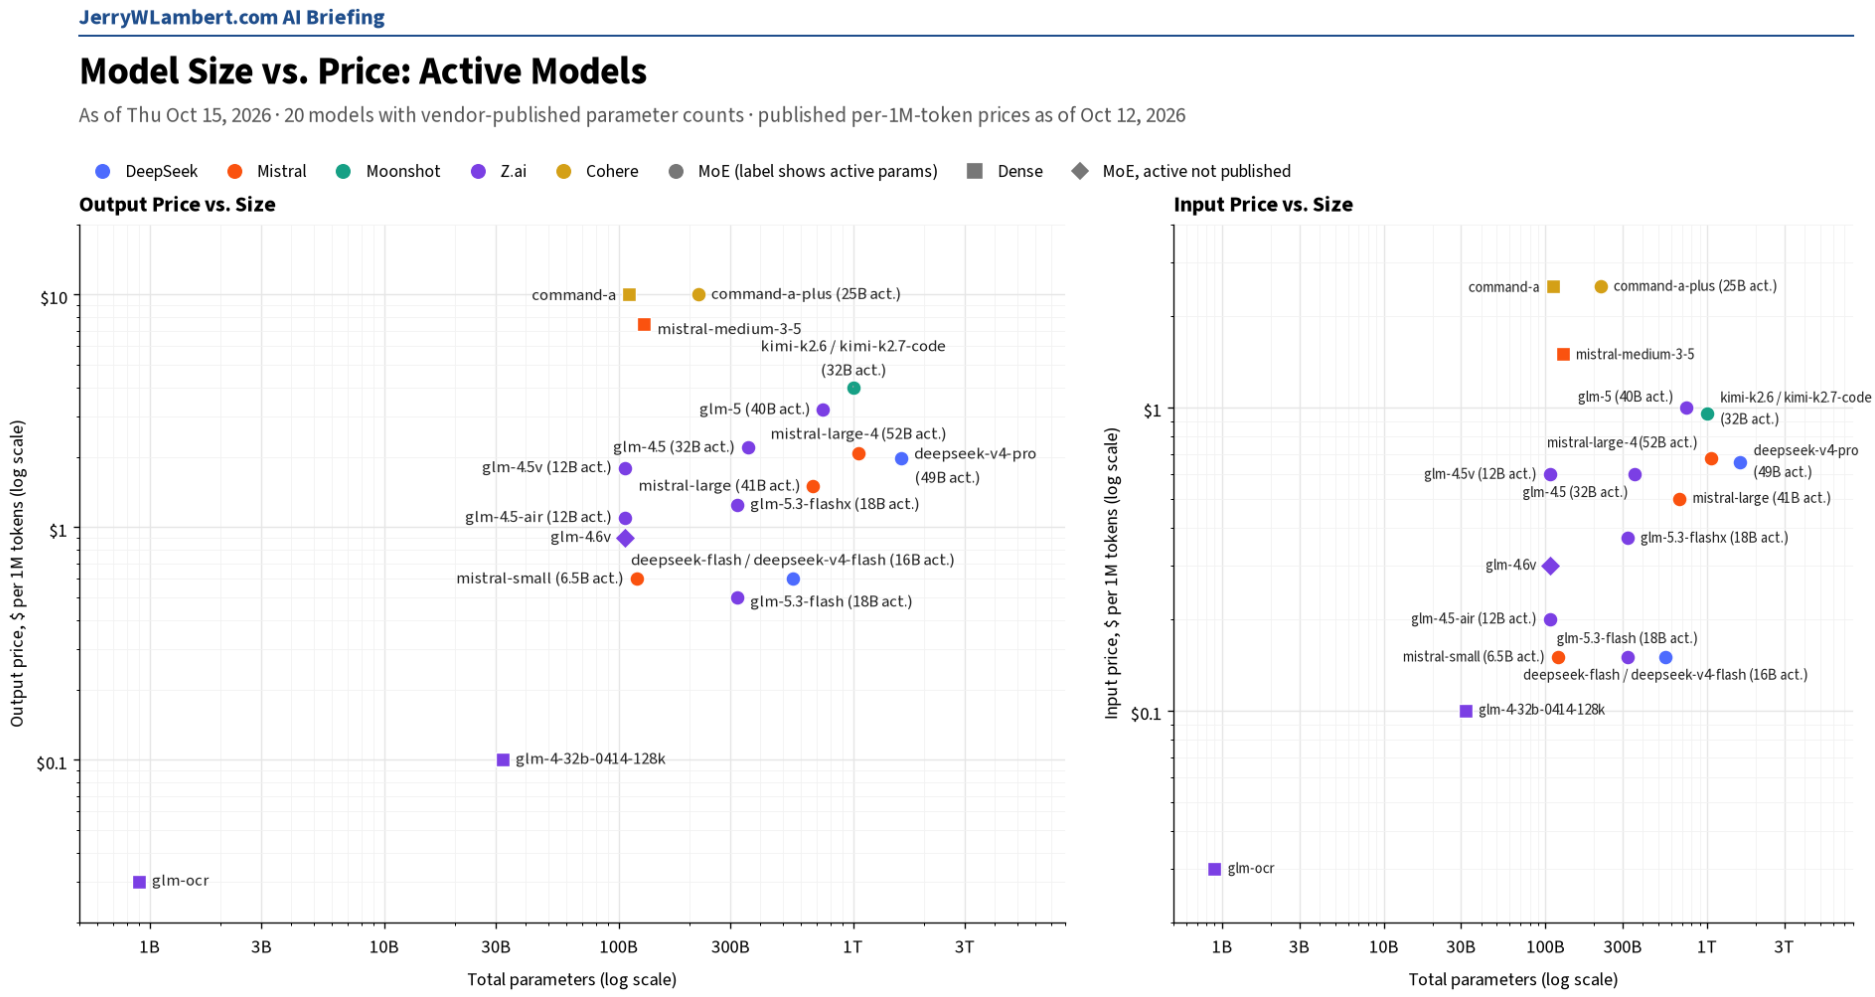
Vendor	Model ID	Shutdown (CT)	Replacement	Note

Cohere	`command-r`
	`command-r-plus`
		Deprecated Sep 15, 2025; no shutdown date set	`command-r-08-2024`
			`command-r-plus-08-2024`
			`command-a-03-2025`
				Closed to new customers; existing users keep access until retirement.

- Perplexity will retire its Agent API and Orchestrator API routes to older OpenAI models (openai/gpt-5.4, gpt-5.4-mini, gpt-5.4-nano, gpt-5.2, gpt-5.1, gpt-5 and gpt-5-mini) at 7:00 PM CT on Fri Oct 23, 2026 (00:00 UTC Oct 24). Unchanged.
- OpenAI's earlier notices are unchanged: gpt-5.3-codex, gpt-5.1 and gpt-5.4-nano shut down April 1, 2027. OpenAI also lists the /v1/audio/translations endpoint shutting down February 26, 2027, replaced by v1/realtime/translations (endpoint only; no model id).
- Google also lists deep-research-preview-04-2026 and deep-research-max-preview-04-2026 with no shutdown date announced. Google's deprecations page shows no shutdown-date changes today.
- Unverified: Google's deprecations page still shows gemini-2.5-flash-image shutting down Mon Mar 15, 2027, with gemini-3.1-flash-lite-image as the replacement, while Google's pricing page says it shut down Oct 2, 2026. Confirm with Google before you budget either date.

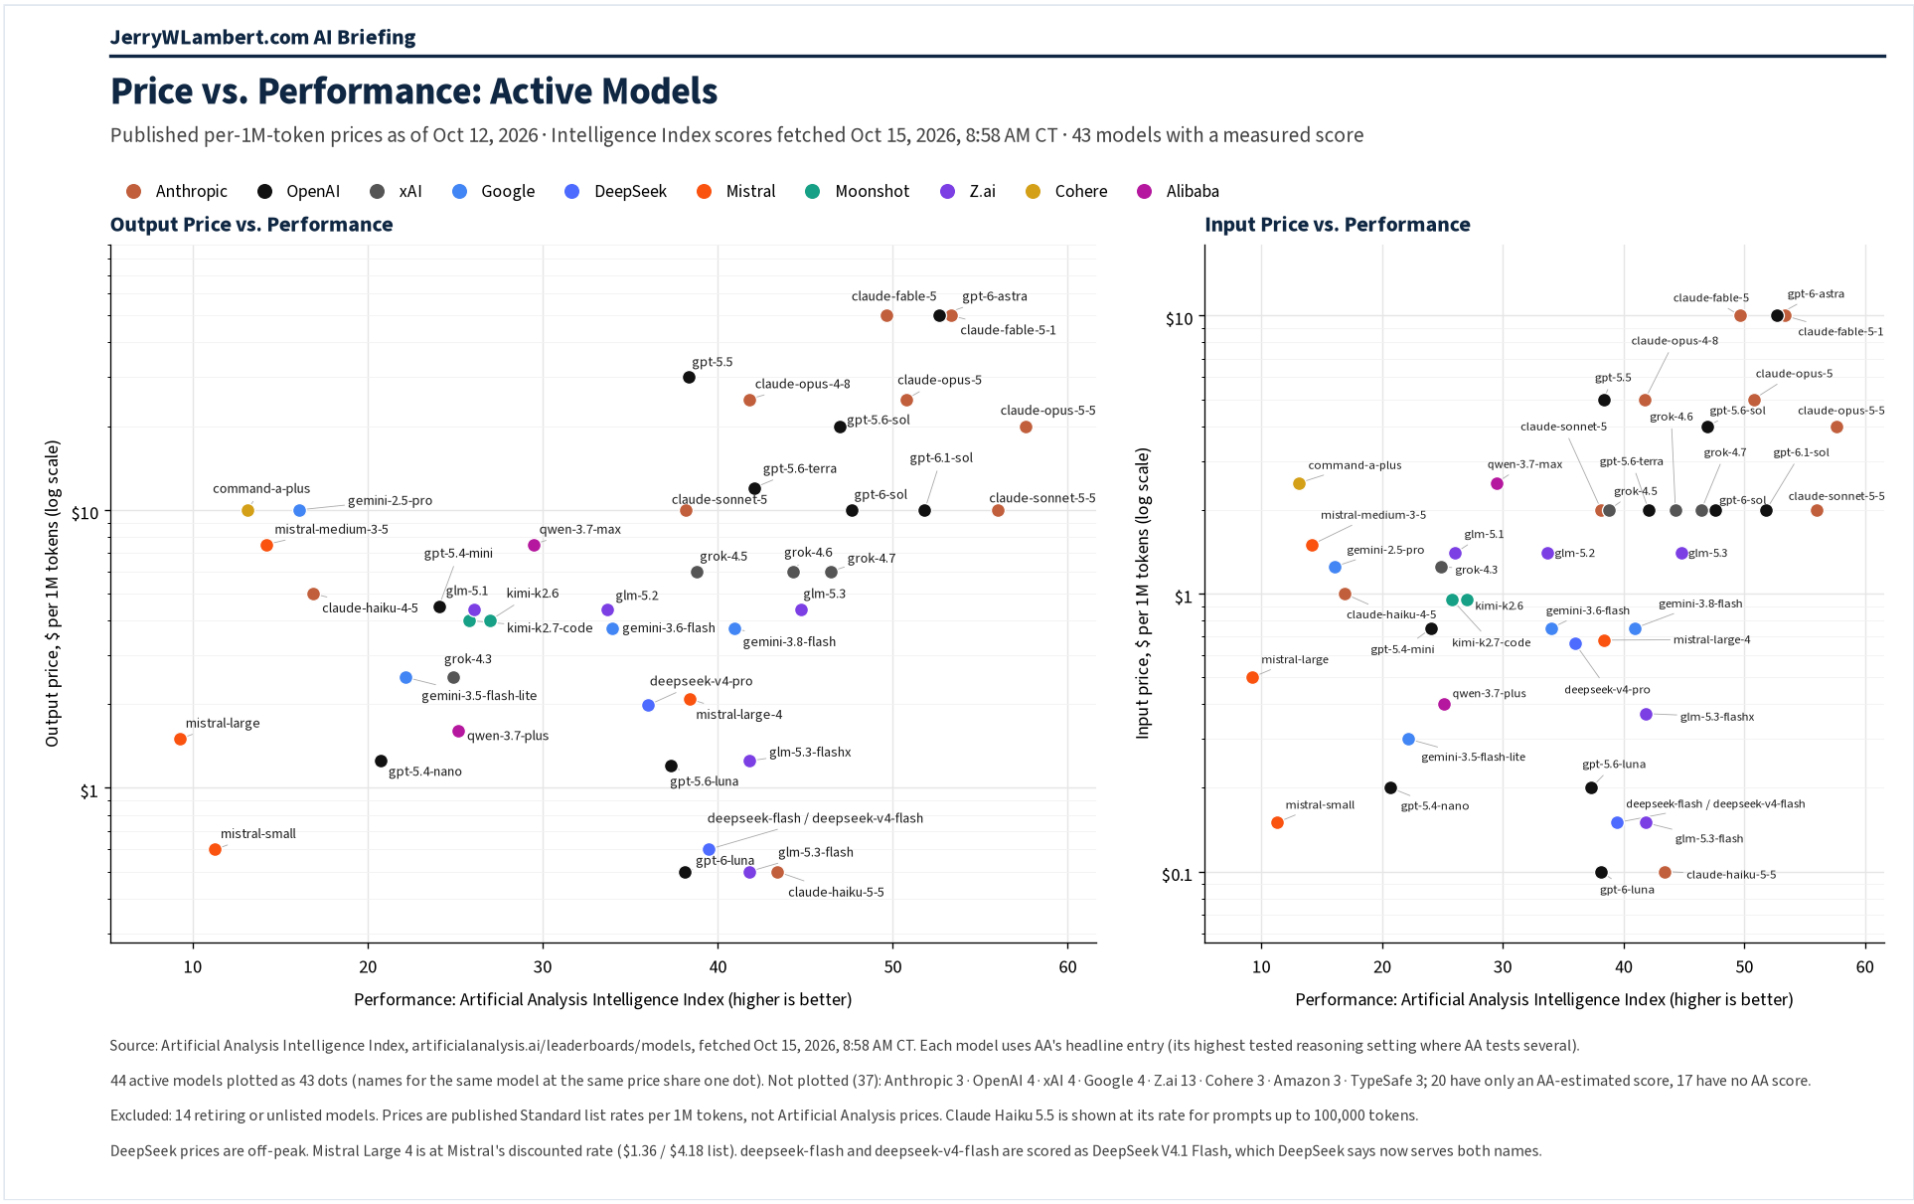
All times Central. Vendor prices, dates and model IDs are taken from first-party vendor pricing, model, deprecation and changelog pages as of 8:26 AM CT on October 15, 2026. Chart prices are as of October 12, 2026. Performance scores are from the Artificial Analysis Intelligence Index, fetched 8:58 AM CT on October 15, 2026. Items that could not be confirmed are marked unverified.
© 2026 Jerry W. Lambert Sr.

Model Size vs. Price



Across the 20 active models whose vendors publish parameter counts, price rises only loosely with total size (rank correlation 0.39 on output price). Anthropic, OpenAI, xAI and Google publish no parameter counts, so their models, including Claude Haiku 5.5 and GPT-6.1 Sol, are not plotted. Google's Gemini 4 Argon and Meta's Muse Wonder are not yet in this chart. Prices as of October 12, 2026.

Price vs. Performance



Across the 44 active models with a measured Artificial Analysis Intelligence Index score (fetched Oct 15, 2026, 8:58 AM CT), price rises with performance only loosely (rank correlation 0.45 on output price). Claude Haiku 5.5 leads at $0.50 or less per million output tokens (43.4); GPT-6.1 Sol scores 51.8 at $10. Google's Gemini 4 Argon (52.6) and Meta's Muse Wonder (52.4) are not yet in this chart. Chart prices as of October 12, 2026.