Daily Model Brief, October 22, 2026

No vendor changed a token price on any model I track today. Alibaba listed Qwen3.8-27B at $0.50 input and $3.00 output per million tokens, and OpenAI's gpt-4.1-nano, o4-mini and gpt-image-1 stop at 7:00 PM CT tonight.

Thursday, October 22, 2026 · Prices as of 8:45 AM CT

Executive Summary

No token price changed today on any model I track. Anthropic, OpenAI, Google, xAI, DeepSeek (peak and off-peak), Moonshot, Z.ai, MiniMax, Mistral, Alibaba, Meta, TypeSafe, Xiaomi, Tencent and StepFun all show the same prices as Wednesday. Mistral Large 4 stays at its $1.36 input, $0.14 cached and $4.18 output list price.

Alibaba's Model Studio page for Qwen3.8-27B (`qwen3.8-27b`) lists international prices of $0.50 input, $3.00 output, $0.10 for implicitly cached input, $0.625 for explicit cache creation and $0.05 for explicit cache reads, per million tokens. Alibaba says those prices exclude limited-time promotions.

Cohere lists token prices only for legacy models, "for existing customers" (Command R 03-2024 at $0.50 input / $1.50 output; Command R+ 08-2024 at $2.50 / $10.00). It shows no token price for Command A, Command A Plus or the current Command R, so I no longer chart Command A or Command A Plus.

Two shutdowns land today. OpenAI's `gpt-4.1-nano`, `o4-mini` and `gpt-image-1` stop at 7:00 PM CT (OpenAI lists October 23; that is midnight UTC). Google lists `gemini-omni-flash-preview` and three Veo 3.1 preview models shutting down today, with no time of day. Microsoft Foundry now dates Llama-4-Maverick-17B-128E-Instruct-FP8 for November 19, 2026 and Llama-3.3-70B-Instruct for January 31, 2027, and moved mistral-small-2503 from December 9, 2026 to February 15, 2027.

Anthropic released API version 2026-10-21 with new SDKs that send it by default (Python SDK 2.0.0, TypeScript SDK 1.0.0 and others), now licensed under Apache 2.0, and organization admins can now turn off web fetch. OpenAI now lets organizations restrict which service tiers each project can use. Z.ai listed GLM-5.3-FlashX-Legal at $0.37 input, $0.075 cached and $1.25 output, with API access only for selected enterprise customers by application. None of these changes a price you pay today.

Prices from Amazon, ByteDance and IBM could not be confirmed today; Microsoft lists no pay-as-you-go prices (IBM's list price is unchanged). The performance chart still uses scores from October 16.

What This Means for Customers

- If you still call `gpt-4.1-nano`, `o4-mini` or `gpt-image-1`, switch before 7:00 PM CT today. OpenAI names `gpt-5.6-luna`, `gpt-5.6-terra`, and `gpt-image-2.5-sunburst` or `gpt-image-2.5-flare` as replacements.
- If you still call `gemini-omni-flash-preview` or a Veo 3.1 preview model, move to `gemini-omni-1.1-flash` now. Google's shutdown date is today and it gives no time of day.
- If you are pricing a mid-size open model, Qwen3.8-27B lists at $0.50 input and $3.00 output on Alibaba's international pricing. Explicit cache reads are $0.05, half the $0.10 implicit-cache rate, so check which cache mode your client uses.
- If you budget Cohere's Command A or Command A Plus, get a written price from Cohere. It lists no public token price for either model.
- If you run Llama-4-Maverick-17B-128E-Instruct-FP8 on Azure Foundry, plan a move before November 19, 2026; Llama-3.3-70B-Instruct follows on January 31, 2027. If you run mistral-small-2503 there, you now have until February 15, 2027.

- If you call Anthropic's API through the new SDKs, they send API version 2026-10-21 by default and carry breaking changes of their own; read Anthropic's migration guide before upgrading.
- Google's deep-research-pro preview ends October 23, and Perplexity retires its routes to older OpenAI models at 7:00 PM CT on October 23.
- If your agents search the web through DeepSeek Harness, Claude Code or Codex, those server-side calls start billing at $0.75 per 1,000 successful Search or Fetch requests at 11:00 AM CT on Saturday, October 31. Track this separately from token spend.
- Claude Haiku 4.5 is still active. Anthropic lists retirement not sooner than October 15, 2026, a date that has passed, and still sets no date. Keep a tested fallback ready.

New Models and Price Changes

Prices in USD per 1M tokens.

Vendor	Model ID	Status	Input	Cached	Output
Alibaba	`qwen3.8-27b`	Newly priced (International). Implicit cache $0.10; explicit cache creation $0.625; explicit cache read $0.05.	$0.50	$0.10	$3.00
Z.ai	`glm-5.3-flashx-legal`	Newly listed. API access for selected enterprise customers by application.	$0.37	$0.075	$1.25
Cohere	`command-a / command-a-plus / command-r`	No public token price from Cohere. Legacy prices only, for existing customers.	Not listed	Not listed	Not listed
Cohere	`Command R 03-2024 (legacy)`	Listed for existing customers.	$0.50	—	$1.50
Cohere	`Command R+ 08-2024 (legacy)`	Listed for existing customers.	$2.50	—	$10.00
Mistral	`mistral-large-4`	Unchanged at list price (sale ended Oct 21).	$1.36	$0.14	$4.18
Mistral	`mistral-large (Mistral Large 3)`	Unchanged.	$0.50	—	$1.50
Alibaba	`qwen3.7-max`	Unchanged. Cached input is 20% of input.	$2.50	$0.50	$7.50
Alibaba	`qwen3.8-max`	Unchanged.	$2.00	—	$6.00
Alibaba	`qwen3.6-35b-a3b`	Unchanged.	$0.375	—	$2.25
Moonshot	`kimi-k3`	Unchanged. 5-minute cache write $3.00; 1-hour cache write $6.00.	$3.00	$0.30	$15.00
DeepSeek	`deepseek-flash`	Unchanged. Off-peak price shown (peak $0.30 / $0.006 / $1.20).	$0.15	$0.003	$0.60
DeepSeek	`deepseek-v4-pro`	Unchanged. Off-peak price shown (peak $1.32 / $0.044 / $3.96).	$0.66	$0.022	$1.98
DeepSeek	`Search API / Fetch API`	Unchanged. Billing for calls through DeepSeek Harness, Claude Code and Codex starts Oct 31, 11:00 AM CT. $0.75 per 1,000 successful requests each.	Per request	—	Per request
Tencent	`Hy3`	Unchanged.	$0.132	$0.033	$0.528
TypeSafe	`jev-1.13.0 / jev-1.14.0`	Unchanged. jev-latest and jev-preview aliases; output free.	$0.042	—	$0
IBM	`granite-4-h-small`	Unchanged. Not confirmed today.	$0.0636	Not listed	$0.265

- No tracked token price changed today. Every price above is current as of this morning.

- Alibaba also lists qwen3.8-2.4t-a95b, qwen3.8-flash and qwen3.8-omni-flash; I have not confirmed their prices and do not show them here.
- Z.ai still marks cached input storage as limited-time free, and Xiaomi still marks cache writes as limited-time free; neither gives an end date.
- Anthropic's October 21 release notes add enforcement_status and enforcement_actions fields to user profiles, so a platform can see whether Anthropic has restricted or banned one of its end customers. No price change.
- Prices from Amazon, ByteDance and IBM could not be confirmed today; Microsoft lists no pay-as-you-go prices. Their last confirmed prices stand.

Deprecations

Vendor	Model ID	Shutdown (CT)	Replacement	Note

OpenAI	`gpt-4.1-nano`
	`o4-mini`
	`gpt-image-1`
		Thu Oct 22, 2026, 7:00 PM CT	`gpt-5.6-luna`
			`gpt-5.6-terra`
			`gpt-image-2.5-sunburst` or `gpt-image-2.5-flare`
				Today. OpenAI lists Oct 23 with no time; 7:00 PM CT Oct 22 is 00:00 UTC Oct 23.

Google	`gemini-omni-flash-preview`	Thu Oct 22, 2026 (time not stated)	`gemini-omni-1.1-flash`	Today.

Google	`veo-3.1-generate-preview`
	`veo-3.1-fast-generate-preview`
	`veo-3.1-lite-generate-preview`
		Thu Oct 22, 2026 (time not stated)	`gemini-omni-1.1-flash`	Today.

Google	`deep-research-pro-preview-12-2025`	Fri Oct 23, 2026 (time not stated)	`deep-research-preview-04-2026`	Unchanged. Tomorrow.
Microsoft (Azure)	`Llama-4-Maverick-17B-128E-Instruct-FP8`	Thu Nov 19, 2026 (time not stated)	`None named`	New date on Microsoft's Foundry retirement page (no date listed Wednesday).
Microsoft (Azure)	`Ministral-3B` `mistral-medium-2505`	Wed Dec 9, 2026 (time not stated)	`ministral-ministral-4b-instruct` `mistral-medium-3-5`	Unchanged. Applies to Azure deployments, not Mistral's own API.
Microsoft (Azure)	`mistral-small-2503`	Mon Feb 15, 2027 (time not stated)	`mistralai-mistral-small-4-119b-2603`	Moved from Dec 9, 2026. Microsoft says it deprecated the model on Oct 9, 2026.
OpenAI	`gpt-4.1 (gpt-4.1-2025-04-14)` `gpt-4.1-mini (gpt-4.1-mini-2025-04-14)`	Wed Jan 13, 2027, 6:00 PM CT	`gpt-6-luna`	Unchanged since Oct 14. OpenAI lists Jan 14, 2027 with no time; 6:00 PM CT Jan 13 is 00:00 UTC Jan 14.

OpenAI	`gpt-4o-2024-08-06`
	`gpt-4o-2024-11-20`
	`gpt-4o-mini-2024-07-18`
		Wed Jan 13, 2027, 6:00 PM CT	`gpt-6-luna`	Unchanged since Oct 14. Fine-tuned versions follow the same schedule.

Microsoft (Azure)	`Phi-4`
	`Phi-4-mini-instruct`
	`Phi-4-mini-reasoning`
	`Phi-4-multimodal-instruct`
		Sun Jan 17, 2027 (time not stated)	`microsoft--phi-4 (Foundry Managed Compute)`	Unchanged since Wednesday. Applies to Azure Foundry deployments.

Microsoft (Azure)	`Llama-3.3-70B-Instruct`	Sun Jan 31, 2027 (time not stated)	`None named`	New date on Microsoft's Foundry retirement page (no date listed Wednesday).

Google	`gemini-3.1-flash-tts-preview`
	`gemini-2.5-flash-preview-tts`
	`gemini-2.5-pro-preview-tts`
		Tue Nov 17, 2026 (time not stated)	`gemini-3.8-flash-tts` or `gemini-3.8-flash-lite-tts`	Unchanged.

Google	`gemini-3.1-flash-live-preview` `gemini-2.5-flash-native-audio-preview-12-2025`	Tue Nov 17, 2026 (time not stated)	`gemini-3.8-live`	Unchanged.

Vendor	Model ID	Shutdown (CT)	Replacement	Note
Anthropic	`claude-sonnet-4-5-20250929`	Mon Nov 30, 2026 (time not stated)	`claude-sonnet-5-5`	Unchanged.
Anthropic	`claude-haiku-4-5-20251001`	Not sooner than Thu Oct 15, 2026; no date set	`None named`	Watch item. The earliest date has passed; still active on Anthropic's deprecations page today with no date set.
Google	`gemini-3.5-flash` `gemini-3.7-flash`	Not shutting down (Google lists N/A)	`gemini-3.6-flash` / `gemini-3.8-flash` (Google auto-route)	Unchanged since Oct 12. Google routes requests automatically.
Cohere	`command-r` `command-r-plus`	Deprecated Sep 15, 2025; no shutdown date set	`command-r-08-2024` `command-r-plus-08-2024` `command-a-03-2025`	Closed to new customers; existing users keep access until retirement.

- Microsoft's Foundry retirement page also lists Cohere-embed-v3-english and Cohere-embed-v3-multilingual as deprecated on November 15, 2026, and still lists claude-opus-5-5-fast (version 2) on Azure retiring Fri Jan 15, 2027.
- Perplexity will retire its Agent API and Orchestrator API routes to older OpenAI models (openai/gpt-5.4, gpt-5.4-mini, gpt-5.4-nano, gpt-5.2, gpt-5.1, gpt-5 and gpt-5-mini) at 7:00 PM CT on Fri Oct 23, 2026 (00:00 UTC Oct 24). Unchanged.
- OpenAI's earlier notices are unchanged: gpt-5.3-codex, gpt-5.1 and gpt-5.4-nano shut down April 1, 2027, and the /v1/audio/translations endpoint shuts down February 26, 2027, replaced by v1/realtime/translations.
- No other shutdown dates moved since Wednesday at Anthropic, Google or OpenAI.

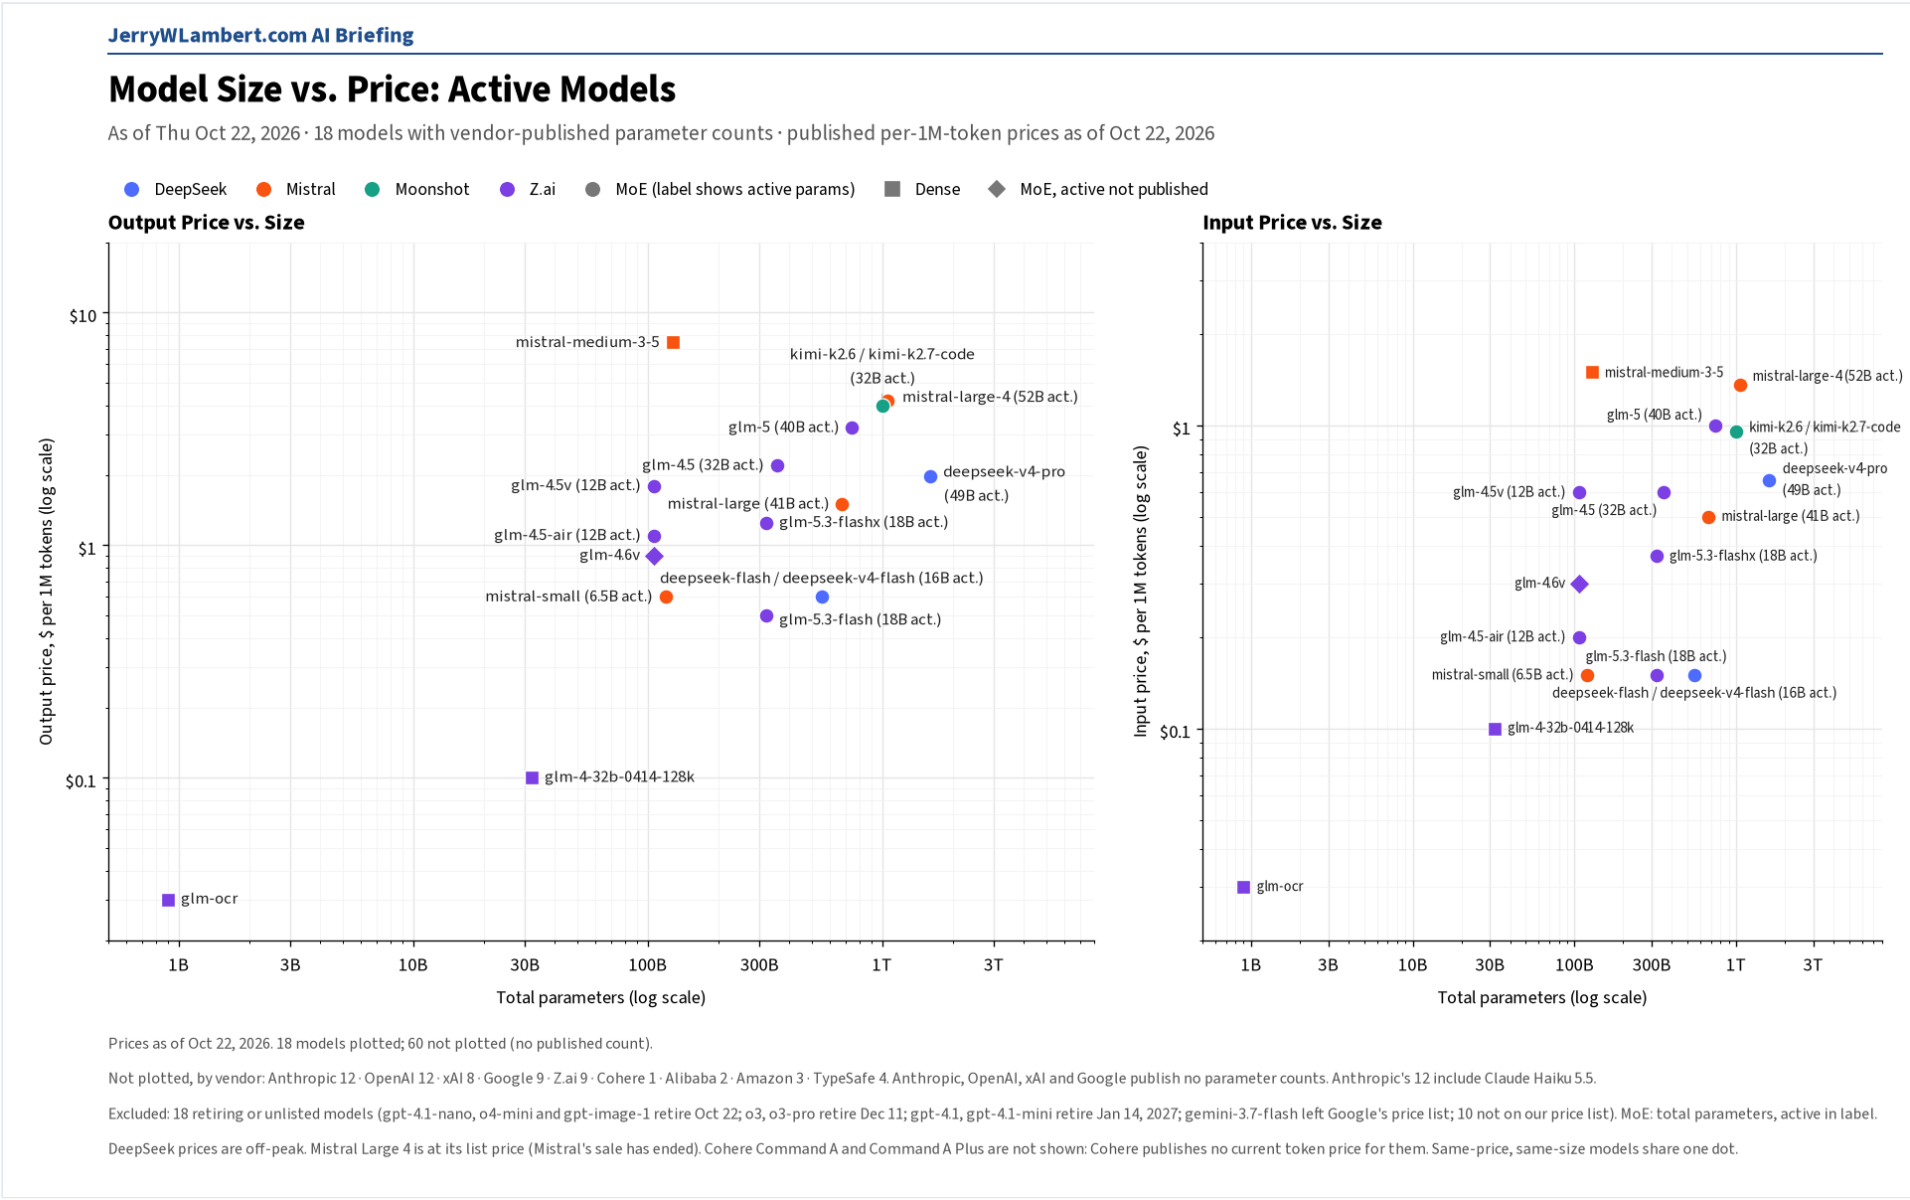
Agent Platforms

Hosted AI agents that come with their own persistent cloud computer.

- **Grok Bot**: its own persistent cloud computer, one per user, with no separate computer charge found; usage comes from a weekly pool, then prepaid top-ups. No change since its October 14 release.
- **dots (OpenAI)**: OpenAI says a dot "has its own cloud computer"; the first dot comes with Pro and Business Premium "at no extra cost." Runs on GPT-6 Astra. No pricing change.
- **Muse (Meta)**: "You and your Muse share your own dedicated computer in the cloud." Runs on Muse Spark 1.3, whose API price is unchanged at $1.25 input, $0.15 cached and $4.25 output per million tokens. Billing is unverified today.
- **Manus**: ordinary tasks get a temporary sandbox; the Cloud Computer add-on is a persistent server billed separately at $30 a month (2 vCPU, 4 GB) or $50 a month (2 vCPU, 8 GB). No change.
- **Genspark Claw**: a "dedicated cloud computer provisioned just for you," billed separately as a Lite, Standard or Powerful Cloud Computer subscription plus credits; Genspark publishes no dollar amounts. Page re-laid out; no substantive change.

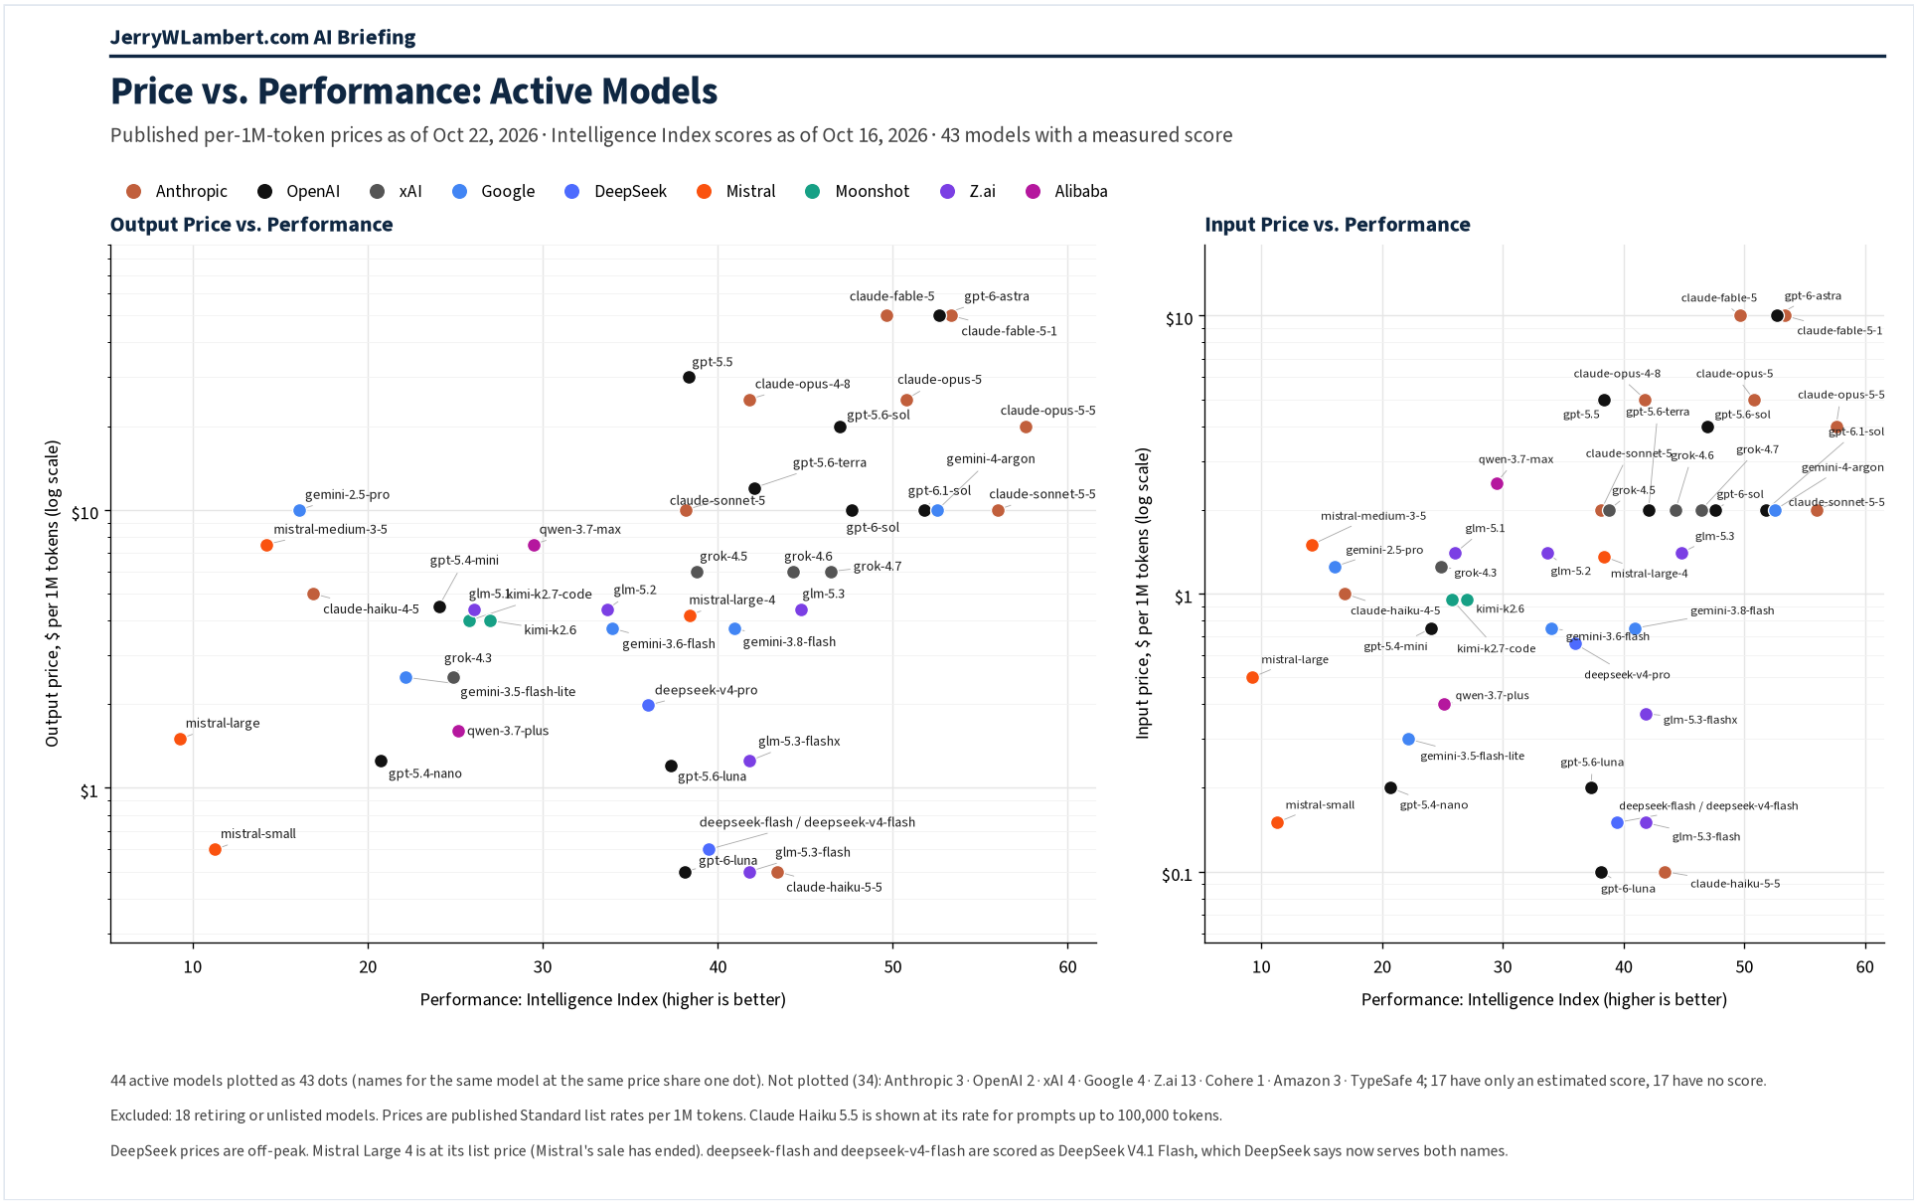
All times Central. Vendor prices, dates and model IDs are as of 8:45 AM CT on October 22, 2026. Chart prices are as of October 22, 2026. Performance scores are as of October 16, 2026. Items that could not be confirmed are marked unverified.
© 2026 Jerry W. Lambert Sr.

Model Size vs. Price



Across the 18 active models whose vendors publish parameter counts, price rises with total size (rank correlation 0.62 on output price). Anthropic, OpenAI, xAI and Google publish no parameter counts, so their models are not plotted. Cohere's Command A and Command A Plus are no longer shown (no current Cohere token price); that removal, not a price change, raised the correlation from 0.41. Prices as of October 22, 2026.

Price vs. Performance



Across the 44 active models with an Intelligence Index score, price rises with performance only loosely (rank correlation 0.49 on output price). Claude Haiku 5.5 leads at $0.50 or less (43.4); GPT-6.1 Sol scores 51.8 and Gemini 4 Argon 52.6, each at $10. Scores as of October 16, 2026; prices as of October 22, 2026.