Daily Model Brief, October 30, 2026

Mistral cut the list price of Mistral Large 4 by about half on Thursday, the only token price change since Thursday's brief. Tomorrow, DeepSeek starts billing Search and Fetch and Moonshot closes K2 to new accounts. Moonshot's Batch API now covers K3 and K3.1 Pro.

Friday, October 30, 2026 · Prices as of 8:47 AM CT · Covers Thursday, October 29 through this morning

Executive Summary

Mistral cut the list price of `mistral-large-4` by about half on Thursday, October 29, after Thursday's brief went out. Input fell from $1.36 to $0.70, cached input from $0.14 to $0.07 and output from $4.18 to $2.10 per million tokens. Mistral shows no sale marker or end date, so this is the new list price.

Both Saturday changes take effect tomorrow, October 31, and I re-checked them this morning. At 11:00 AM CT, DeepSeek starts billing Search and Fetch calls made through DeepSeek Harness, Claude Code and Codex at $0.75 per 1,000 successful requests each, a dollar price DeepSeek confirms today. The same day, Moonshot stops letting newly registered accounts use its K2 models, and it still gives no exact retirement date beyond the end of December.

Moonshot's Batch API now covers `kimi-k3.1-pro` and `kimi-k3`; before, it covered only the K2 models. Batch pricing stays at 60% of standard rates, and batch jobs on these two models also pay cache-write fees.

Google made its Computer Use tool generally available on Friday, October 30. It carries no separate fee: it is charged as regular tokens at the model's price.

No other vendor I track changed a token price or launched a new model since Thursday's brief: Anthropic, OpenAI, Google, xAI (Grok), DeepSeek, Moonshot, Z.ai, MiniMax, Alibaba, Cohere, Meta, TypeSafe, Xiaomi, StepFun and Perplexity. `gpt-6.1-terra` still costs $1.00 input, $0.05 cached and $5.00 output per million tokens, half of `gpt-6.1-sol`.

Not confirmed today: Amazon and ByteDance prices, Tencent prices (last confirmed Thursday afternoon, October 29), Moonshot's per-search web search prices, IBM's model list, and Microsoft, which lists no pay-as-you-go prices. Their last confirmed prices stand. Performance scores in the chart are still as of October 16 and have not been refreshed since.

What This Means for Customers

- If you use `mistral-large-4`, your token cost roughly halves from Thursday, October 29, with nothing to change on your side. If you ruled it out on price, re-run the comparison at $0.70 input and $2.10 output per million tokens.
- If your agents search the web through DeepSeek Harness, Claude Code or Codex, DeepSeek starts billing those Search and Fetch calls at 11:00 AM CT tomorrow, Saturday, October 31, at $0.75 per 1,000 successful requests each. Set a budget line for it today, separate from token spend.
- If you run `kimi-k2.6` or `kimi-k2.7-code`, any Moonshot account registered from tomorrow cannot use K2, so new projects on new accounts should start on `kimi-k3.1-pro`. Existing accounts keep access until retirement, which Moonshot still dates only as the end of December.
- If you run large offline jobs on `kimi-k3.1-pro` or `kimi-k3`, the Batch API now bills them at 60% of standard rates. Budget for the cache-write fees batch jobs on these models still pay, and expect cache hits only on a best-effort basis.
- If your app grounds Gemini answers with Google Search, check your daily peak against the 1,500-request daily cap Google lists, and spread large batch jobs across days if you are near it. The cost per search is unchanged.

- If you use `gpt-6.1-sol`, test `gpt-6.1-terra` on the same work: every token price is half. Check output quality first, and check prompt sizes, because a request over 272,000 input tokens bills entirely at the higher long-context rate.
- If you still run `gpt-4.1` or `gpt-4.1-mini`, shutdown is January 14, 2027, and the replacement is `gpt-6-luna`.
- Claude Haiku 4.5 is still active. Anthropic lists its retirement as not sooner than October 15, 2026, a date that has passed, and still sets no date. Keep a tested fallback ready.

New Models and Price Changes

Prices in USD per 1M tokens.

Vendor	Model ID	Status	Input	Cached	Output
Mistral	`mistral-large-4`	Price cut Thu Oct 29 (was $1.36 / $0.14 / $4.18). No sale marker or end date shown. No cache-write price published.	$0.70	$0.07	$2.10
Moonshot	`Batch API (kimi-k3.1-pro, kimi-k3)`	New for these two models (before, K2 models only). Billed at 60% of standard rates, plus cache-write fees.	60% of standard	Not stated	60% of standard
Google	`Computer Use tool`	Generally available Fri Oct 30. No separate fee; charged as regular tokens at the model's price.	Model price	Model price	Model price
DeepSeek	`Search API / Fetch API`	Billing for calls through DeepSeek Harness, Claude Code and Codex starts tomorrow, Sat Oct 31, 11:00 AM CT, at $0.75 per 1,000 successful requests each. Dollar price confirmed today.	Per request	Not applicable	Per request
DeepSeek	`deepseek-flash (also serves deepseek-v4-flash)`	Off-peak price, confirmed today in dollars; unchanged. Peak-hour prices are double.	$0.15	$0.003	$0.60
DeepSeek	`deepseek-v4-pro`	Off-peak price, confirmed today in dollars; unchanged. Peak-hour prices are double.	$0.66	$0.022	$1.98
OpenAI	`gpt-6.1-terra`	Unchanged (launched Mon Oct 26). Context window: 1,050,000 tokens. Cache writes: $1.25 (30-minute) or $2.50 (12-hour). Prompts over 272K input tokens: 2x input and cache, 1.5x output, for the whole request.	$1.00	$0.05	$5.00
OpenAI	`gpt-6.1-sol`	Unchanged. Terra costs half on every line.	$2.00	$0.10	$10.00
OpenAI	`gpt-5.6-terra`	Unchanged. Older model; not the same as gpt-6.1-terra.	$2.00	$0.20	$12.00
Google	`Grounding with Google Search (Gemini 3 and newer)`	Unchanged. Daily cap of 1,500 requests; 5,000 free searches a month, then $14 per 1,000.	Per request	Not applicable	Per request
Moonshot	`Web search (Responses API and Kimi K3 built-in tool)`	Terms unchanged. Billed per search the server runs, plus model tokens. Per-call prices still did not display today; not confirmed.	Per call	Not applicable	Per call
Moonshot	`kimi-k3.1-pro`	Unchanged (standard rate). Cache writes: $2.00 (5-minute) or $4.00 (1-hour).	$2.00	$0.20	$10.00

Vendor	Model ID	Status	Input	Cached	Output
Moonshot	`kimi-k3.1-pro-highspeed`	Unchanged. Cache writes: $4.00 (5-minute) or $8.00 (1-hour).	$4.00	$0.20	$20.00
Moonshot	`kimi-k3`	Unchanged (standard rate).	$3.00	$0.30	$15.00
Moonshot	`kimi-k2.7-code`	Unchanged. Closed to new accounts from tomorrow, Oct 31; retiring by end of December.	$0.95	$0.19	$4.00
Moonshot	`kimi-k2.6`	Unchanged. Closed to new accounts from tomorrow, Oct 31; retiring by end of December.	$0.95	$0.16	$4.00
Google	`gemini-3.8-flash (also serves gemini-3.7-flash)`	Unchanged through Dec 31; $1.50 / $0.15 / $7.50 from Jan 1, 2027.	$0.75	$0.075	$3.75
Alibaba	`qwen3.8-27b`	Unchanged (International).	$0.50	Not listed	$3.00
Alibaba	`qwen3.8-max`	Unchanged.	$2.00	Not listed	$6.00
MiniMax	`MiniMax-M3.1-Flash`	Unchanged.	$0.15	$0.015	$0.60
Meta	`muse-spark-1.3`	Unchanged (Standard mode).	$1.25	$0.15	$4.25
Cohere	`Command R+ 08-2024 (legacy)`	Unchanged. No cached-token price published.	$2.50	Not listed	$10.00
Tencent	`Hy3`	Last confirmed Thursday afternoon (Oct 29). Not confirmed today.	$0.132	$0.033	$0.528
TypeSafe	`jev-1.14.0`	Unchanged. jev-latest points here. Output free.	$0.042	Not listed	Free
IBM	`granite-4-h-small`	Last confirmed price. Not confirmed today.	$0.0636	Not listed	$0.265

- Mistral's cut is the only public token price change since Thursday's brief at any vendor I could confirm, and no new model launched.
- Mistral Large 4's new price is about 49% lower on input and 50% lower on cached input and output.
- DeepSeek's dollar prices were confirmed again today and match the last confirmed figures. The table shows off-peak rates; peak-hour rates are double.
- Anthropic now says Enterprise and API customers can get volume and committed-spend discounts through its sales team. It publishes no discount rates, and its list prices did not change.
- Terra's long-context rule applies to the whole request: a 300,000-token prompt bills every input token at $2.00 per million, not just the tokens past 272,000.
- Prices from Amazon, ByteDance and Tencent could not be confirmed today, Moonshot's per-search prices did not display, IBM's model list could not be confirmed, and Microsoft lists no pay-as-you-go prices. Their last confirmed prices stand, unverified today.

Deprecations

Vendor	Model ID	Shutdown (CT)	Replacement	Note
Google	`gemini-3.7-flash`	No shutdown	`gemini-3.8-flash`	Google routes API requests to gemini-3.8-flash automatically, at the same price.

Vendor	Model ID	Shutdown (CT)	Replacement	Note

Moonshot	`kimi-k2.6`
	`kimi-k2.7-code`
	`(all K2 models)`
		By end of December 2026 (no exact date)	`kimi-k3.1-pro`	Unchanged. Closed to newly registered accounts from tomorrow, Sat Oct 31; existing accounts keep access until retirement.

Google	`gemini-3.1-flash-tts-preview`
	`gemini-2.5-flash-preview-tts`
	`gemini-2.5-pro-preview-tts`
		Tue Nov 17, 2026 (time not stated)	`gemini-3.8-flash-tts` or `gemini-3.8-flash-lite-tts`	Unchanged.

Google	`gemini-3.1-flash-live-preview` `gemini-2.5-flash-native-audio-preview-12-2025`	Tue Nov 17, 2026 (time not stated)	`gemini-3.8-live`	Unchanged.
Microsoft (Azure)	`Llama-4-Maverick-17B-128E-Instruct-FP8`	Thu Nov 19, 2026 (time not stated)	`None named`	Unchanged.
Anthropic	`claude-sonnet-4-5-20250929`	Mon Nov 30, 2026 (time not stated)	`claude-sonnet-5-5`	Unchanged.
Microsoft (Azure)	`Ministral-3B` `mistral-medium-2505`	Wed Dec 9, 2026 (time not stated)	`ministral-ministral-4b-instruct` `mistral-medium-3-5`	Unchanged. Applies to Azure deployments, not Mistral's own API.

OpenAI	`gpt-5`
	`gpt-5-mini`
	`gpt-5-nano`
	`gpt-5-pro`
	`o3`
	`o3-pro (dated snapshots)`
		Fri Dec 11, 2026 (time not stated)	`gpt-5.6-sol`
			`gpt-5.6-terra`
			`gpt-5.6-luna`
				Unchanged (announced June 11, 2026). gpt-5 and o3 move to gpt-5.6-sol, gpt-5-mini to gpt-5.6-terra, gpt-5-nano to gpt-5.6-luna, and gpt-5-pro and o3-pro to gpt-5.6-sol in pro reasoning mode.

OpenAI	`gpt-4.1`
	`gpt-4.1-mini`
	`gpt-4o-2024-08-06`
	`gpt-4o-2024-11-20`
	`gpt-4o-mini-2024-07-18`
		Wed Jan 13, 2027, 6:00 PM CT	`gpt-6-luna`	Unchanged. OpenAI lists Jan 14, 2027 with no time; 6:00 PM CT Jan 13 is 00:00 UTC Jan 14.

Microsoft (Azure)	`Llama-3.3-70B-Instruct`	Sun Jan 31, 2027 (time not stated)	`None named`	Unchanged.
Microsoft (Azure)	`cohere-command-a` `mistral-small-2503`	Mon Feb 15, 2027 (time not stated)	`Cohere-command-a-plus-05-2026` `mistralai-mistral-small-4-119b-2603`	Unchanged.
OpenAI	`gpt-5.4-nano` `gpt-5.4-nano-2026-03-17`	Thu Apr 1, 2027 (time not stated)	`gpt-6-luna`	Unchanged. Also retiring that day: gpt-5.3-codex and gpt-5.1, with gpt-6-sol as their replacement.
Microsoft (Azure)	`Kimi-K2.6` `Kimi-K2.7-Code`	Fri Apr 16, 2027 (time not stated)	`None named`	Unchanged. Azure's date is later than Moonshot's own K2 retirement.
Microsoft (Azure)	`embed-v-4-0`	Tue Jun 1, 2027 (time not stated)	`Cohere-embed-v5-fast`, `Cohere-embed-v5-pro`	Unchanged.
Microsoft (Azure)	`Cohere-rerank-v4.0-fast` `Cohere-rerank-v4.0-pro`	Mon Sep 13, 2027 (time not stated)	`Cohere-rerank-v5-fast` `Cohere-rerank-v5-pro`	Unchanged.
Microsoft (Azure)	`gpt-oss-120b`	Wed Sep 15, 2027 (time not stated)	`None named`	Unchanged.
Microsoft (Azure)	`DeepSeek-V3.2` `DeepSeek-V3.2-Speciale`	Mon Nov 1, 2027 (time not stated)	`DeepSeek-V4.1-Flash` `DeepSeek-V4-Pro`	Unchanged.

Vendor	Model ID	Shutdown (CT)	Replacement	Note
Anthropic	`claude-haiku-4-5-20251001`	Not sooner than Thu Oct 15, 2026; no date set	`None named`	Watch item. Still active with no date set.

Cohere	`command-r`
	`command-r-plus`
		Deprecated Sep 15, 2025; no shutdown date set	`command-r-08-2024`
			`command-r-plus-08-2024`
			`command-a-03-2025`
				Closed to new customers; existing users keep access until retirement.

- No new retirement and no moved shutdown date since Thursday's brief at any vendor I track.
- OpenAI's transcription notice is unchanged: whisper-1, gpt-4o-transcribe, gpt-4o-mini-transcribe, gpt-4o-transcribe-diarize and the /v1/audio/translations endpoint shut down February 26, 2027.
- Moonshot still gives a month, not a day, for the K2 retirement, so treat December 1 as the safe planning date.

Price Tracker

Shares of the full tracked population.

- Tracked models: 1,934 (count as of October 25).
- Audited pricing, confirmed against the vendor's own published prices: 13.1% of the full tracked population.
- Unaudited pricing: 5.3%.
- The rest have no public per-token price (47.6%: self-hosted open weights, retired or unpublished) or are duplicate listings of another tracked model (34.0%).

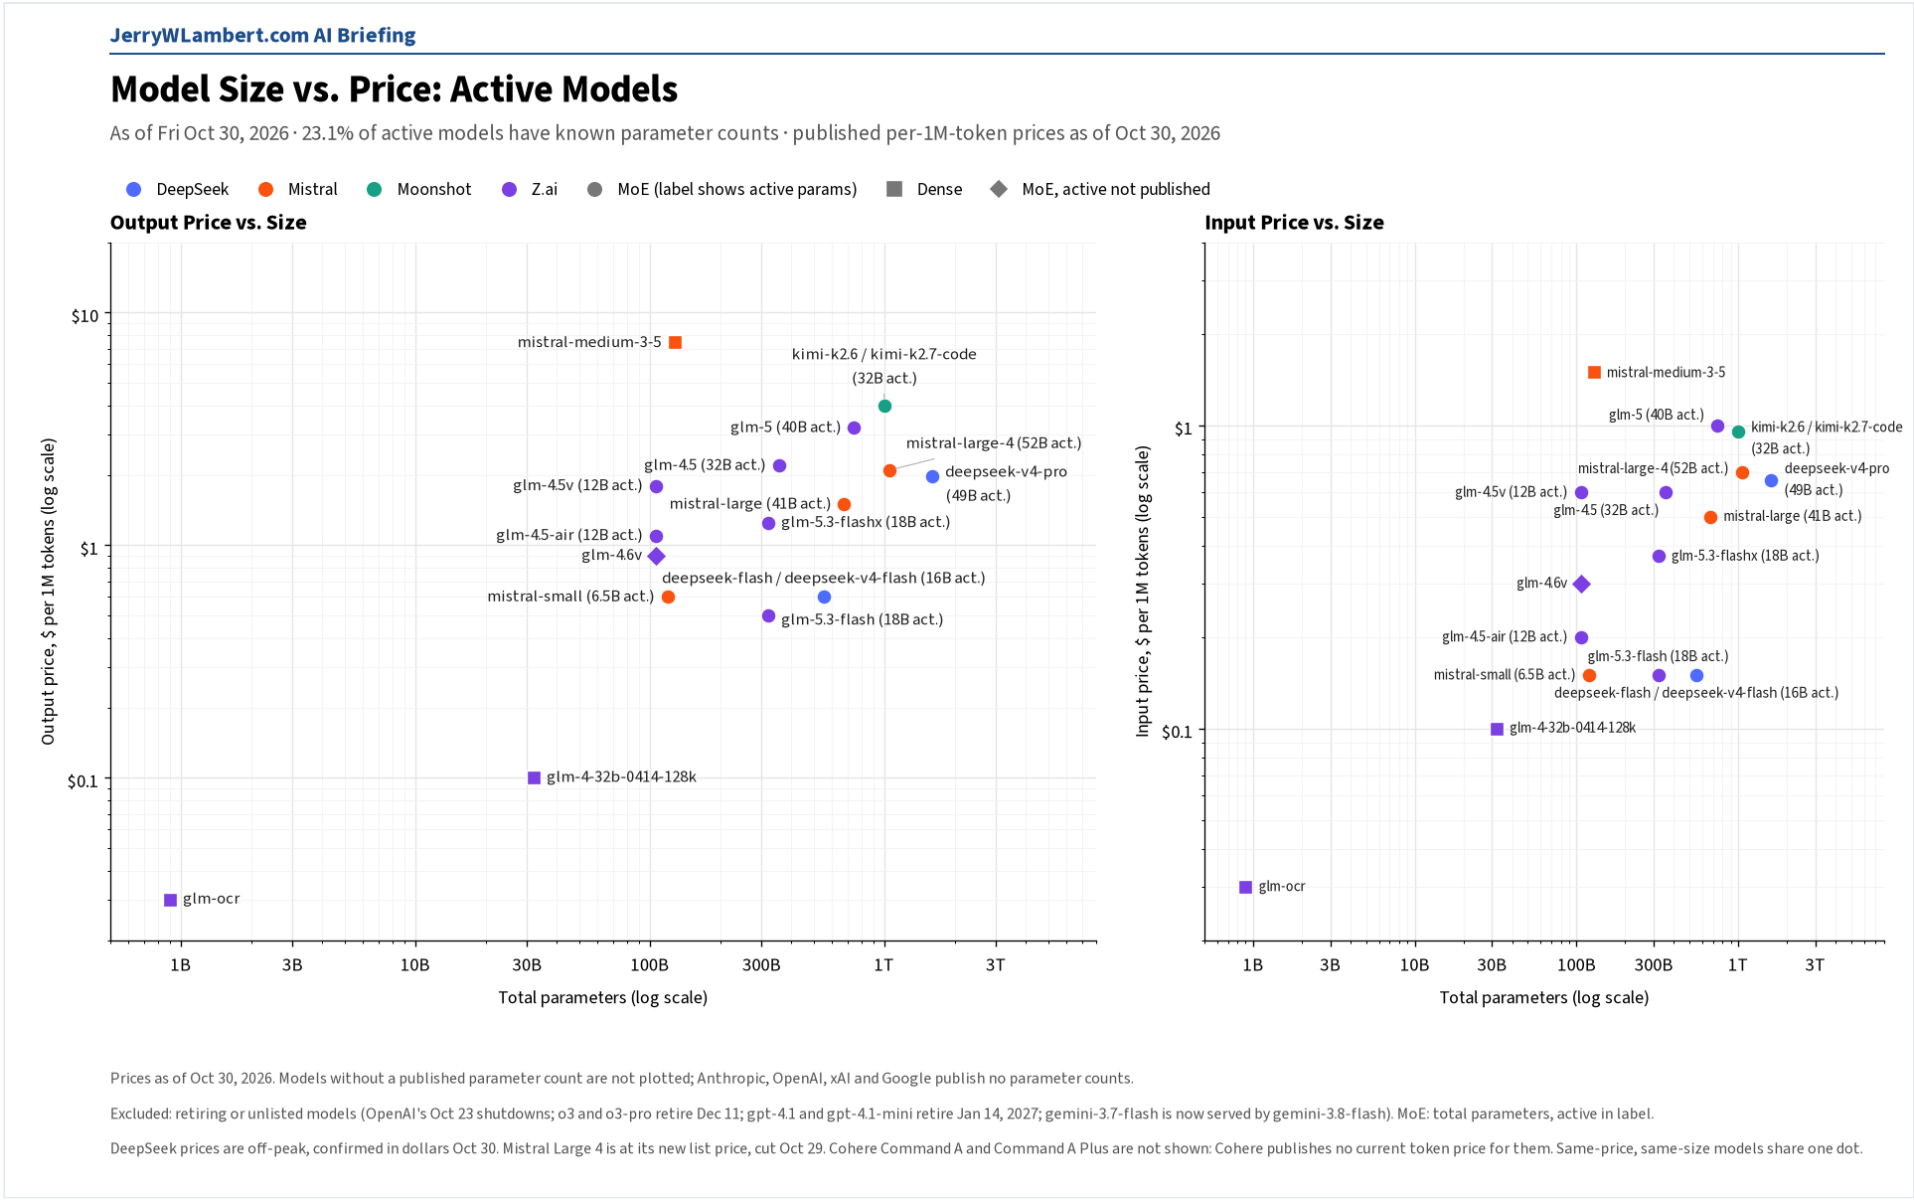
Agent Platforms

Hosted AI agents that come with their own persistent cloud computer.

- **Grok Bot (xAI)**: its own persistent cloud computer. On the X Subscription it is included on Premium, Super and Ultra, draws from the same weekly usage pool as Cursor and the Grok apps, and prepaid credit covers usage past the pool. No change since Thursday: the latest release is still v0.78.1 (dated October 27), which added shopping product cards with an approval step before any order, and plan prices are unchanged.
- **dots (OpenAI)**: not verified today. Last confirmed: a dot "has its own cloud computer," and the first comes with Pro and Business Premium at no extra cost.
- **Muse (Meta)**: "You and your Muse share your own dedicated computer in the cloud." Runs on Muse Spark 1.3, whose API price is unchanged at $1.25 input, $0.15 cached and $4.25 output per million tokens. Billing is unverified today.
- **Manus**: the Cloud Computer add-on is a persistent server billed separately at $30 a month (2 vCPUs, 4 GB RAM, 70 GB storage) or $50 a month (2 vCPUs, 8 GB RAM, 120 GB storage). No change.
- **Genspark Claw**: runs "on a dedicated cloud computer provisioned just for you," which requires a Lite, Standard or Powerful Cloud Computer subscription; Genspark publishes no dollar amounts. Help-page layout changes only; no price change.

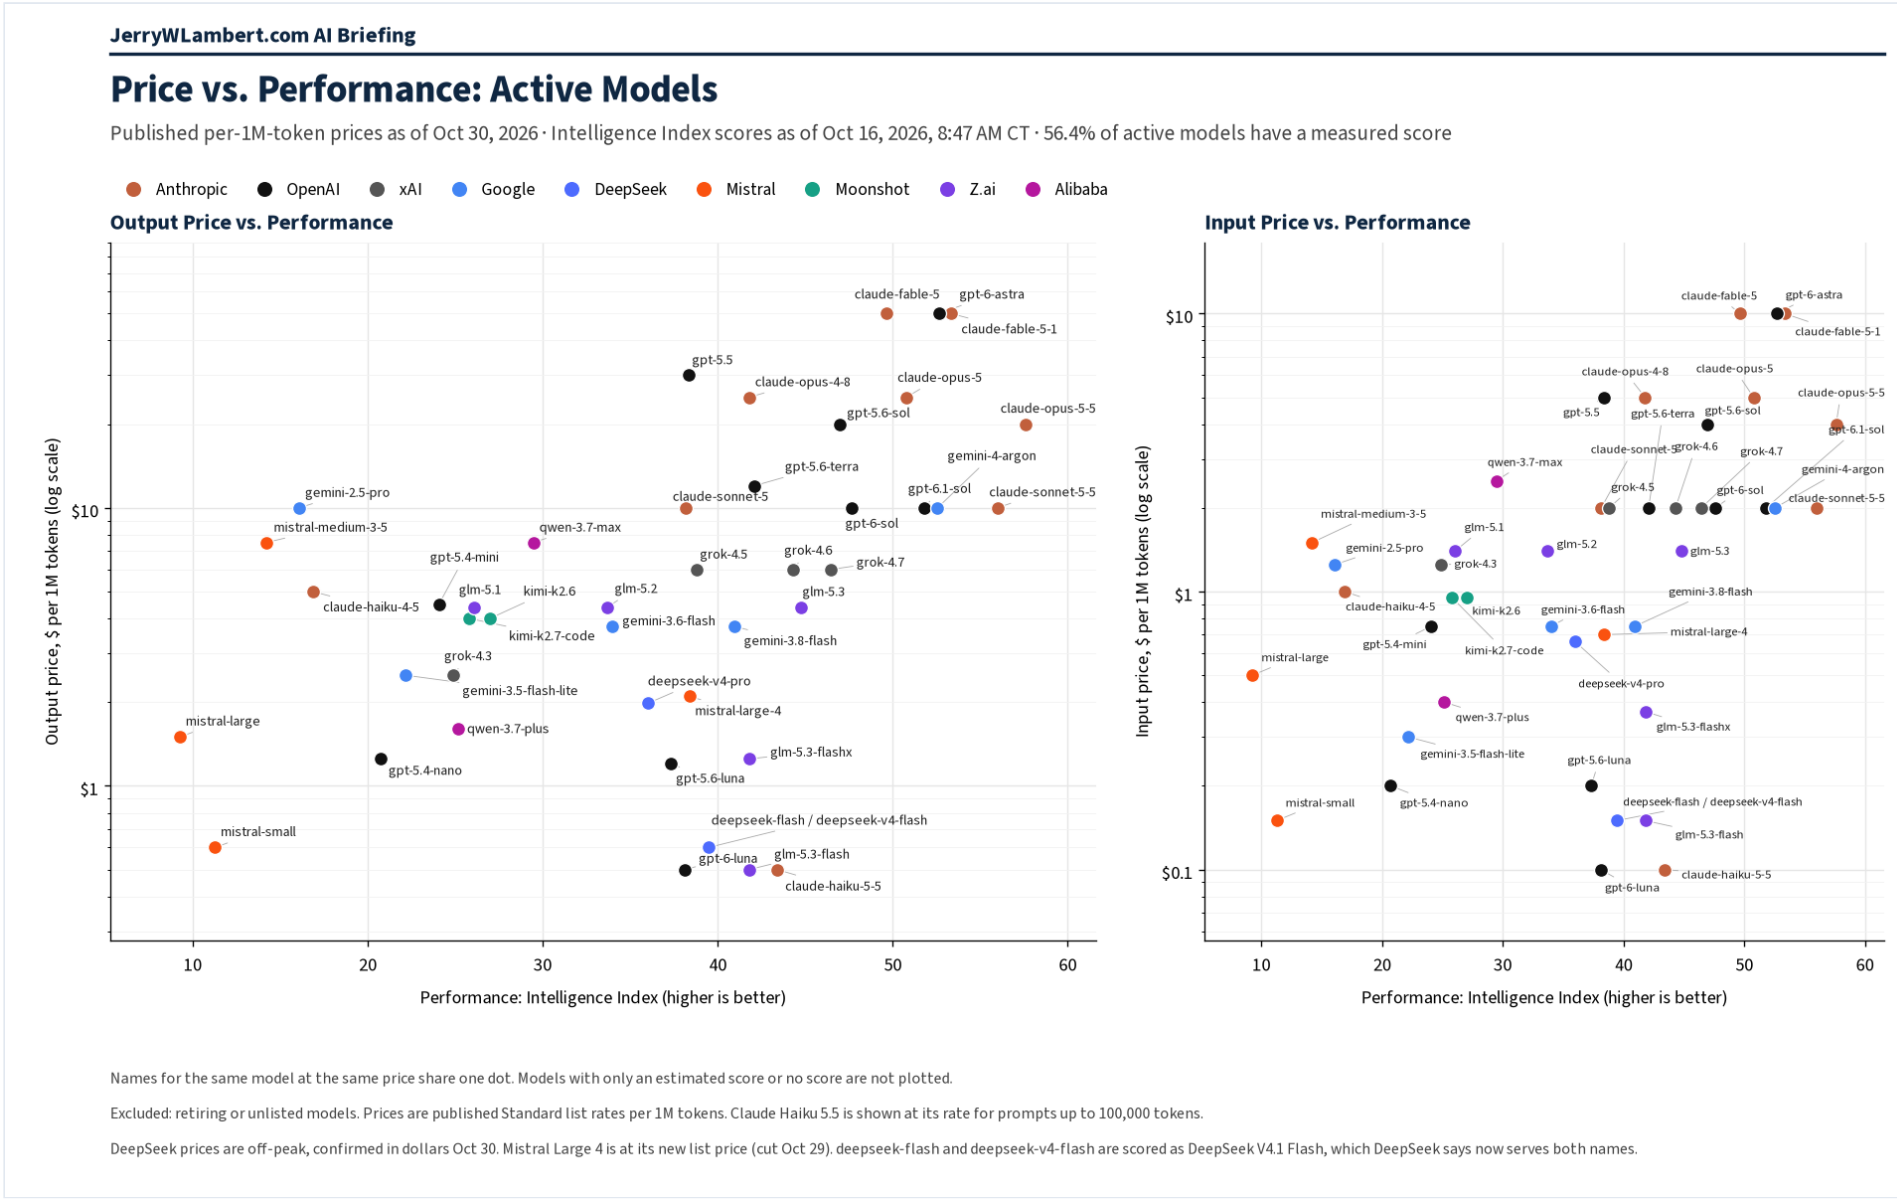
All times Central. Vendor prices, dates and model IDs are as of 8:47 AM CT on October 30, 2026. Chart prices are as of October 30, 2026. Performance scores are as of October 16, 2026 and have not been refreshed since. Price Tracker percentages are shares of the full tracked population; chart percentages are shares of active models. Items that could not be confirmed are marked unverified.
© 2026 Jerry W. Lambert Sr.

Model Size vs. Price



Among active models with a known parameter count (23.1% of active models), price rises with total size (rank correlation 0.59 on output price). Mistral Large 4 moved down to its new list price; no other plotted price changed since Thursday. Anthropic, OpenAI, xAI and Google publish no parameter counts, so their models are not plotted. Prices as of October 30, 2026.

Price vs. Performance



Among active models with a performance score (56.4% of active models), price rises with performance only loosely (rank correlation 0.49 on output price). Claude Haiku 5.5 leads at $0.50 or less (43.4). At the frontier, Gemini 4 Argon scores 52.6 and GPT-6.1 Sol 51.8, each at $10 output, while xAI's Grok 4.7 scores 46.5 at $6. Scores as of October 16, 2026 (not refreshed since); prices as of October 30, 2026.