WHITE PAPER · VERSION 1.0

How Do You Know Your AI Bill Is Correct?

The 35 AI Charging Methods

> Every way I can document that AI vendors turn usage into a charge, numbered 1 to 35 and ranked by how often each one shows up across my full tracked population of 1,934 models, with what to watch, how to optimize and what to avoid for each one.

35	1,934	12.8%
MECHANISMS	MODELS TRACKED	NOW RETIRED

Jerry W. Lambert, Sr.

October 24, 2026 · jerrywlambert.com

© 2026 jerrywlambert.com

Executive Summary

AI pricing keeps getting more complicated. I now track 1,934 models and count 35 distinct ways a vendor turns usage into a charge. They range from the per-token meter to HDR surcharges, download fees, capacity units, router models with no price of their own, and license terms that flip from free to paid at a revenue line.

Vendors are also cutting their old models loose. 12.8% of my full tracked population is now retired. Another 22.2% are open-weights models you host yourself, with no vendor price at all, and 10.8% have no public price. A retired model takes its price with it.

The per-token meter still dominates. 82.7% of priced models bill input and output tokens, and the priced vendors that skip it are all media vendors. Half of priced models (50.4%) discount cached input, and 63.6% of priced vendors do. Reasoning billed at the output price reaches 42.5% of priced models. Media units rank third by vendor share. 54.5% of priced vendors bill by the image, the second, the minute or the character, and 35.0% of priced models carry one of those units.

The twists are where budgets break. Server tool fees cover 34.3% of priced models. Web search costs $10 per 1,000 calls at Anthropic and OpenAI, and Google charges $14 per 1,000 after 5,000 free requests a month. Batch discounts cover 33.1%. Regional price lists cover 29.5%, and prompt-size bands that reprice a whole request cover 23.2%. Speed tiers run the other way. OpenAI Fast is 2x standard and OpenAI Ultrafast reaches 6x.

The new mechanisms add meters that a token-only bill misses. Luma prices HDR video at 2x SDR and HDR plus EXR at 3x. xAI charges $0.20 per GiB to download stored files. Luma sells Uni-1.1 throughput in capacity units with a minimum of four. Liquid AI lets businesses under $10M in annual revenue run its open-weights models in production at no cost. xAI sells Grok 4.7 Fast only through Cursor and Grok Build. Alibaba will not combine batch and cache discounts, and xAI applies the cache discount before the priority multiplier.

My operating rule is simple. Every mechanism needs a meter, a dated price and a stacking rule. If a billing system cannot meter a mechanism, it should refuse the request that triggers it instead of billing it at the wrong price. Bill a 1-hour cache write on Moonshot Kimi K3.1 Pro at the 5-minute price and you collect $2.00 of the $4.00 owed on that line. Let a web search through unmetered and you lose $10 for every 1,000 calls.

> **HOW I COUNTED**
>
> I track 1,934 models. For every model with a published price, I recorded each mechanism that the vendor's public announcements show for that model.

© 2026 jerrywlambert.com

Every percentage is a share of the full tracked population's priced models, or of the vendors behind them. I rank by the model share, most common first, and break ties on the vendor share. Media-only vendors cannot show token mechanisms such as cache discounts or reasoning tokens, which holds those shares down.

All 35 are numbered in one list. Numbers 1 to 27 are ranked by how often they show up on a model's own price row. Numbers 28 to 35 are plan and contract-level charges. They live on plan pages, in contracts or behind access tiers, so a share of priced models does not describe them, and the table marks them "Plan or contract level" instead.

Prices are US dollars per million tokens unless stated. A pair such as $4/$20 means $4 per million input tokens and $20 per million output tokens. Any claim I could not confirm against a vendor's public announcement carries the label UNVERIFIED.

All 35 Mechanisms, Ranked

Share of priced models and priced vendors in my full tracked population of 1,934 models.

#	MECHANISM	PRICED MODELS	PRICED VENDORS
1	Per Input and Output Token	82.7%	81.8%
2	Cached Input Read Discount	50.4%	63.6%
3	Reasoning (Thinking) Tokens Billed as Output	42.5%	36.4%
4	Per Image, Per Second, Per Minute, Per Character	35.0%	54.5%
5	Per-Call Server Tool Fees	34.3%	30.3%
6	Batch Discount	33.1%	27.3%
7	Data Residency and Regional Price Lists	29.5%	27.3%
8	Prompt-Size (Long Context) Bands	23.2%	27.3%
9	Speed, Priority and "Highspeed" Tiers	22.8%	27.3%
10	Free Tiers and $0 Price Lines	22.4%	18.2%
11	Cache and File Storage by Time	15.0%	15.2%
12	Flex (Slow, Cheap Synchronous) Tier	15.0%	12.1%
13	Separate Cache Write Charge	13.8%	15.2%
14	Billing-Event Rules	10.2%	18.2%
15	Cache Write Priced by Retention Time (TTL)	9.1%	9.1%
16	Promotional Pricing and Dated Step-Ups	8.7%	21.2%
17	Data-for-Discount Tiers	8.7%	6.1%
18	Hidden Token Overhead	7.9%	9.1%

© 2026 jerrywlambert.com

#	MECHANISM	PRICED MODELS	PRICED VENDORS
19	Fine-Tuning, Training and Fine-Tuned Inference	7.5%	12.1%
20	Container and Session Runtime Fees	6.3%	9.1%
21	Marketplace Consumption Units	5.5%	3.0%
22	Minimums, Top-Ups, Credits and Expiring Prepaid Balances	3.5%	15.2%
23	Time-of-Day and Day-of-Week Bands	2.0%	6.1%
24	Resale Markup and Reseller Chains	1.6%	9.1%
25	Subscription With Bundled, Rationed Usage	1.2%	9.1%
26	Legacy Model Names Billed as a Different Model	1.2%	6.1%
27	Commits, Take-or-Pay and Negotiated Volume Discounts	0.4%	3.0%
	Plan and Contract-Level Charges		
28	Per-User License, Platform Token Fees and Bring-Your-Own-Key Fees	Plan or contract level	
29	Access-Gated and Channel-Only Models	Plan or contract level	
30	Router Models Billed at the Routed Model	Plan or contract level	
31	Discount-Stacking Rules	Plan or contract level	
32	Feature and Format Surcharges	Plan or contract level	
33	Data Egress Fees	Plan or contract level	
34	Provisioned Capacity Units	Plan or contract level	
35	Open-Weights Licenses Free Below a Revenue Threshold	Plan or contract level	

© 2026 jerrywlambert.com

The Mechanisms

01 Per Input and Output Token

NAME · MECHANISM 1 OF 35

82.7% of priced models

DESCRIPTION

This is the base meter. The vendor prices input (prompt) tokens and output (completion) tokens separately, per million. Output almost always costs more.

EXAMPLE

Claude Opus 5.5 lists $4.00 input and $20.00 output. OpenAI gpt-5.5-pro lists $30.00 input and $180.00 output. Some vendors flip the pattern. TypeSafe Jev and Liquid d1 price output at $0.00 and bill input only.

WHO USES IT

Seen on 82.7% of priced models and 81.8% of priced vendors in my full tracked population. Vendors include AI21, Alibaba, Amazon, Anthropic, Baidu, ByteDance, Cohere, Cursor, DeepSeek, Google, IBM, Inception, Meta, Microsoft, MiniMax, Mistral, Moonshot, OpenAI, Perplexity, StepFun, Tencent, Thinking Machines, TypeSafe, Upstage, xAI, Xiaomi and Z.ai. The priced vendors without a token price are all media vendors.

THINGS TO WATCH

- Output-heavy work such as code generation and long reports is priced on the expensive side. At a 5x ratio, one output token costs as much as five input tokens.
- Token counts do not compare cleanly across vendors, or even across generations of one vendor's models (see Hidden Token Overhead).
- An output price of $0 moves the whole bill onto the prompt. TypeSafe Jev bills $0.042 input and $0 output.

WAYS TO OPTIMIZE

- Cap output length with a max-tokens setting and ask for structured, terse output.
- Route by task. Claude Haiku 5.5 at $0.10/$0.50 against Opus 5.5 at $4/$20 is a 40x spread inside one vendor.

THINGS TO AVOID

- Comparing models on input price alone.
- Treating a list price as the effective price. Nearly every mechanism that follows changes it.

© 2026 jerrywlambert.com

NAME · MECHANISM 2 OF 35

02 Cached Input Read Discount

50.4% of priced models

DESCRIPTION

When a request reuses a prompt prefix the vendor already holds in cache, those tokens bill at a cache-hit price below the normal input price.

EXAMPLE

VENDOR AND MODEL	CACHE READ PRICE
Anthropic, most models	0.1x input
Anthropic Opus 5.5 and Sonnet 5.5	0.05x input
Anthropic Fable 5.1	0.025x input ($0.25 against $10.00)
DeepSeek Flash, off-peak	$0.003 cached against $0.15 uncached (98 percent off)

WHO USES IT

Seen on 50.4% of priced models and 63.6% of priced vendors in my full tracked population. Vendors include Alibaba, Amazon, Anthropic, ByteDance, Cursor, DeepSeek, Google, Inception, Meta, MiniMax, Mistral, Moonshot, OpenAI, Perplexity, StepFun, Tencent, Thinking Machines, Upstage, xAI, Xiaomi and Z.ai.

THINGS TO WATCH

- Discounts vary 30-fold, from 0.0083x (Xiaomi MiMo-V2.6-Pro) to 0.25x. A blanket "90 percent off cached input" assumption is wrong for many models, and some models price a cache read at the full input price.
- On Anthropic, a cache hit also refreshes the cache timer.
- Some models cache only implicitly. Google says Gemini 4 Argon supports "implicit caching only (explicit caching is not supported)."

WAYS TO OPTIMIZE

- Put stable content first (system prompt, tool definitions, reference documents) and variable content last, so the prefix matches.
- Keep a conversation on one model. Caches do not carry across models.

THINGS TO AVOID

- Reordering tools or injecting a timestamp at the top of the prompt. Either one breaks the prefix.
- Billing a model with no published cache price as if it had one. A cache price with no first-party basis is an invented number.

© 2026 jerrywlambert.com

NAME · MECHANISM 3 OF 35

03 Reasoning (Thinking) Tokens Billed as Output

42.5% of priced models

DESCRIPTION

Internal reasoning tokens bill at the output price, even when the vendor hides or encrypts them.

EXAMPLE

- Google labels every Gemini output price "including thinking tokens." Gemini 4 Argon output is $10.00.
- Alibaba Qwen 3.6-35B-A3B lists one $2.25 output price for "thinking and non-thinking."
- xAI's Responses API always returns encrypted reasoning content for grok-4.7.

WHO USES IT

Seen on 42.5% of priced models and 36.4% of priced vendors in my full tracked population. Vendors include Alibaba, Anthropic, Baidu, DeepSeek, Google, Meta, MiniMax, Moonshot, OpenAI, Thinking Machines, xAI and Z.ai.

THINGS TO WATCH

- Reasoning effort settings change output volume, often by multiples, at the same unit price.
- DeepSeek defaults to thinking mode.

WAYS TO OPTIMIZE

- Set the lowest reasoning effort that passes your quality bar.
- Use non-thinking modes for simple tasks.

THINGS TO AVOID

- Forecasting output cost from the length of the visible answer.

© 2026 jerrywlambert.com

NAME · MECHANISM 4 OF 35

04 Per Image, Per Second, Per Minute, Per Character

35.0% of priced models

DESCRIPTION

The vendor prices media outside the token meter, or converts each asset into a fixed number of tokens.

EXAMPLE

UNIT	PUBLISHED EXAMPLES
Per image	xAI grok-imagine-image $0.02 and image-2.0 $0.04 per image. Z.ai GLM-Image $0.015. MiniMax image-01 $0.0035. Gemini 2.5 Flash Image $0.039 per 1024×1024 image (1,290 tokens).
Per second or per video	xAI video $0.020 to $0.080 per second. MiniMax H3 $0.025 to $0.13 per output second, by resolution. Z.ai CogVideoX-3 $0.20 per video.
Per megapixel	Black Forest Labs prices FLUX.2 by megapixel and bills FLUX 3 Video at the delivered resolution class.
Tokens per asset	Tencent TokenHub and BytePlus Seedance bill an image or video as a fixed token count times the token rate.
Per minute or per hour	OpenAI GPT-Live $0.05 per minute, billed per second. OpenAI gpt-transcribe $0.0045 per minute. xAI speech-to-text $0.10 per hour (REST). Xiaomi ASR $0.074 per hour. Gemini bills audio at 25 tokens per second.
Per character	OpenAI tts-1 $15 per million characters. xAI text-to-speech $15 per million characters. MiniMax speech $60 to $100 per million characters.
Modality uplift inside a token model	OpenAI gpt-realtime-2.1 audio input $32 against text input $4.

WHO USES IT

Seen on 35.0% of priced models and 54.5% of priced vendors in my full tracked population. Vendors include Alibaba, Amazon, Black Forest Labs, ByteDance, Google, Meta, Microsoft, MiniMax, Mistral, OpenAI, Pruna, Recraft, Runway, Stability AI, Tencent, Vidu, xAI and Z.ai.

THINGS TO WATCH

- Audio sent to a text-priced model can bill at the vendor's higher audio rate. Gemini Flash audio input runs about 2x text. The rate pair is UNVERIFIED for this article.

WAYS TO OPTIMIZE

- Pick resolution and duration deliberately. Output resolution drives video price.

THINGS TO AVOID

- Routing audio or image content through an endpoint priced only for text.

© 2026 jerrywlambert.com

NAME · MECHANISM 5 OF 35

05 Per-Call Server Tool Fees

34.3% of priced models

DESCRIPTION

When the vendor runs a tool on its own servers (web search, grounding, file search), it charges per call or per 1,000 calls. That fee sits on top of the tokens the tool's results add to the prompt.

EXAMPLE

VENDOR	TOOL	PRICE	NOTES
Anthropic	Web search	$10 per 1,000 searches plus tokens	A failed search is not billed. Web fetch has no extra fee.
OpenAI	Web search	$10 per 1,000 calls plus content tokens	The non-reasoning preview tool is $25 per 1,000 with free content tokens.
OpenAI	File search	$2.50 per 1,000 calls	
Google	Grounding with Search or Maps (Gemini 3 and newer)	$14 per 1,000 after 5,000 free requests per month	The allowance is shared across models.
Google	Grounding on Gemini 2.5	$35 per 1,000 grounded prompts (Search), $25 per 1,000 (Maps)	After the free daily allowance.
Z.ai	Web search	$0.01 per use	
MiniMax	web_search	$0.01 per request	
Xiaomi	Overseas search	$5 per 1,000	
DeepSeek	Search and Fetch	$0.75 per 1,000 successful requests	Free until 00:00 Beijing time on November 1, 2026 (16:00 UTC, 11:00 AM Central, on October 31).
xAI	Web, X search, code and collections tools	$2.50 to $10 per 1,000	UNVERIFIED. X Search bills per returned item.

On Google, one customer request can trigger several billed searches. xAI X Search bills per item returned.

WHO USES IT

Seen on 34.3% of priced models and 30.3% of priced vendors in my full tracked population. Vendors include Amazon, Anthropic, DeepSeek, Google, Meta, OpenAI, Perplexity, StepFun, xAI and Z.ai.

© 2026 jerrywlambert.com

THINGS TO WATCH

- The fee is per call, not per answer. One prompt can fan out into many searches.
- The tool's results come back as input tokens and stay in context on later turns.
- Free allowances are shared across models and reset monthly or daily, depending on the model family.

WAYS TO OPTIMIZE

- Cap tool calls per request.
- Prefer token-only tools where the vendor offers them, such as Anthropic web fetch and Google URL context.

THINGS TO AVOID

- Letting a server tool through a billing layer that cannot count it. If the call count is not metered, refuse the tool instead of billing tokens only.

© 2026 jerrywlambert.com

NAME · MECHANISM 6 OF 35

06 Batch Discount

33.1% of priced models

DESCRIPTION

You submit work asynchronously and accept a delay, typically up to 24 hours. In exchange you pay roughly half.

EXAMPLE

MODEL	STANDARD	BATCH
Anthropic Opus 5.5	$4/$20	$2/$10 (50 percent off input and output)
OpenAI gpt-6-astra	$10/$50	$5/$25
Google Gemini 4 Argon	$2/$10	$1/$5
Xiaomi MiMo-V2.6-Pro	$0.435/$0.87	$0.2175/$0.435

WHO USES IT

Seen on 33.1% of priced models and 27.3% of priced vendors in my full tracked population. Vendors include Alibaba, Amazon, Anthropic, ByteDance, Google, Moonshot, OpenAI, xAI and Xiaomi.

THINGS TO WATCH

- On Anthropic, the batch discount stacks with caching.
- Not every model or mode qualifies. Xiaomi Ultraspeed, Anthropic fast mode and Anthropic hosted sessions have no batch option.
- OpenAI's batch table omits cached-input discounts for several older models, gpt-4.1 among them.

WAYS TO OPTIMIZE

- Move evaluations, backfills, classification and nightly reports to batch.

THINGS TO AVOID

- Using batch for anything a user is waiting on.

© 2026 jerrywlambert.com

NAME · MECHANISM 7 OF 35

07 Data Residency and Regional Price Lists

29.5% of priced models

DESCRIPTION

The vendor adds a surcharge to keep inference inside a named geography, or publishes a separate price list for each region.

EXAMPLE

- Anthropic's inference_geo: "us" setting is 1.1x on every token category for Claude 4.6 and later.
- Claude on Amazon Bedrock regional endpoints and Google Cloud multi-region endpoints carries a 10 percent premium over global endpoints.
- OpenAI regional and FedRAMP endpoints carry a 10 percent uplift for models released on or after March 5, 2026.
- Xiaomi publishes separate China (RMB) and overseas (USD) tables. MiMo-V2.6-Pro input is ¥3.00 on one and $0.435 on the other. The Tencent Hunyuan prices in my Tariff are Singapore-region prices. Alibaba Model Studio publishes an International (Singapore) list.
- Tencent stopped new purchases on its old Hunyuan platform and moved buyers to TokenHub, which has its own price list.

WHO USES IT

Seen on 29.5% of priced models and 27.3% of priced vendors in my full tracked population. Vendors include Alibaba, Amazon, Anthropic, Cursor, IBM, OpenAI, Tencent, xAI and Xiaomi.

THINGS TO WATCH

- The surcharge multiplies every other mechanism, including fast mode and long-context bands.

WAYS TO OPTIMIZE

- Use residency only for data that requires it, and route the rest globally.

THINGS TO AVOID

- Accepting a residency parameter on a route that bills global rates.

© 2026 jerrywlambert.com

NAME · MECHANISM 8 OF 35

08 Prompt-Size (Long Context) Bands

23.2% of priced models

DESCRIPTION

When the prompt crosses a threshold, the whole request reprices at a higher band. The higher price applies to every token in the request, not only the tokens above the line.

EXAMPLE

MODEL	THRESHOLD	STANDARD	ABOVE THRESHOLD
OpenAI gpt-5.5	Above 272,000 input tokens	$5/$30	$10/$45
xAI grok-4.7	200,000 prompt tokens or more	$2.00/$6.00	$4.00/$12.00, "for all tokens in the request"
Google Gemini 2.5 Pro	Above 200,000	$1.25/$10	$2.50/$15
Anthropic Claude Haiku 5.5	Above 100,000 (cache reads count toward it)	$0.10/$0.50	$0.50/$2.50 (5x)
MiniMax M3	Above 512,000	$0.30/$1.20	$0.60/$2.40
Alibaba Qwen 3.7 Plus	Above 256,000	$0.40/$1.60	$1.20/$4.80

The counter-example is Anthropic. Claude 4.6 and later, except Haiku 5.5, bill the full 1M-token window at standard rates.

WHO USES IT

Seen on 23.2% of priced models and 27.3% of priced vendors in my full tracked population. Vendors include Alibaba, Anthropic, Baidu, Black Forest Labs, ByteDance, Google, MiniMax, OpenAI and xAI.

THINGS TO WATCH

- One token over the line can double the cost of the entire request.
- Thresholds differ by vendor, and so does the boundary. OpenAI's line is ">272K" and xAI's is "≥200K".

WAYS TO OPTIMIZE

- Trim or summarize context to stay under the threshold.
- Retrieve the passages a request needs instead of stuffing whole documents into the prompt.

THINGS TO AVOID

- Estimating cost from a blended average. Price each request against its own band.

© 2026 jerrywlambert.com

NAME · MECHANISM 9 OF 35

09 Speed, Priority and "Highspeed" Tiers

22.8% of priced models

DESCRIPTION

You pay more for faster output or for priority admission. Vendors sell this either as a request flag on the same model or as a separate model ID.

EXAMPLE

MODEL	TIER	PRICE	STANDARD
OpenAI gpt-6-astra	Fast (renamed from Priority on July 30, 2026), 2x	$20/$100	$10/$50
OpenAI gpt-6-astra	Ultrafast, 6x	$60/$300	$10/$50
Anthropic Opus 5.5	Fast mode, 2x	$8/$40	$4/$20
Google Gemini 4 Argon	Priority, 1.8x	$3.60/$18.00	$2.00/$10.00
MiniMax M3	Priority, 1.5x	$0.45/$1.80	$0.30/$1.20
Moonshot kimi-k3.1-pro-highspeed	Separate model ID	$4.00 input, $0.20 cached, $4.00 5-minute write, $20.00 output	$2.00, $0.20, $2.00, $10.00
Xiaomi mimo-v2.6-pro-ultraspeed	Separate model ID, 10x	$4.35/$8.70	$0.435/$0.87
xAI	Priority, 2x on every token type	2x standard	Standard

Anthropic fast mode on Opus 5 and Opus 4.8 is $10/$50. It stacks with caching and data residency, and it is not available with batch.

WHO USES IT

Seen on 22.8% of priced models and 27.3% of priced vendors in my full tracked population. Vendors include Amazon, Anthropic, Cursor, Google, MiniMax, Moonshot, OpenAI, xAI and Xiaomi.

© 2026 jerrywlambert.com

THINGS TO WATCH

- The multiplier is not uniform. I see 1.5x, 1.8x, 2x, 6x and 10x in published prices.
- Moonshot Highspeed doubles input and output but keeps the cached read at $0.20.
- A tier flag on the request can silently raise the price if a client library sets it.

WAYS TO OPTIMIZE

- Reserve speed tiers for user-facing paths where latency matters, and keep them off background work.

THINGS TO AVOID

- Passing a tier flag through to a vendor when the platform bills at standard rates. Either price the multiplier or reject the flag.

© 2026 jerrywlambert.com

10 Free Tiers and $0 Price Lines

NAME · MECHANISM 10 OF 35

22.4% of priced models

DESCRIPTION

The vendor prices a model, a unit or a usage band at zero, either permanently or for a limited time.

EXAMPLE

- **Fully free models.** Google Gemma 4 12B and 31B are free tier only and listed as "Not available" on paid. Z.ai lists GLM-4.7-Flash, GLM-4.5-Flash and GLM-4.6V-Flash as free. NVIDIA runs free, rate-limited endpoints with no per-token billing.
- **Free units.** TypeSafe Jev output is $0. Amazon Nova cache writes are $0. Z.ai and Xiaomi cache writes are free for a limited time.
- **Free allowances.** Anthropic code execution includes 1,550 free hours per organization per month. Google grounding includes 5,000 requests per month. OpenAI includes 1 GB of file storage, and OpenAI moderation is free.
- **Free trial credit.** Anthropic gives new users "a small amount of free credits."

WHO USES IT

Seen on 22.4% of priced models and 18.2% of priced vendors in my full tracked population. Vendors include AI21, Alibaba, Google, MiniMax, Stability AI and Z.ai. NVIDIA's build endpoints and the free variants on one large model aggregator also price at $0.

THINGS TO WATCH

- A free tier can carry a data cost. Google's free tier states "Content used to improve our products."
- A $0 line with no first-party source looks exactly like a missing price.

WAYS TO OPTIMIZE

- Use free allowances for development and evaluation traffic.

THINGS TO AVOID

- Recording an unknown price as $0. A $0 line should carry its source and the vendor's own word "free."

© 2026 jerrywlambert.com

11

NAME · MECHANISM 11 OF 35

Cache and File Storage by Time

15.0% of priced models

DESCRIPTION

The vendor charges to keep data available, priced per token-hour or per GB-day.

EXAMPLE

ITEM	PRICE
Gemini context cache storage	$0.50 per million tokens per hour through December 31, 2026, then $1.00 from January 1, 2027
OpenAI file search storage	$0.10 per GB per day after 1 GB free
OpenAI ChatKit upload storage	$0.10 per GB per day after 1 GB free
Z.ai cached input storage	"Limited-time Free"

WHO USES IT

Seen on 15.0% of priced models and 15.2% of priced vendors in my full tracked population. Vendors include ByteDance, Google, StepFun, xAI and Z.ai.

THINGS TO WATCH

- A 1M-token Gemini cache held for 24 hours costs $12 in storage today and $24 from January 1, 2027, before a single read.

WAYS TO OPTIMIZE

- Set short cache lifetimes.
- Delete stores you no longer query.

THINGS TO AVOID

- Long-lived explicit caches on low-traffic prompts.

© 2026 jerrywlambert.com

12

NAME · MECHANISM 12 OF 35

Flex (Slow, Cheap Synchronous) Tier

15.0% of priced models

DESCRIPTION

The request stays synchronous but accepts slower or interruptible service in exchange for a batch-like discount.

EXAMPLE

MODEL	STANDARD	FLEX
OpenAI gpt-6-astra	$10/$50	$5/$25
Google Gemini 4 Argon	$2/$10	$1/$5

WHO USES IT

Seen on 15.0% of priced models and 12.1% of priced vendors in my full tracked population. Vendors include Amazon, ByteDance, Google and OpenAI.

THINGS TO WATCH

- Flex trades availability for price. Expect slower responses or capacity errors.

WAYS TO OPTIMIZE

- Use Flex for non-urgent synchronous jobs that cannot be restructured as batch.

THINGS TO AVOID

- Retrying a failed Flex request on Fast without a cost check.

© 2026 jerrywlambert.com

13 Separate Cache Write Charge

NAME · MECHANISM 13 OF 35

13.8% of priced models

DESCRIPTION

The vendor bills the cache write as its own line. Depending on the vendor, the write costs more than input, the same as input, less than input, or nothing.

EXAMPLE

TREATMENT	EXAMPLE
Premium	Anthropic 5-minute write at 1.25x input: Opus 5.5 $5.00 against $4.00
Premium	OpenAI newer models, 30-minute write at 1.25x: gpt-6-astra $12.50 against $10.00
Discount	DeepSeek V4 Pro write $0.435 against $1.32 peak input. UNVERIFIED. DeepSeek's current page publishes no write price.
Free	Z.ai lists cached input storage as "Limited-time Free" on all paid GLM models
Free	Xiaomi lists "Cache Write: Limited-time Free"
Folded in	StepFun has no separate write line. The cache-miss input price includes the write.

WHO USES IT

Seen on 13.8% of priced models and 15.2% of priced vendors in my full tracked population. Vendors include Alibaba, Anthropic, Moonshot, OpenAI and StepFun.

THINGS TO WATCH

- A write premium means caching loses money if nothing reads the prefix again. Anthropic says a 5-minute write pays off after one read and a 1-hour write after two.
- "Limited-time free" has no end date on either the Z.ai page or the Xiaomi page.

WAYS TO OPTIMIZE

- Cache only the prefixes you expect to reuse within the retention window.

THINGS TO AVOID

- Treating a free write as permanent. Record it as a promotion and check it again at every price review.

© 2026 jerrywlambert.com

14 Billing-Event Rules

NAME · MECHANISM 14 OF 35

10.2% of priced models

DESCRIPTION

These rules decide which events are billable, separate from the price itself.

EXAMPLE

- DeepSeek bills Search and Fetch per successful request.
- Anthropic does not bill a web search that errors.
- MiniMax does not deduct points for failed or security-blocked videos, bills the original inputs again when a video is regenerated, and charges $1.50 per cloned voice and $3 per designed voice at first use.
- Meta bills SAM 3.1 video at a minimum of 50 frames. UNVERIFIED.
- xAI charges $0.05 per request for a policy violation, even when it blocks the request.
- Kimi and Perplexity bill only non-empty results. Vidu, BytePlus, Tencent and Black Forest Labs do not charge for failed or moderated jobs.
- Runway sets minimum charges of 56 credits on aleph2 and 64 on seedance2_mini. Vidu Q2 charges 55 credits for the first second, then 15 per second.

WHO USES IT

Seen on 10.2% of priced models and 18.2% of priced vendors in my full tracked population. Vendors include ByteDance, Meta, MiniMax, Tencent, Vidu and xAI.

THINGS TO WATCH

- The same price can produce different bills depending on success rules and minimums.

WAYS TO OPTIMIZE

- Count only billable events in your own meter, so invoices reconcile.

THINGS TO AVOID

- Billing your customers for failed calls the vendor did not charge you for.

© 2026 jerrywlambert.com

NAME · MECHANISM 15 OF 35

15 Cache Write Priced by Retention Time (TTL)

9.1% of priced models

DESCRIPTION

The longer the vendor holds the cache, the more the write costs.

EXAMPLE

MODEL	SHORT RETENTION WRITE	LONG RETENTION WRITE
Anthropic Opus 5.5	5-minute: $5 (1.25x input)	1-hour: $8 (2x input)
OpenAI gpt-6-astra	30-minute: $12.50	12-hour: $25.00 (2x the 30-minute)
Moonshot kimi-k3.1-pro	5-minute: $2.00	1-hour: $4.00
Moonshot kimi-k3.1-pro Highspeed	5-minute: $4.00	1-hour: $8.00
Moonshot kimi-k3	5-minute: $3.00	1-hour: $6.00

Moonshot selects the retention time with a request option, prompt_cache_options.ttl.

WHO USES IT

Seen on 9.1% of priced models and 9.1% of priced vendors in my full tracked population. Vendors include Anthropic, Moonshot and OpenAI.

THINGS TO WATCH

- A platform that stores only one write column bills the long-retention write at the short-retention price. On Kimi K3.1 Pro that is $2.00 billed against $4.00 owed.
- Each vendor uses a different request field to select retention, so a check written for one vendor's field misses another's.

WAYS TO OPTIMIZE

- Use the long retention only when the gaps between reuses exceed the short window. On Anthropic it needs two reads to break even.

THINGS TO AVOID

- Metering one retention time and passing the other through unpriced. Either price every retention column or refuse the long-retention option.

© 2026 jerrywlambert.com

NAME · MECHANISM 16 OF 35

16 Promotional Pricing and Dated Step-Ups

8.7% of priced models

DESCRIPTION

The vendor publishes a price with an announced end date, or schedules an increase in advance.

EXAMPLE

MODEL	CURRENT PRICE	WHAT CHANGES
Google Gemini 4 Argon	$2/$10 through November 30, 2026	$4/$20 from December 1, 2026
Google Gemini 3.8, 3.7 and 3.6 Flash	$0.75/$3.75 through December 31, 2026	$1.50/$7.50 from January 1, 2027
OpenAI GPT-5.6 Sol	Promotional pricing "available at least through November 21, 2026"	No successor price stated
MiniMax M3	"Permanent 50% off" list	No end date
DeepSeek Search and Fetch	Free	$0.75 per 1,000 successful requests from November 1, 2026
Inception Mercury 3, Upstage Solar Mini 4	50 percent off at launch	Launch discount

WHO USES IT

Seen on 8.7% of priced models and 21.2% of priced vendors in my full tracked population. Vendors include Alibaba, Google, MiniMax, OpenAI, Upstage, Xiaomi and Z.ai.

THINGS TO WATCH

- Promotions without end dates, such as the Z.ai and Xiaomi "limited-time free" lines, can end without notice.

WAYS TO OPTIMIZE

- Load the successor price with its start date ahead of time, so billing changes on the right day.

THINGS TO AVOID

- Carrying promotional prices into budgets or contracts past the stated end date.

© 2026 jerrywlambert.com

17 Data-for-Discount Tiers

NAME · MECHANISM 17 OF 35

8.7% of priced models

DESCRIPTION

The vendor lowers the price in exchange for the right to use your data.

EXAMPLE

- Google's invite-only Gemini API Contributor Tier prices Gemini 4 Argon at $1/$5, half of standard, with "Used to improve our products: Yes."
- OpenAI o4-mini fine-tuned inference with data sharing is $2/$8 against $4/$16.
- Mistral Labs models are free in exchange for mandatory use of your data for training.

WHO USES IT

Seen on 8.7% of priced models and 6.1% of priced vendors in my full tracked population. Vendors include Google and Meta.

THINGS TO WATCH

- You pay for the discount with data rights. Check contracts and regulated data before you opt in.

WAYS TO OPTIMIZE

- Use these tiers only for public or synthetic workloads.

THINGS TO AVOID

- Sending customer or regulated data through a contributor tier.

© 2026 jerrywlambert.com

18 Hidden Token Overhead

NAME · MECHANISM 18 OF 35

7.9% of priced models

DESCRIPTION

These are tokens you pay for but did not write: tool system prompts, tool definitions, tokenizer changes and fixed blocks.

EXAMPLE

- Anthropic adds a tool-use system prompt of 286 tokens (Opus 5.5) to 675 tokens (Opus 4.7) whenever tools are present. The computer-use toolset adds about 4,500 input tokens and the browser toolset about 6,600.
- Claude 4.7 and later use a tokenizer that "produces approximately 30% more tokens for the same text."
- OpenAI bills web search content for gpt-4o-mini and gpt-4.1-mini as a fixed block of 8,000 input tokens per call.
- Meta runs the other way. It adds its own steering tokens and leaves them off the bill.

WHO USES IT

Seen on 7.9% of priced models and 9.1% of priced vendors in my full tracked population. Vendors include Anthropic, Meta and OpenAI.

THINGS TO WATCH

- A model with a lower list price can cost more per task if it emits 30 percent more tokens.

WAYS TO OPTIMIZE

- Compare models on cost per completed task, not cost per token.
- Declare only the tools a request needs.

THINGS TO AVOID

- Carrying old per-token benchmarks across a tokenizer change.

© 2026 jerrywlambert.com

NAME · MECHANISM 19 OF 35

19 Fine-Tuning, Training and Fine-Tuned Inference

7.5% of priced models

DESCRIPTION

The vendor charges for training, then charges a higher inference price for the customized model.

EXAMPLE

ITEM	PRICE
OpenAI gpt-4.1 fine-tuning	$25 per million training tokens
OpenAI gpt-4.1 fine-tuned inference	$3.00/$12.00 against $2.00/$8.00 for the base model
OpenAI o4-mini reinforcement fine-tuning	$100 per training hour

OpenAI is winding the platform down and has closed it to new users. Amazon Nova bills custom-model inference at the base-model rate.

WHO USES IT

Seen on 7.5% of priced models and 12.1% of priced vendors in my full tracked population. Vendors include Amazon, Microsoft, OpenAI and Thinking Machines.

THINGS TO WATCH

- In reinforcement fine-tuning, grader tokens bill at the grader model's rate.
- Fine-tuned models retire with their base model.

WAYS TO OPTIMIZE

- Try prompting and caching before training.

THINGS TO AVOID

- Building on a fine-tuning platform with an announced wind-down.

© 2026 jerrywlambert.com

NAME · MECHANISM 20 OF 35

20 Container and Session Runtime Fees

6.3% of priced models

DESCRIPTION

The vendor charges by time for a sandbox or a hosted session, separately from tokens.

EXAMPLE

PRODUCT	PRICE	RULES
Anthropic code execution	$0.05 per container-hour after 1,550 free hours per month	5-minute minimum. Free when used with the newer web search or fetch tools. Files in the request start billing even if the tool is never called.
Anthropic hosted sessions	$0.08 per session-hour	Counted only while the session is "running."
OpenAI Hosted Shell and Code Interpreter	$0.03 (1 GB) to $1.92 (64 GB) per 20-minute session	Billed by the minute, with a 5-minute minimum.
Google code execution	No runtime charge	Tokens only.

WHO USES IT

Seen on 6.3% of priced models and 9.1% of priced vendors in my full tracked population. Vendors include Anthropic, Perplexity and Thinking Machines.

THINGS TO WATCH

- Memory size multiplies OpenAI's price 64-fold.
- Minimum billable durations make many short calls expensive.

WAYS TO OPTIMIZE

- Reuse a container across steps.
- Pick the smallest memory class that does the job.

THINGS TO AVOID

- Attaching files "just in case" on Anthropic. The attachment starts the meter.

© 2026 jerrywlambert.com

NAME · MECHANISM 21 OF 35

21 Marketplace Consumption Units

5.5% of priced models

DESCRIPTION

The vendor prices usage in dollars, then converts it into the marketplace's own billing unit.

EXAMPLE

Claude Platform on AWS and Claude in Microsoft Foundry bill in Claude Consumption Units at $0.01 per unit (100 units equal $1.00), metered hourly. Discounts show up as fewer units.

WHO USES IT

Seen on 5.5% of priced models and 3.0% of priced vendors in my full tracked population. Vendors include Anthropic.

THINGS TO WATCH

- The cloud bill shows a single unit line, so per-model detail lives in the vendor console.

WAYS TO OPTIMIZE

- Use marketplace billing to draw down existing cloud commitments.

THINGS TO AVOID

- Reconciling marketplace units without the vendor's per-model breakdown.

© 2026 jerrywlambert.com

NAME · MECHANISM 22 OF 35

22 Minimums, Top-Ups, Credits and Expiring Prepaid Balances

3.5% of priced models

DESCRIPTION

Vendors set spend floors, require unlock payments, price in credits and let prepaid balances lapse.

EXAMPLE

- Moonshot Kimi K3.1 Pro is "unlocked after a successful top-up (minimum $1)."
- Anthropic code execution and OpenAI containers bill a 5-minute minimum.
- MiniMax video points "expire with the package and do not roll over."
- DeepSeek deducts from a granted balance before a topped-up balance.
- Runway, Black Forest Labs, Stability AI and Ideogram price in credits at $0.01 each. Vidu prices a credit at ¥0.03125. StepFun excludes vouchers from tier top-ups.

WHO USES IT

Seen on 3.5% of priced models and 15.2% of priced vendors in my full tracked population. Vendors include MiniMax, Recraft, Runway, Stability AI and Vidu.

THINGS TO WATCH

- Prepaid credit is often non-refundable.
- Expiry turns unused prepay into a price increase.

WAYS TO OPTIMIZE

- Buy prepay in amounts you will consume within the validity window.

THINGS TO AVOID

- Holding large prepaid balances with vendors that apply expiry rules.

© 2026 jerrywlambert.com

NAME · MECHANISM 23 OF 35

23 Time-of-Day and Day-of-Week Bands

2.0% of priced models

DESCRIPTION

The same model costs different amounts depending on when the request runs.

EXAMPLE

DeepSeek's off-peak rates are half of peak. Peak runs 01:00 to 04:00 and 06:00 to 10:00 UTC, Monday through Friday, excluding Chinese public holidays. Weekends and Chinese public holidays are off-peak all day.

MODEL	PEAK	OFF-PEAK
deepseek-flash	$0.30/$1.20	$0.15/$0.60
deepseek-v4-pro	$1.32/$3.96	$0.66/$1.98

Vidu offers off-peak pricing on viduq3-pro and viduq3-turbo only when the audio parameter is true.

WHO USES IT

Seen on 2.0% of priced models and 6.1% of priced vendors in my full tracked population. Vendors include DeepSeek and Vidu. Tencent TokenHub and BytePlus publish time bands on resold DeepSeek models.

THINGS TO WATCH

- Peak falls in the US evening and overnight: 8:00 PM to 11:00 PM and 1:00 AM to 5:00 AM Central time, Sunday night through Friday morning, under daylight time.
- The holiday calendar is Chinese, not American. A meter that ignores it overbills on those days.

WAYS TO OPTIMIZE

- Schedule DeepSeek batch-style jobs for US business hours and weekends.

THINGS TO AVOID

- Pricing all DeepSeek traffic at one band.

© 2026 jerrywlambert.com

NAME · MECHANISM 24 OF 35

24 Resale Markup and Reseller Chains

1.6% of priced models

DESCRIPTION

A reseller prices on top of vendor cost, and resellers can sit under other resellers.

EXAMPLE

Illustrative example, not a vendor price.

LEVEL	COST TO THIS LEVEL	SELLS AT	MARGIN
Vendor list		$8.00	
Platform, 25 percent markup	$8.00	$10.00	$2.00
Reseller, 20 percent over its own cost	$10.00	$12.00	$2.00

Each level earns $2.00. Each level's markup applies to the price that level pays, never to vendor list.

Resellers also retire model IDs early, and some sell their own hosted copy and a "Vendor Direct" copy of one model at different prices.

WHO USES IT

Seen on 1.6% of priced models and 9.1% of priced vendors in my full tracked population. Vendors include MiniMax, Moonshot and Runway. Tencent, Perplexity and BytePlus resell other vendors' models, and Cursor sells third-party models at list price.

THINGS TO WATCH

- A model a reseller has not priced should be refused at that level, not fall through to the parent's price.

WAYS TO OPTIMIZE

- Keep a dated price trail from vendor list to end customer for every call.

THINGS TO AVOID

- Applying a markup to vendor list when the reseller's actual cost is the parent's price.

© 2026 jerrywlambert.com

NAME · MECHANISM 25 OF 35

25 Subscription With Bundled, Rationed Usage

1.2% of priced models

DESCRIPTION

The customer pays a flat monthly fee that includes usage measured in opaque limits instead of tokens.

EXAMPLE

- Claude Pro is $20 per month ($17 on annual billing). Max starts at $100 with 5x or 20x Pro usage per five-hour session, plus weekly limits, and includes $100 to $200 per month of API credits. Overage on paid plans can continue "at standard API rates."
- Claude Duet is a research preview, by waitlist, for US Pro and Max subscribers only. It runs on Claude Opus 5.5 with no model choice. It is included at no extra cost and draws on the plan's five-hour and weekly limits, and its background checks count too.
- MiniMax Audio Subscription runs $5 to $999 per month for 100,000 to 20,000,000 audio points. MiniMax Video Packages run $1,000 to $6,000 per month, and unused points expire.
- Xiaomi says Token Plan quota "is not interoperable" with a pay-as-you-go balance.
- Cursor bills cloud runs from "your plan's usage … at the selected model's rate." Hark Pro is $20 (2x usage) and $100 (10x). Both are UNVERIFIED.

WHO USES IT

Seen on 1.2% of priced models and 9.1% of priced vendors in my full tracked population. Vendors include Cursor, MiniMax and xAI. Cognition, Ideogram and Krea also sell plan-based usage.

THINGS TO WATCH

- Limits are expressed as multiples of another plan, not in tokens.
- Background jobs spend the same pool as chat.

WAYS TO OPTIMIZE

- Measure the cost per unit of work under the plan against API rates before you commit.

THINGS TO AVOID

- Assuming "included" means unlimited.

© 2026 jerrywlambert.com

NAME · MECHANISM 26 OF 35

26 Legacy Model Names Billed as a Different Model

1.2% of priced models

DESCRIPTION

An old model name keeps working, but a different model serves the request and sets the bill.

EXAMPLE

- DeepSeek: "The legacy names deepseek-v4-flash and deepseek-v4-flash-vision-exp are still accepted, but the corresponding models have been retired, their requests are served by the DeepSeek-V4.1-Flash-1017 model and billed at the Flash price."
- xAI points unversioned names and "-latest" names at newer releases.
- Google routes Gemini 3.5 Flash to 3.6 Flash. BytePlus serves glm-4-7 with a ByteDance dola-seed model, a substitution across vendors.

WHO USES IT

Seen on 1.2% of priced models and 6.1% of priced vendors in my full tracked population. Vendors include ByteDance and Perplexity. DeepSeek, xAI, MiniMax and Google all publish rules that point an old or unversioned name at a newer model.

THINGS TO WATCH

- The price can stay the same while the model changes, or the reverse.

WAYS TO OPTIMIZE

- Pin dated model versions in production.

THINGS TO AVOID

- Billing one model's price for another model's output without recording the substitution.

© 2026 jerrywlambert.com

NAME · MECHANISM 27 OF 35

27 Commits, Take-or-Pay and Negotiated Volume Discounts

0.4% of priced models

DESCRIPTION

The customer promises to spend a set amount in return for a lower price, and the vendor still bills any unused commitment.

EXAMPLE

- Anthropic's sales-assisted Enterprise plan lists "Tiered incentives on committed spend" and negotiated volume discounts.
- Google Enterprise offers provisioned throughput and volume-based discounts.
- Upstage publishes committed-use pricing, and AI21 and Reka sell committed use through sales.
- Illustrative example, not a vendor price: a $1,000 monthly take-or-pay commit with $700 used produces a $300 shortfall line, with no rollover.

WHO USES IT

Seen on 0.4% of priced models and 3.0% of priced vendors in my full tracked population. Vendors include Upstage. AI21 and Reka offer committed use through sales.

THINGS TO WATCH

- Whether the vendor discount reaches the end customer or stays with the reseller as margin.
- When the shortfall is billed, monthly or at the end of the term.

WAYS TO OPTIMIZE

- Commit at the floor of forecast usage, not the midpoint.

THINGS TO AVOID

- Commits tied to a single model generation. Prices fall and models retire within a term.

© 2026 jerrywlambert.com

Plan and Contract-Level Charges

These eight mechanisms live on plan pages, in contracts or behind access tiers rather than on a model's own price row. They are numbered 28 to 35.

NAME · MECHANISM 28 OF 35

28 Per-User License, Platform Token Fees and Bring-Your-Own-Key Fees

Plan or contract level

WHERE IT SHOWS UP

Plan page and contract: per-seat plans and platform fees on usage, including BYOK.

DESCRIPTION

The customer pays a fixed fee per user, a platform fee per token on top of list price, or a fee when it supplies its own vendor API key (bring your own key, or BYOK).

EXAMPLE

- Claude Enterprise costs US$20 per seat per month, billed annually, with a 20-seat minimum, plus usage at API rates.
- One large model aggregator charges a BYOK fee of 5 percent of what the same model would cost on its platform. On self-serve plans the fee applies after a $25,000 monthly list-price allowance.
- Cursor adds a $0.25 per million token rate on eligible third-party models on Teams and Enterprise, including BYOK usage.

WHO USES IT

Anthropic, Cursor and a large model aggregator publish these fees as plan and platform terms.

THINGS TO WATCH

- With BYOK, the vendor bills the customer directly and the platform's fee sits on top.
- A request that falls back from the customer's key to the platform's shared capacity is billed to the platform account.

WAYS TO OPTIMIZE

- Match the pricing model to volume. Per-user pricing favors heavy users, and a percentage fee favors light ones.

THINGS TO AVOID

- Mixing resale and BYOK billing in the same billing cycle. A BYOK request that lands on the platform's own key should be blocked, not silently billed either way.

© 2026 jerrywlambert.com

NAME · MECHANISM 29 OF 35

29 Access-Gated and Channel-Only Models

Plan or contract level

WHERE IT SHOWS UP

Access tier: vetted, invited or single-channel access.

DESCRIPTION

A vendor prices a model or tier but sells it only to vetted or invited customers, or only through one channel.

EXAMPLE

- OpenAI limits its Rosalind models to approved research through OpenAI's trusted-access program.
- Anthropic lists Claude Mythos 5 and 5.1 at $10 input as "limited availability."
- Google's Gemini API Contributor Tier is a private preview by invitation.
- xAI sells Grok 4.7 Fast, the same Grok 4.7 model on faster infrastructure at 2x standard token rates (1.5x for long context), only in Cursor and Grok Build, billed through the plan there.
- Cursor's Plus plan runs only Cursor's own models at fixed settings.

WHO USES IT

OpenAI, Anthropic, Google, xAI and Cursor publish these terms in their public announcements.

THINGS TO WATCH

- A published price you cannot buy is not a benchmark for your own costs.
- A channel-only model bills through someone else's plan, so it never shows up on the vendor's own invoice.

WAYS TO OPTIMIZE

- Check eligibility before you design around a gated model.
- Compare the channel price against the vendor's standard and priority rates.

THINGS TO AVOID

- Treating a gated price list as a public offer in a forecast or a customer quote.

© 2026 jerrywlambert.com

NAME · MECHANISM 30 OF 35

30 Router Models Billed at the Routed Model

Plan or contract level

WHERE IT SHOWS UP

Plan page: the router ID carries no price of its own.

DESCRIPTION

A model ID has no price of its own. The vendor bills each request at the rate of whichever model the router picks.

EXAMPLE

- Runway: "Generations routed through a Model Router are billed at the standard rate of whichever model the router selects." The response reports the model used and the cost in credits.
- A large model aggregator lists its auto, fusion, pareto-code and Body Builder routers with no fixed rate.
- TypeSafe Jev Router and NVIDIA Switchyard are listed as routers with no fixed rate. Cursor Router (Balance and Intelligence) bills at the routed model's list price.

WHO USES IT

Runway, Cursor, TypeSafe, NVIDIA and a large model aggregator publish router billing in their public announcements.

THINGS TO WATCH

- The same prompt can cost different amounts on different days as the router changes its pick.
- A router ID tells you nothing about price until the response comes back.

WAYS TO OPTIMIZE

- Read the model and cost fields in each response and meter on those.
- Pin a specific model for work with a fixed budget.

THINGS TO AVOID

- Pricing a router ID at a single blended rate.

© 2026 jerrywlambert.com

31 Discount-Stacking Rules

NAME · MECHANISM 31 OF 35

Plan or contract level

WHERE IT SHOWS UP

Plan page: the vendor's pricing terms set how discounts combine.

DESCRIPTION

The vendor publishes whether discounts and multipliers combine, and in what order.

EXAMPLE

- Alibaba: batch cuts input and output by 50 percent and context cache discounts input only, and the two discounts cannot be combined.
- xAI applies prompt caching discounts before its 2x priority multiplier, across input, output, cached and reasoning tokens.
- Anthropic stacks cache multipliers on fast mode and data residency on fast mode, combines batch and caching, and does not offer fast mode with batch.

WHO USES IT

Alibaba, xAI and Anthropic publish stacking rules in their public announcements.

THINGS TO WATCH

- The order matters. A discount applied before a multiplier gives a different price than the same discount applied after it.
- Two discounts that look additive may be mutually exclusive.

WAYS TO OPTIMIZE

- Write each vendor's stacking order into your price logic and test it with worked cases.

THINGS TO AVOID

- Multiplying every discount together by default.

© 2026 jerrywlambert.com

NAME · MECHANISM 32 OF 35

32 Feature and Format Surcharges

Plan or contract level

WHERE IT SHOWS UP

Plan page: request options listed as multipliers or add-ons.

DESCRIPTION

A request option or output format multiplies the base price or adds a surcharge to it.

EXAMPLE

VENDOR	OPTION	EFFECT
Luma	HDR output	2x SDR pricing
Luma	HDR plus EXR output	3x SDR pricing
Runway	ProRes or PNG sequence	5 credits per second on top of the model rate
Runway	10-bit and HDR profiles	20 credits per second, 40 above about 4K
Alibaba	prompt_extend on one z-image variant	$0.03 per image against $0.015 with it off
Vidu	Audio output on image-to-video and reference-to-video	15 extra credits

WHO USES IT

Luma, Runway, Alibaba and Vidu publish these surcharges in their public announcements.

THINGS TO WATCH

- A request parameter left at its default can turn a surcharge on without anyone choosing it.
- Surcharges stack on per-second and resolution pricing.

WAYS TO OPTIMIZE

- Set format and feature flags explicitly on every request.
- Generate in SDR and convert only the shots that need HDR.

THINGS TO AVOID

- Passing format options through a billing layer that prices only the base rate.

© 2026 jerrywlambert.com

NAME · MECHANISM 33 OF 35

33 Data Egress Fees

Plan or contract level

WHERE IT SHOWS UP

Plan page: a storage and download fee schedule.

DESCRIPTION

The vendor charges per GiB downloaded for data you stored with it.

EXAMPLE

ITEM	PRICE
xAI file downloads	$0.20 per GiB downloaded
xAI collection downloads	$0.20 per GiB downloaded
xAI file storage	$0.025 per GiB per day
xAI collection storage	$0.10 per GiB per day

WHO USES IT

xAI publishes egress fees in its public announcements.

THINGS TO WATCH

- Storage and download are separate meters. Data you pay to keep, you pay again to take out.

WAYS TO OPTIMIZE

- Keep a local copy of every file you upload.

THINGS TO AVOID

- Using vendor file storage as a general archive.

© 2026 jerrywlambert.com

NAME · MECHANISM 34 OF 35

34 Provisioned Capacity Units

Plan or contract level

WHERE IT SHOWS UP

Contract: a capacity purchase with a minimum unit count.

DESCRIPTION

The customer buys throughput in units, sized in requests per minute with a minimum purchase, instead of paying per request.

EXAMPLE

For example, Luma sells Uni-1.1 capacity in units. One unit equals 1 request per minute on Base or 0.4 requests per minute on Max, with a minimum of 4 units.

WHO USES IT

Luma publishes capacity units with a 4-unit minimum; AWS Bedrock, Azure and Google also sell provisioned throughput.

THINGS TO WATCH

- You pay for the capacity whether you use it or not.
- The same unit buys less throughput on the larger model.

WAYS TO OPTIMIZE

- Size units to steady traffic and send bursts to pay-per-request endpoints.

THINGS TO AVOID

- Buying units for traffic you have not measured.

© 2026 jerrywlambert.com

NAME · MECHANISM 35 OF 35

35 Open-Weights Licenses Free Below a Revenue Threshold

Plan or contract level

WHERE IT SHOWS UP

Contract: a license that applies above a revenue threshold.

DESCRIPTION

The vendor lets you run its open-weights models for free until your company's revenue passes a threshold, then requires a paid license.

EXAMPLE

Liquid AI lets developers, researchers, non-profits and businesses with less than $10M in annual revenue use its open-weights models in production at no cost. Businesses over $10M annual revenue need a paid license.

WHO USES IT

Liquid AI publishes this license in its public announcements.

THINGS TO WATCH

- The trigger is your revenue, not your usage. Growth alone can turn a free model into a licensed one.

WAYS TO OPTIMIZE

- Track the threshold against your revenue forecast and budget the license before you cross it.

THINGS TO AVOID

- Building a product on free weights without reading the revenue terms.

© 2026 jerrywlambert.com

How the Mechanisms Stack

Vendors publish how their multipliers combine. On Anthropic, cache multipliers stack on fast mode, data residency stacks on fast mode, batch and caching combine, and fast mode is not available with batch.

Anthropic's own worked example: a one-hour hosted session on Opus 5.5 with 50,000 input tokens and 15,000 output tokens costs $0.58 ($0.20 input, $0.30 output, $0.08 runtime). If 40,000 of those input tokens are cache reads, the session costs $0.428.

The order I use for a billing system: select the prompt-size band, apply multipliers (speed, residency), apply resale rules, then volume tiers and discounts, then credits, then draw down any commit.

Methodology

Built from vendors' public announcements across my full tracked population of 1,934 models, checked October 2026. Percentages are shares of the priced models and priced vendors in that population.

Appendix B. Claims Labeled UNVERIFIED

1. UNVERIFIED. xAI server tool fees of $2.50 to $10 per 1,000 calls (Mechanism 5).
2. UNVERIFIED. DeepSeek V4 Pro cache write of $0.435. DeepSeek's current page publishes no write price (Mechanism 13).
3. UNVERIFIED. Mistral batch and priority rates, and xAI batch rates (Mechanism 6).
4. UNVERIFIED. Which OpenAI models carry the 10 percent residency uplift for models released on or after March 5, 2026 (Mechanism 7).
5. UNVERIFIED. The Gemini Flash audio-input rate pair, about 2x text (Mechanism 4).
6. UNVERIFIED. Fine-tuned model hosting fees at vendors other than OpenAI (Mechanism 19).
7. UNVERIFIED. Cursor "plan usage at the selected model's rate" and the Hark $20 and $100 plan multiples (Mechanism 25).
8. UNVERIFIED. Meta SAM 3.1 minimum of 50 frames (Mechanism 14).
9. UNVERIFIED. Gemini 3.7 Flash batch, Flex and Priority availability. It is not shown as its own section on Google's page, and I leave it out of the Gemini figures (Mechanism 12).

© 2026 jerrywlambert.com

JERRYWLAMBERT.COM

Get the Next Analysis First

I publish what AI models cost, how vendors bill and what changed overnight. Subscribe to get it in your inbox.

Subscribe

https://jerrywlambert.com/newsletter/

Jerry W. Lambert, Sr.
jerrywlambert.com · Articles · Briefings · Writing · LinkedIn
© 2026 Jerry W. Lambert Sr.