

Jerry W. Lambert's Overnight Model Brief

Jerry W. Lambert Sr. · Saturday, 5 September 2026 · America/Chicago

OpenAI cleared the GPT-6 Astra Trusted Access banner. The price stayed put.

Astra is still a separate product — not a rename of Sol.

Subscriber cut · Newsletter & public archive

Executive summary

OpenAI took down the Trusted Access rollout notice for GPT-6 Astra on its public pricing page. The price did not move. Astra is still \$10 in and \$50 out for standard short context, against Sol at a \$4 / \$20 promo. Clearing a banner is not the same as opening the model to every plan. Treat Astra as a separate product, not a rename of Sol. Elsewhere the board was quiet on dollars. Watch the September clocks: Zhipu's Flash promo ends the 9th, Amazon Nova Premier ends the 14th, and Perplexity Sonar ends the 27th.

What moved

Window is Friday morning through Saturday morning, America/Chicago.

On Friday, OpenAI's pricing page said Astra was rolling out for enterprises in Trusted Access, with Plus, Pro, Business, and Enterprise coming in the coming days. That sentence is gone today. Standard short is still \$10 in / \$1 cached / \$12.50 cache writes / \$50 out per million tokens. Long context is still \$20 / \$2 / \$25 / \$75. Sol's short-context promo is still \$4 in / \$20 out through at least 21 November. Removing the banner does not mean Astra is generally available on every plan. It does not mean Sol got replaced.

No other major vendor rewrote a published token price overnight.

Prices that held

Anthropic Opus 5 stays \$5 in / \$25 out. Fable 5.1 stays \$10 / \$50. Google Gemini 3.8 Flash intro stays \$0.75 / \$3.75 through year end, then steps up. Google Lyria 3.5 stays \$0.08 per song. xAI Grok 4.6 stays \$2 / \$6 under 200k context. Zhipu GLM-5.3 stays \$1.4 / \$4.4. MiniMax M3 stays \$0.30 / \$1.20 at or under 512k. Alibaba's Qwen 3.8 sibling in Singapore stays \$2 / \$6.

Around the market

This week OpenAI put GPT-6 Astra on the flagship table at roughly 2.5× Sol's promo input rate, with Responses API tool calling and mid-turn steering. Google shipped Gemini 3.8 Flash at Flash intro dollars and Lyria 3.5 as a music product, not a chat substitute. Anthropic's Fable 5.1 remains the \$10 / \$50 frontier line beside Opus 5 at \$5 / \$25. Meta still shows Muse Spark 1.3 by name without a public dollar rate. Moonshot retired kimi-k2.5 at the end of August. Traffic to that name should stop, not get forwarded to kimi-k3. DeepSeek, Mistral, Cohere, Perplexity, Amazon, MiniMax, and Zhipu left published rates alone overnight. Microsoft did not publish a usable AI hub update in this window.

What to watch

- Zhipu's GLM Flash half-off promo ends 9 September at 11:00am Central.
- Amazon Nova Premier leaves support on 14 September.
- Perplexity Sonar leaves support on 27 September.
- Google still has no plain billed product named Gemini 3.1 Flash. Do not invent a substitute from 3.6, 3.7, or 3.8 Flash.
- Astra is not Sol. Do not route Sol work onto Astra because a banner disappeared.

Published list prices for Sol and Sonnet 5 still sit above the street promos many buyers actually see (\$4 / \$20 and \$2 / \$10). Read the promo line, not only the list line.

Bottom line

The overnight move was access messaging on Astra, not a price cut and not a forced upgrade path. Hold Astra, Lyria, Gemini 3.8 Flash, Fable 5.1, and the Qwen 0902 snapshot as separate products until you choose to adopt them. Keep retired Moonshot names out of production traffic.

Prices are per million tokens unless the sentence says otherwise.

Get Jerry W. Lambert's Overnight Model Brief by email

<https://jerrywlambert.com/newsletter/>

[Unsubscribe](#)